

CS 8395-03
Automatic
Machine Learning

Slides adapted from the Advanced Course on Data Science & Machine Learning

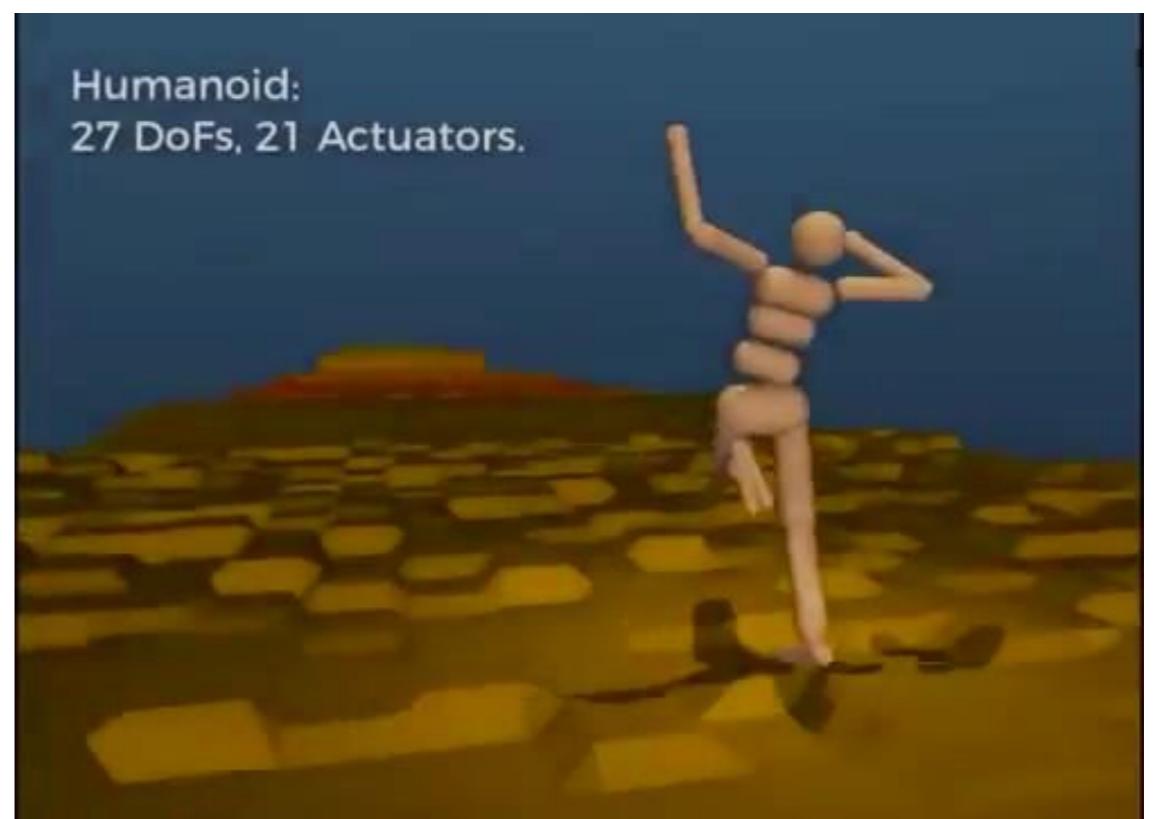
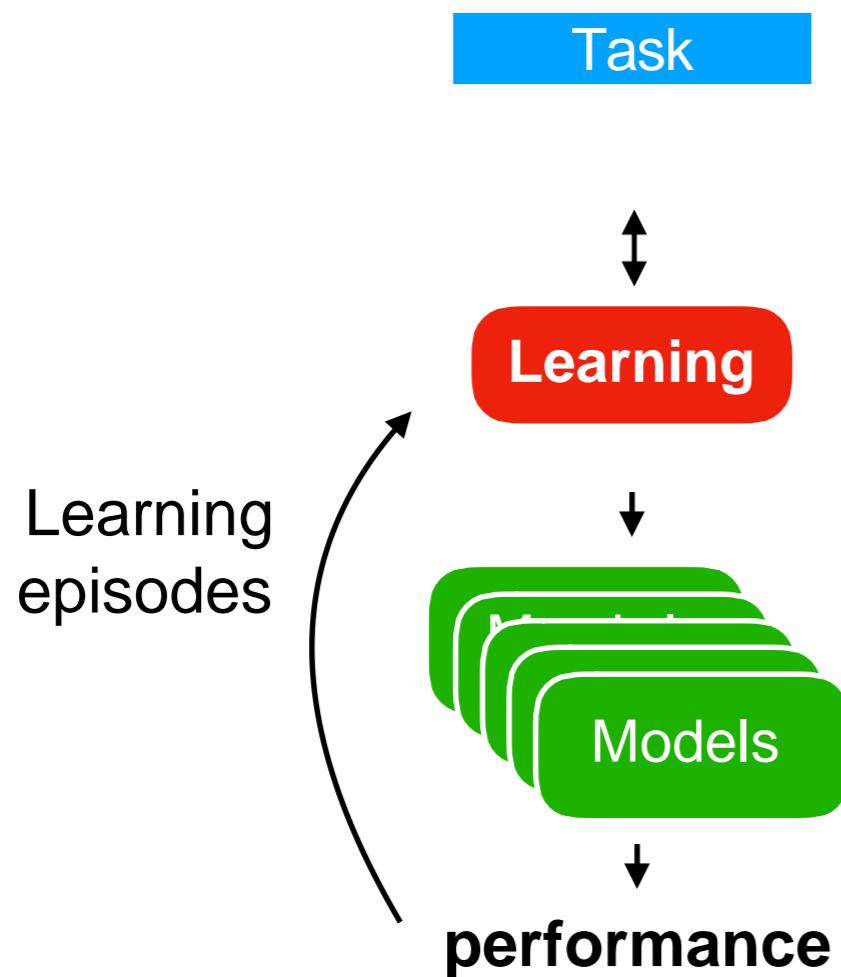
**Joaquin Vanschoren
Siena, 2019**

<http://bit.ly/openml>

AutoML

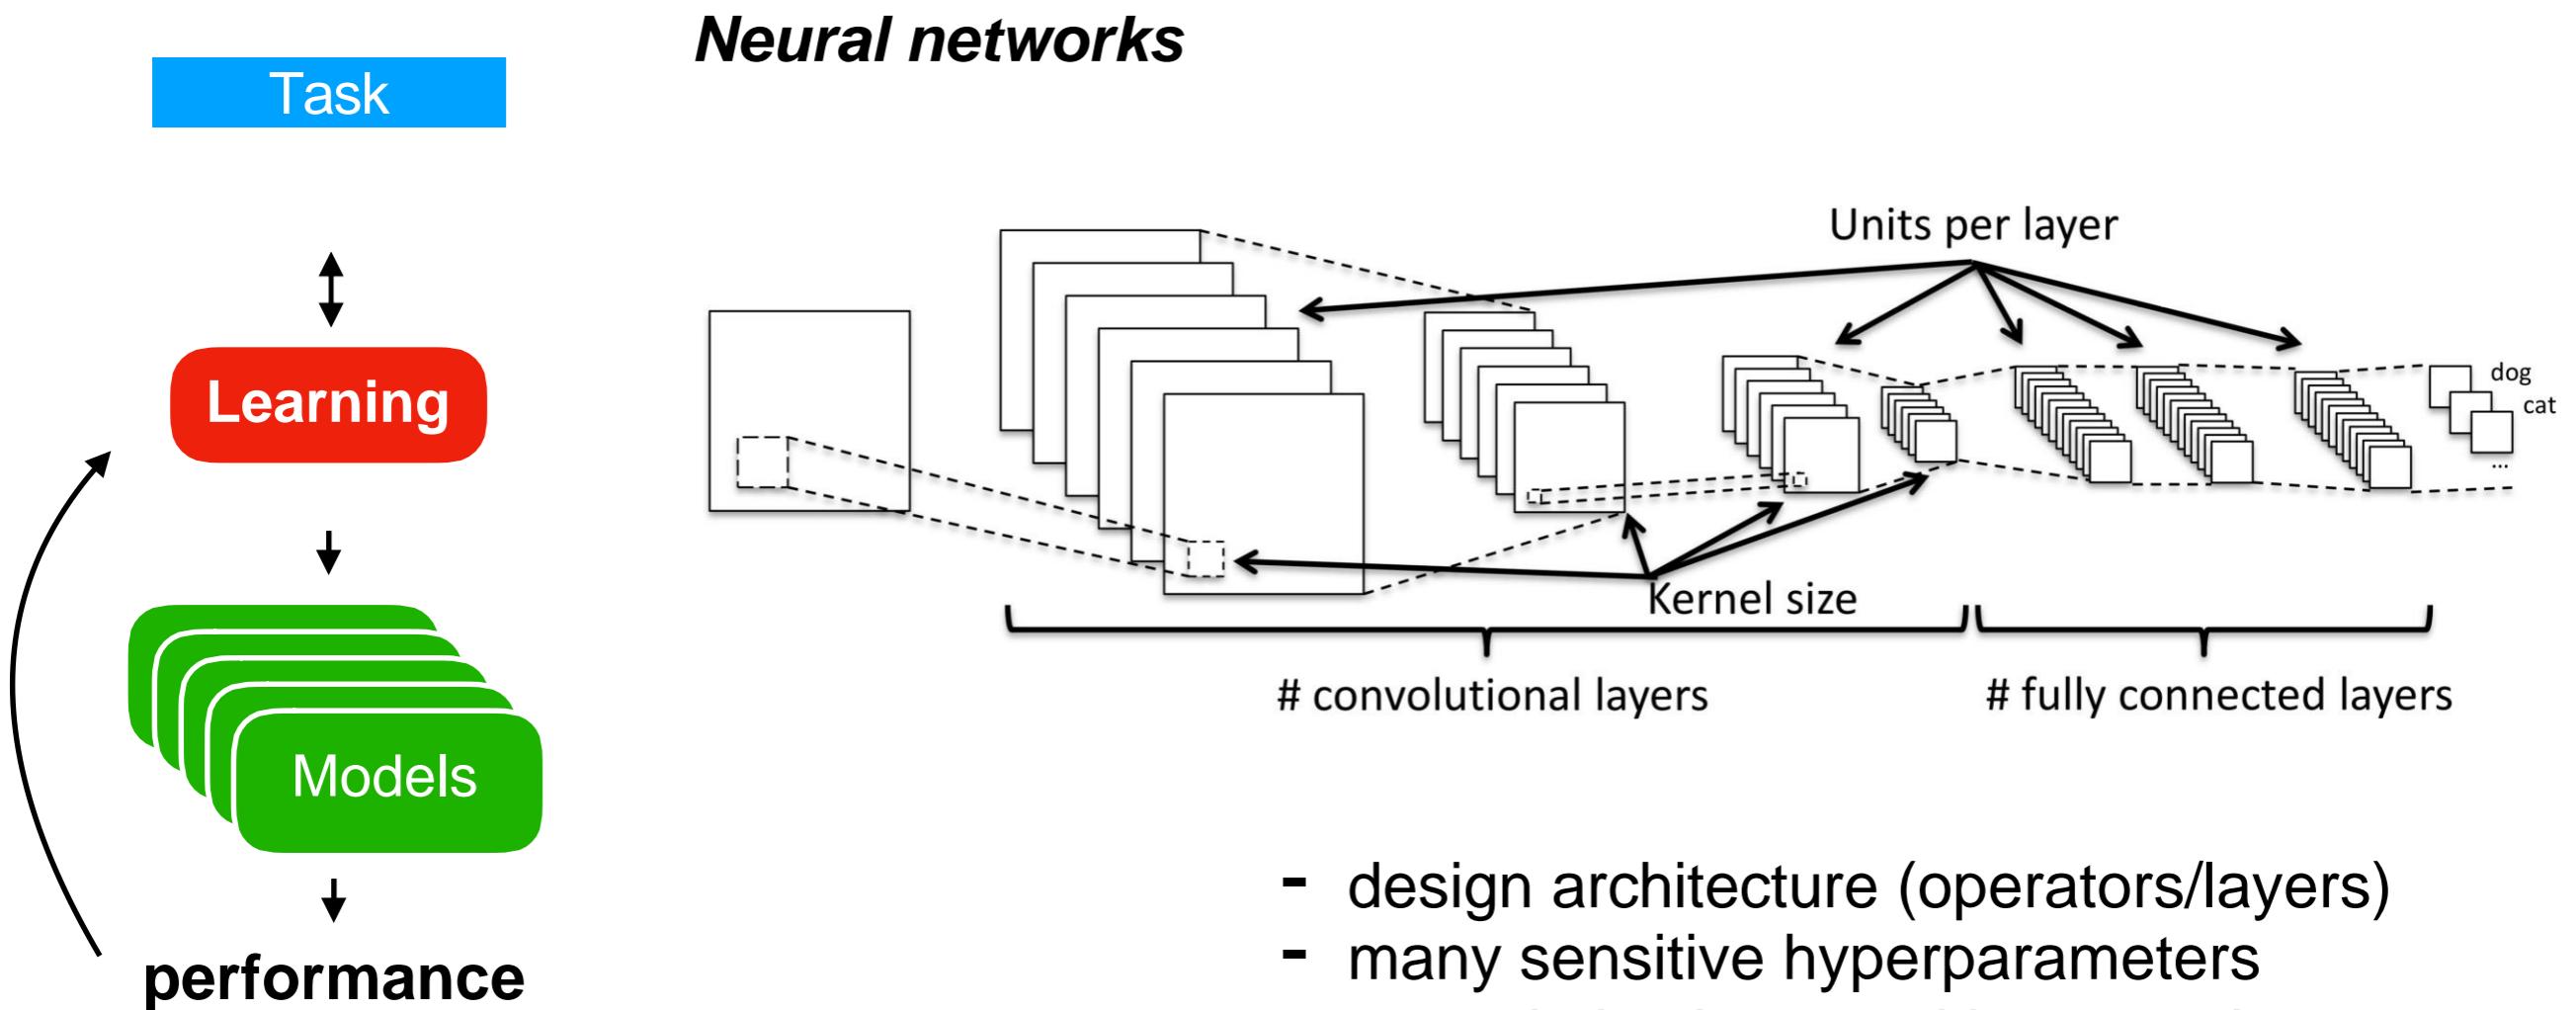
Learning takes experimentation

It can take a *lot* of trial and error



Learning to learn takes experimentation

Building machine learning systems requires a *lot* of expertise and trials



- design architecture (operators/layers)
- many sensitive hyperparameters
 - optimization algorithm, learning rate, ...
 - regularization, dropout, ...
 - batch size, epochs, early stopping, ...
 - data augmentation, ...

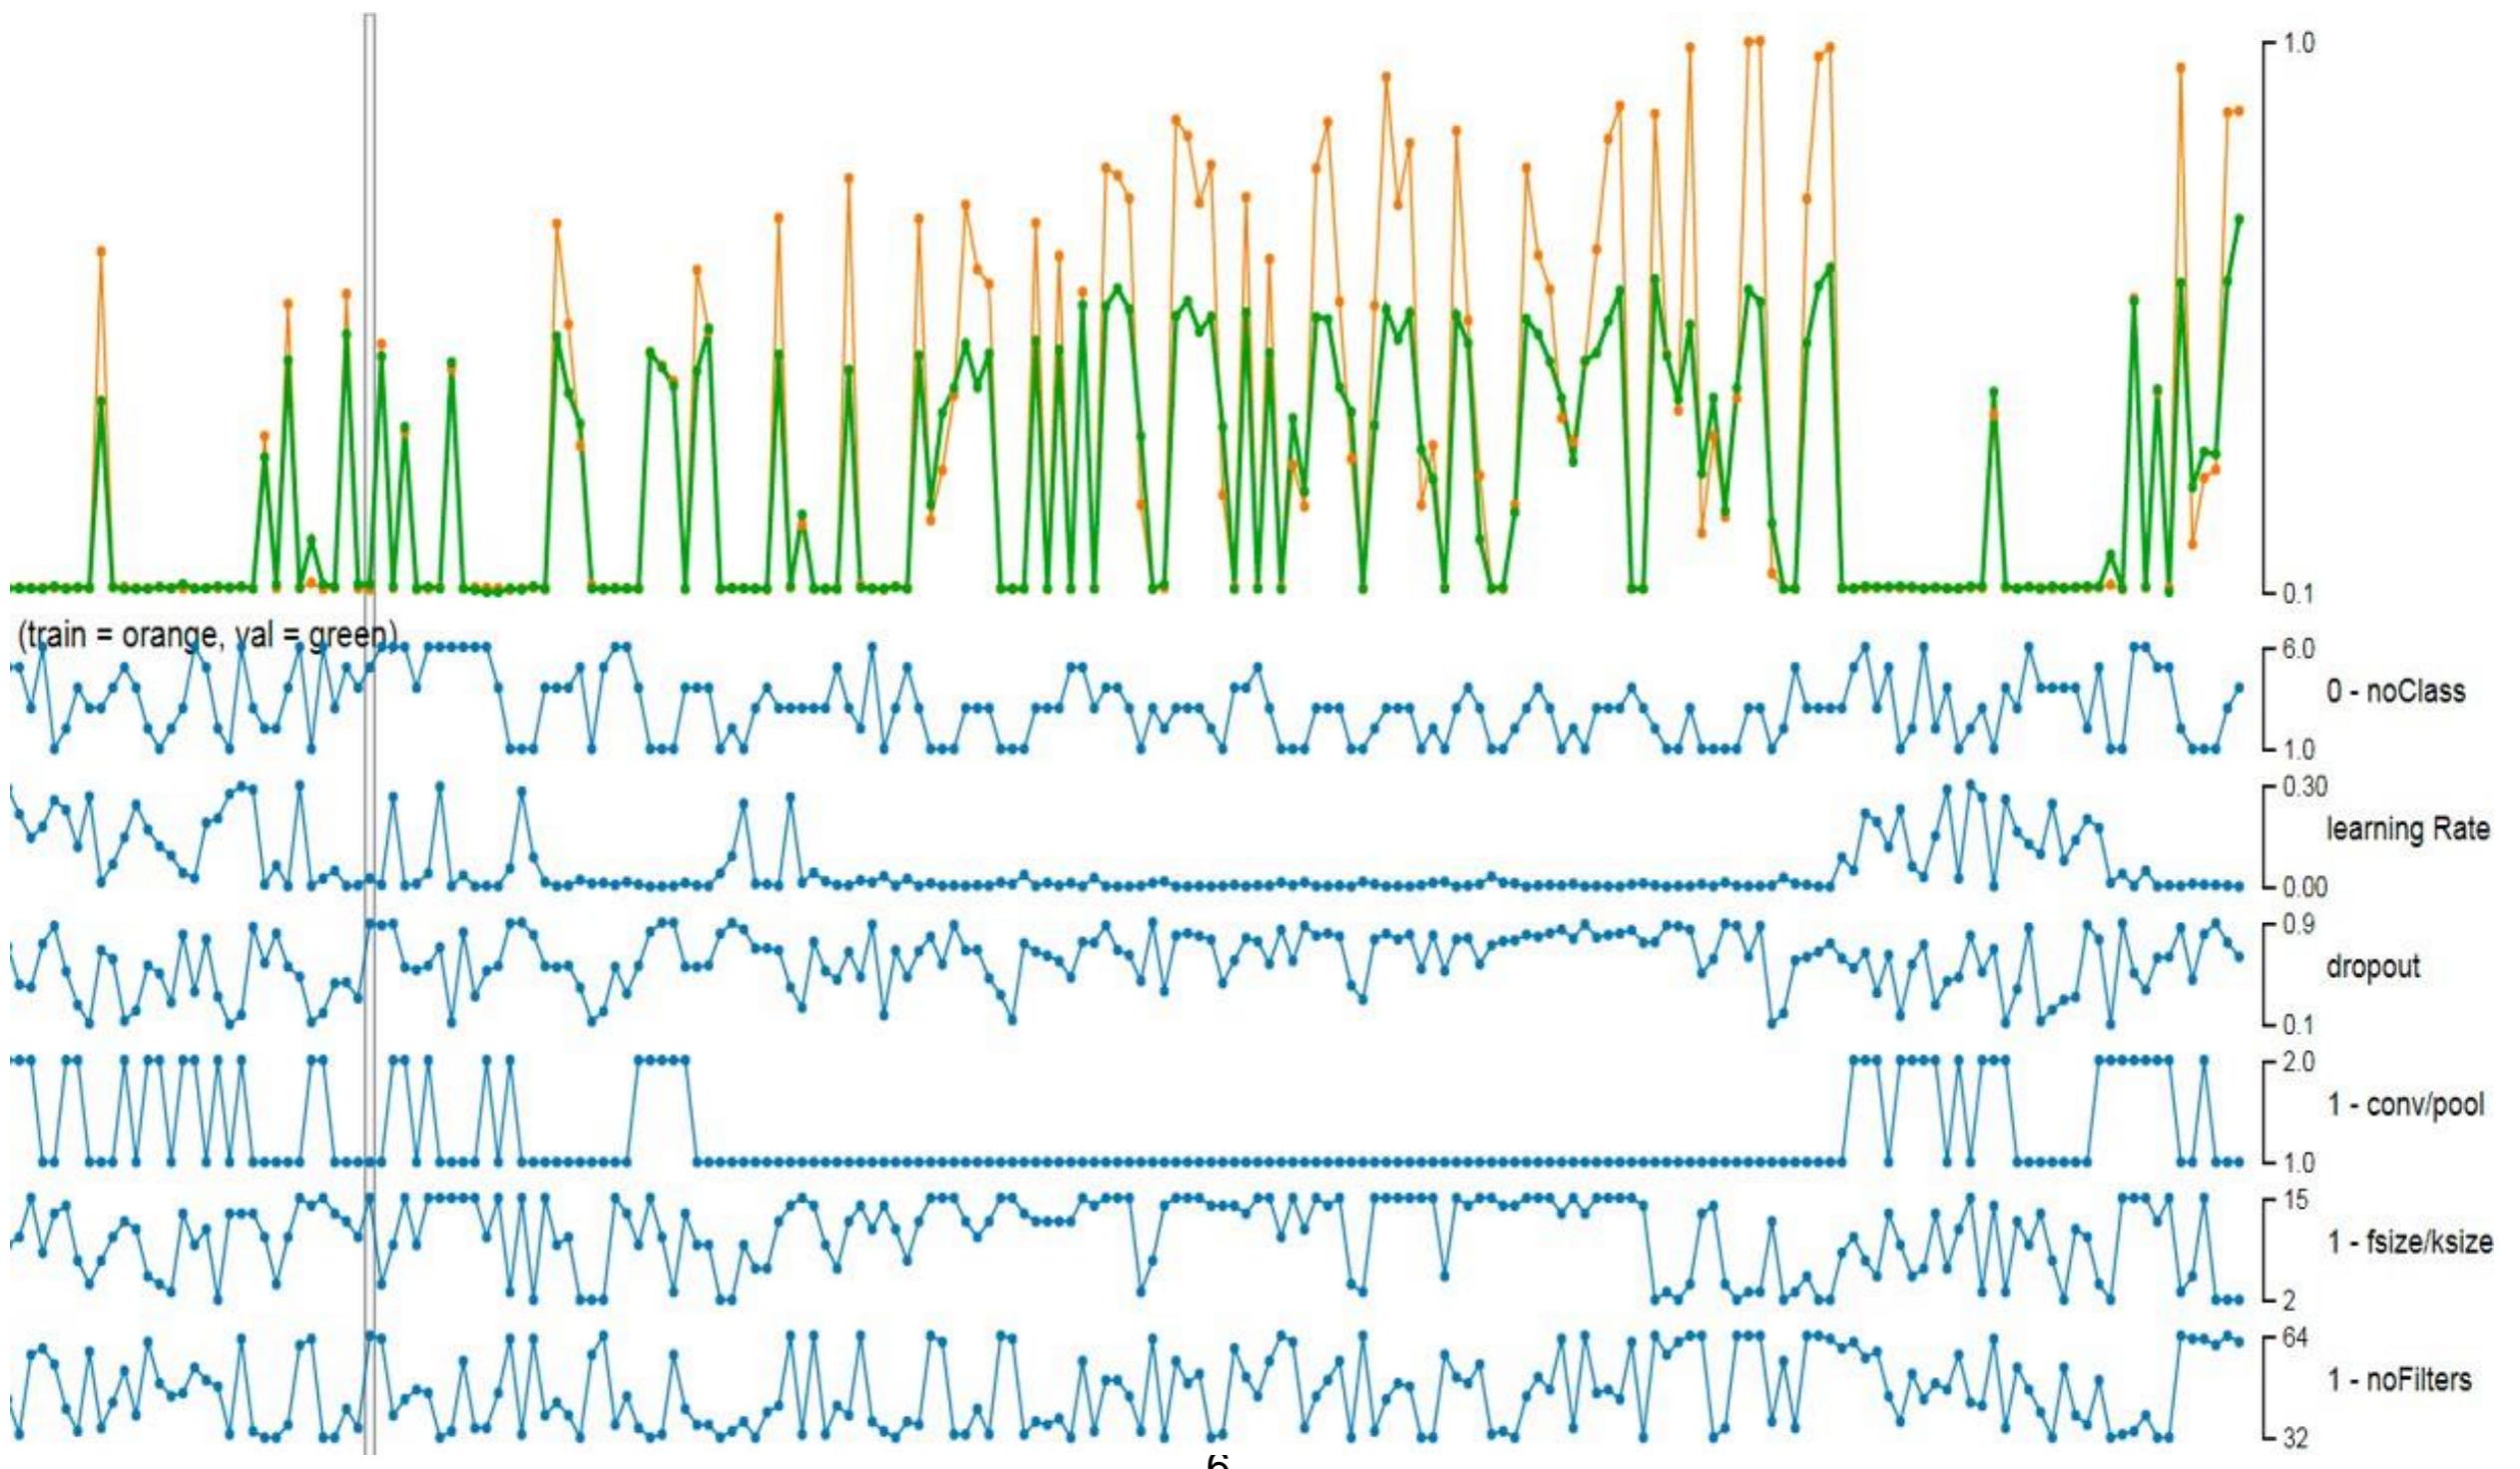
Hyperparameters

Every design decision made by the *user (architecture, operators, tuning, ...)*

	Name	Range	Default	log scale	Type	Conditional
Network hyperparameters	batch size	[32, 4096]	32	✓	float	-
	number of updates	[50, 2500]	200	✓	int	-
	number of layers	[1, 6]	1	-	int	-
	learning rate	[10^{-6} , 1.0]	10^{-2}	✓	float	-
	L_2 regularization	[10^{-7} , 10^{-2}]	10^{-4}	✓	float	-
	dropout output layer	[0.0, 0.99]	0.5	✓	float	-
	solver type	{SGD, Momentum, Adam, Adadelta, Adagrad, smorm, Nesterov }	smorm3s	-	cat	-
Conditioned on solver type	lr-policy	{Fixed, Inv, Exp, Step}	fixed	-	cat	-
	β_1	[10^{-4} , 10^{-1}]	10^{-1}	✓	float	✓
	β_2	[10^{-4} , 10^{-1}]	10^{-1}	✓	float	✓
	ρ	[0.05, 0.99]	0.95	✓	float	✓
Conditioned on lr-policy	momentum	[0.3, 0.999]	0.9	✓	float	✓
	γ	[10^{-3} , 10^{-1}]	10^{-2}	✓	float	✓
	k	[0.0, 1.0]	0.5	-	float	✓
Per-layer hyperparameters	s	[2, 20]	2	-	int	✓
	activation-type	{Sigmoid, TanH, ScaledTanH, ELU, ReLU, Leaky, Linear}	ReLU	-	cat	✓
	number of units	[64, 4096]	128	✓	int	✓
	dropout in layer	[0.0, 0.99]	0.5	-	float	✓
	weight initialization	{Constant, Normal, Uniform, Glorot-Uniform, Glorot-Normal, He-Normal, He-Uniform, Orthogonal, Sparse}	He-Normal	-	cat	✓
	std. normal init.	[10^{-7} , 0.1]	0.0005	-	float	✓
	leakiness	[0.01, 0.99]	$\frac{1}{3}$	-	float	✓
	tanh scale in	[0.5, 1.0]	2/3	-	float	✓
	tanh scale out	[1.1, 3.0]	1.7159	✓	float	✓

Hyperparameters

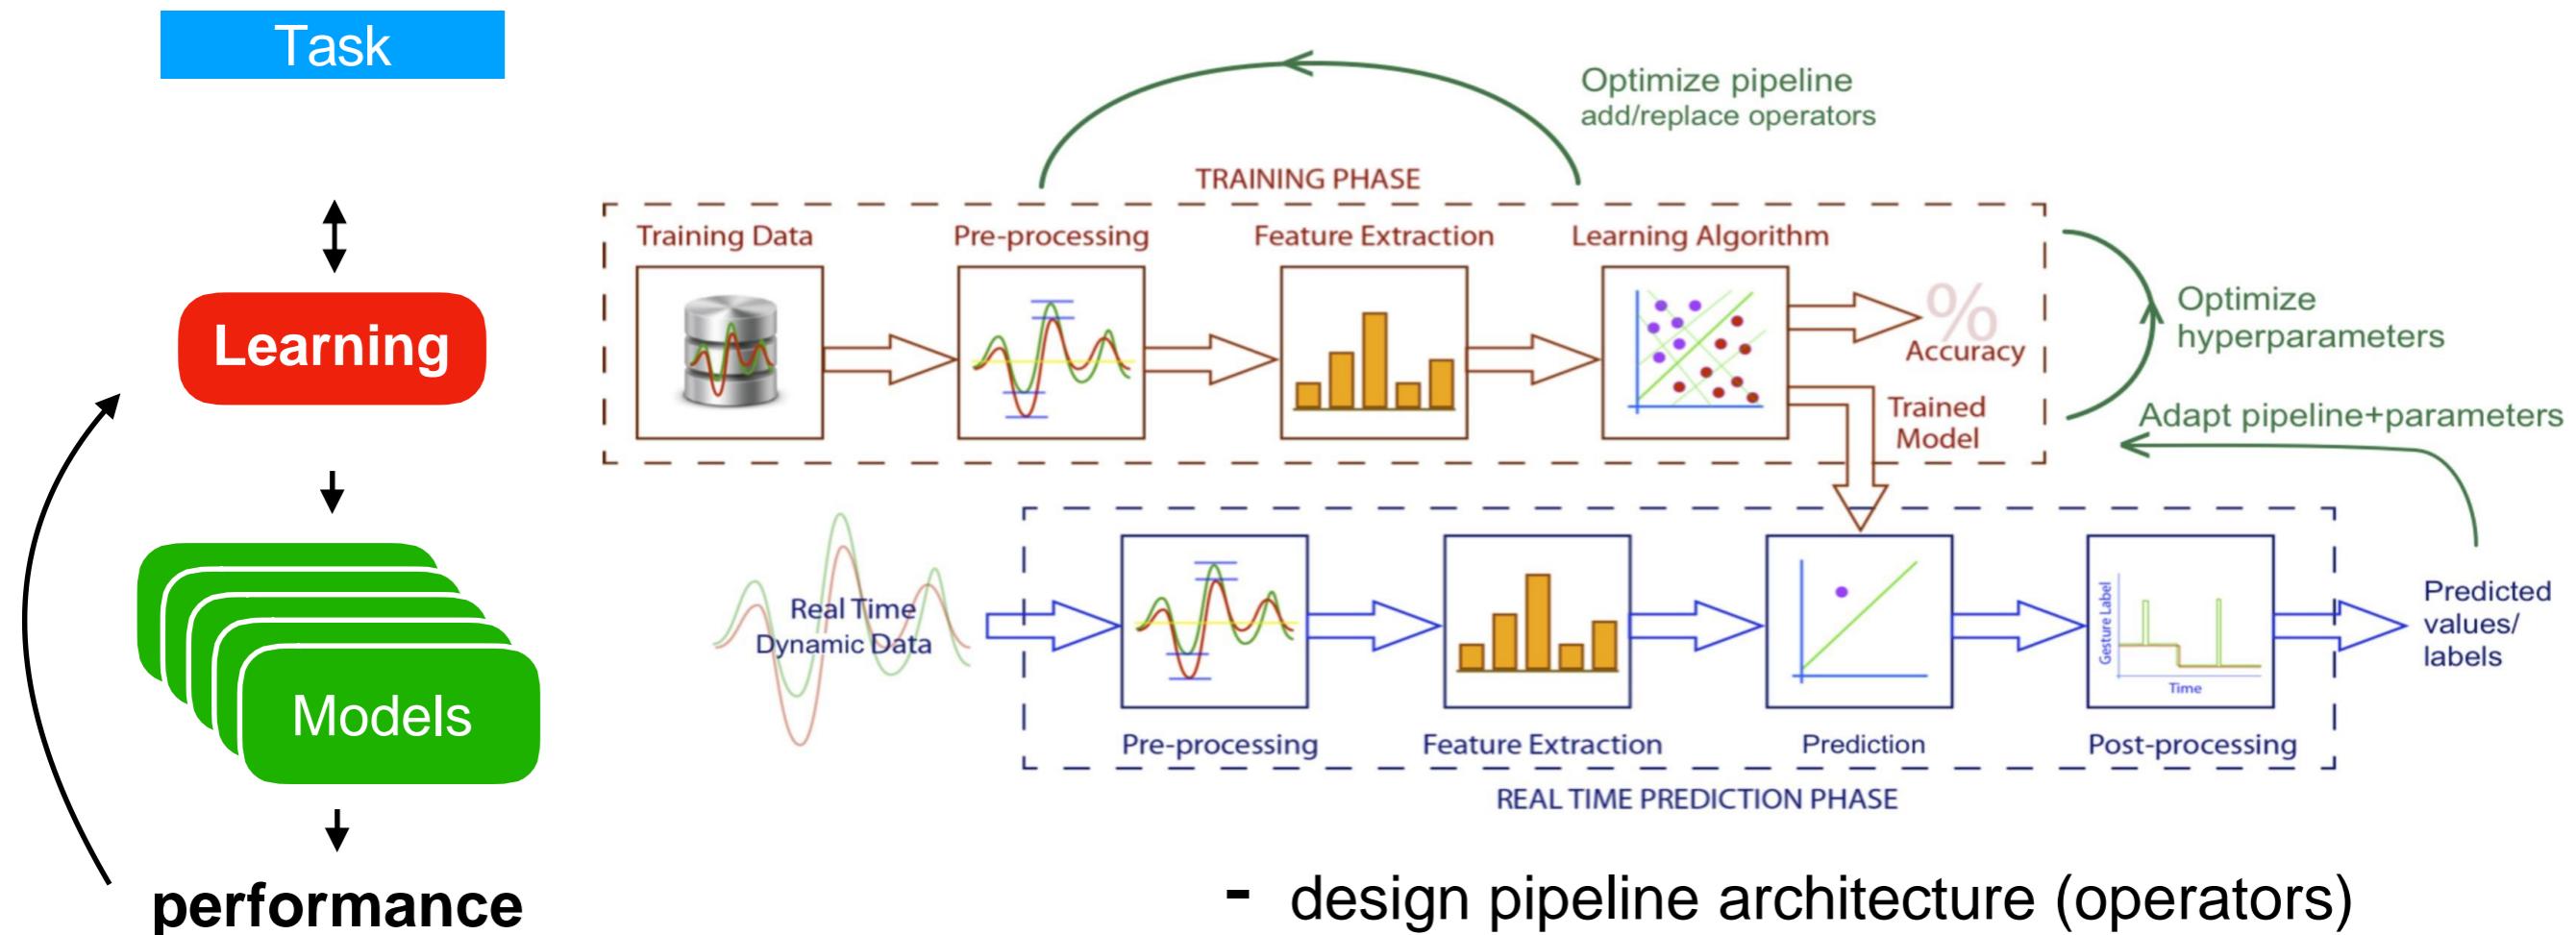
Can be very sensitive



Learning to learn takes experimentation

Building machine learning systems requires a lot of expertise and trials

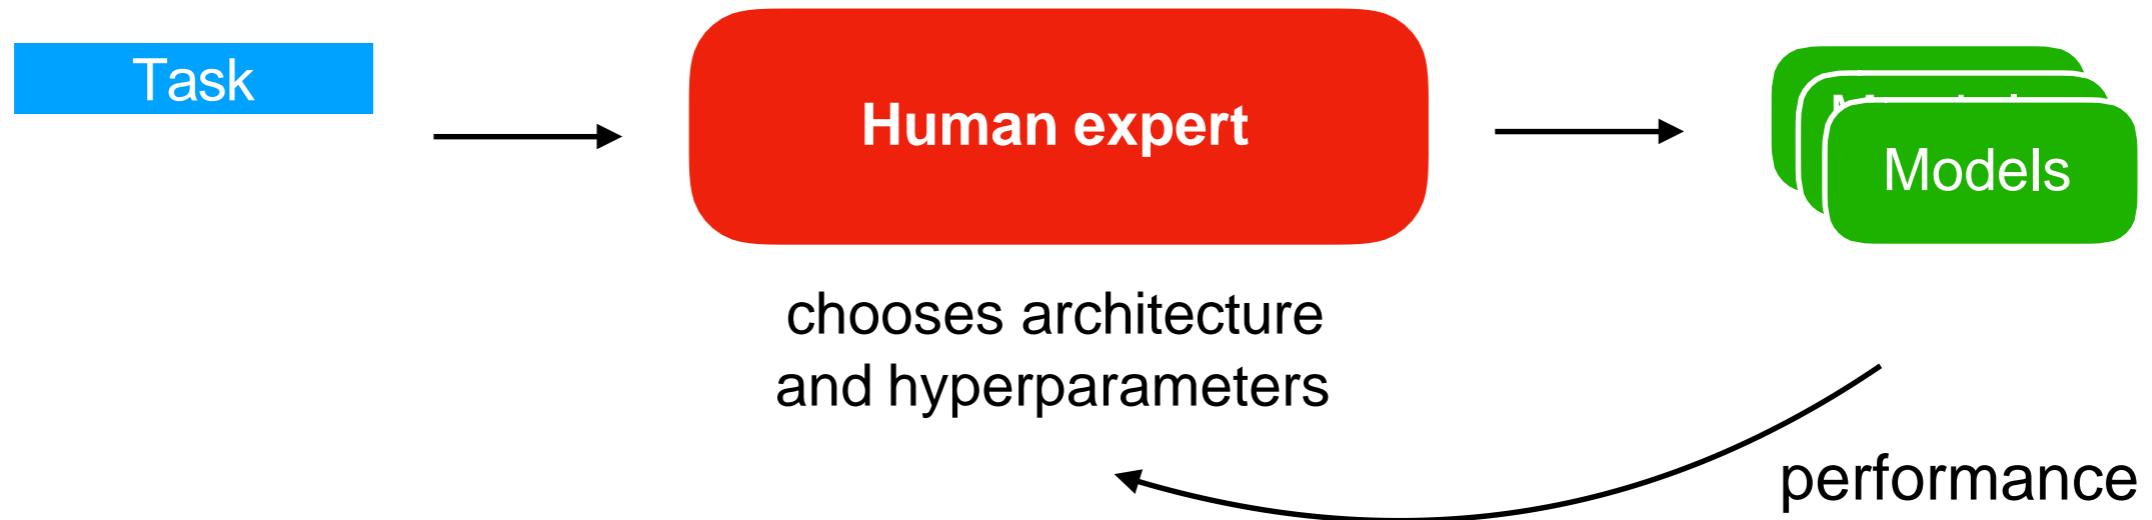
Machine learning pipelines



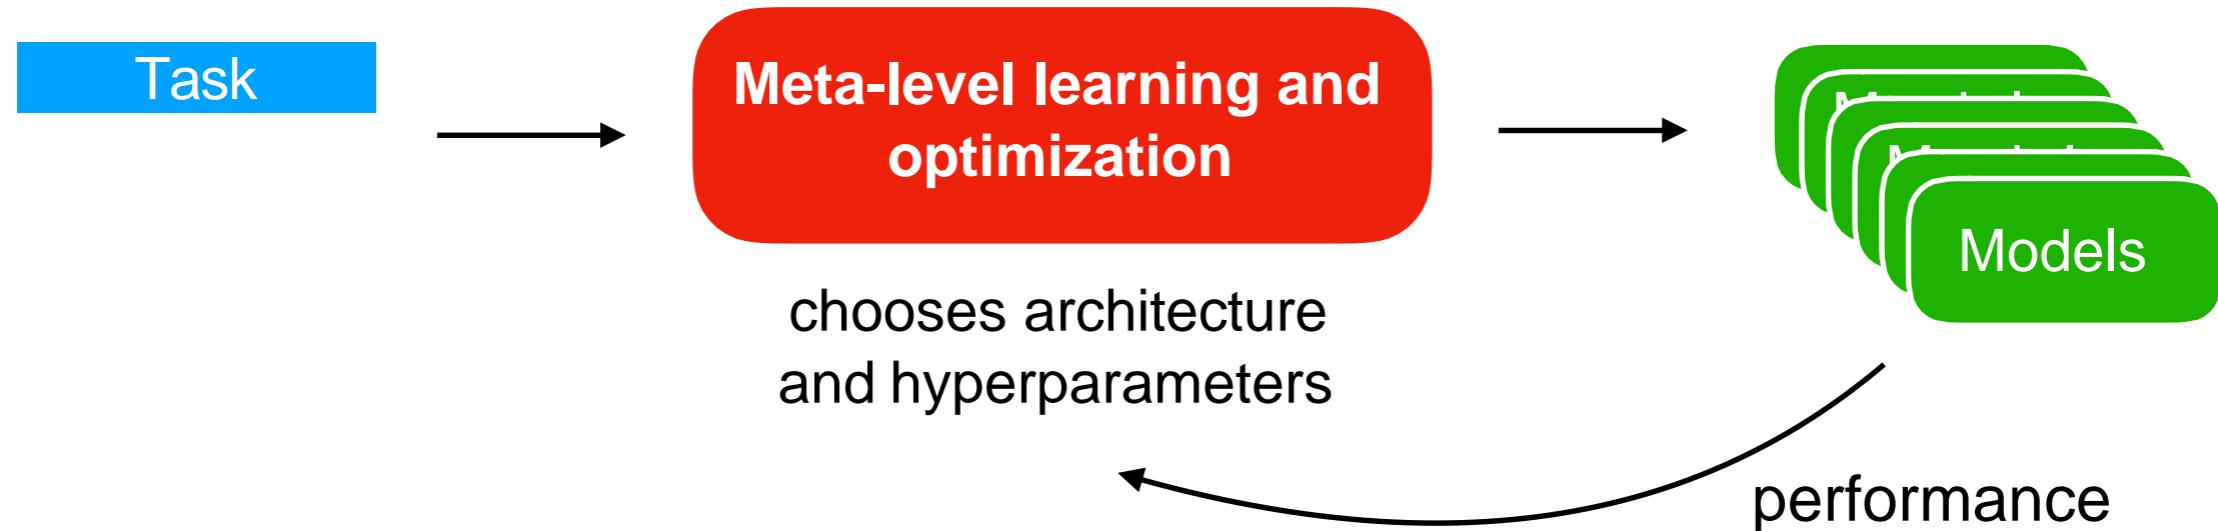
- design pipeline architecture (operators)
- clean, preprocess data (!)
- select and/or engineer features
- select and tune models
- adjust to evolving input data (concept drift),...

Automated machine learning

Current practice



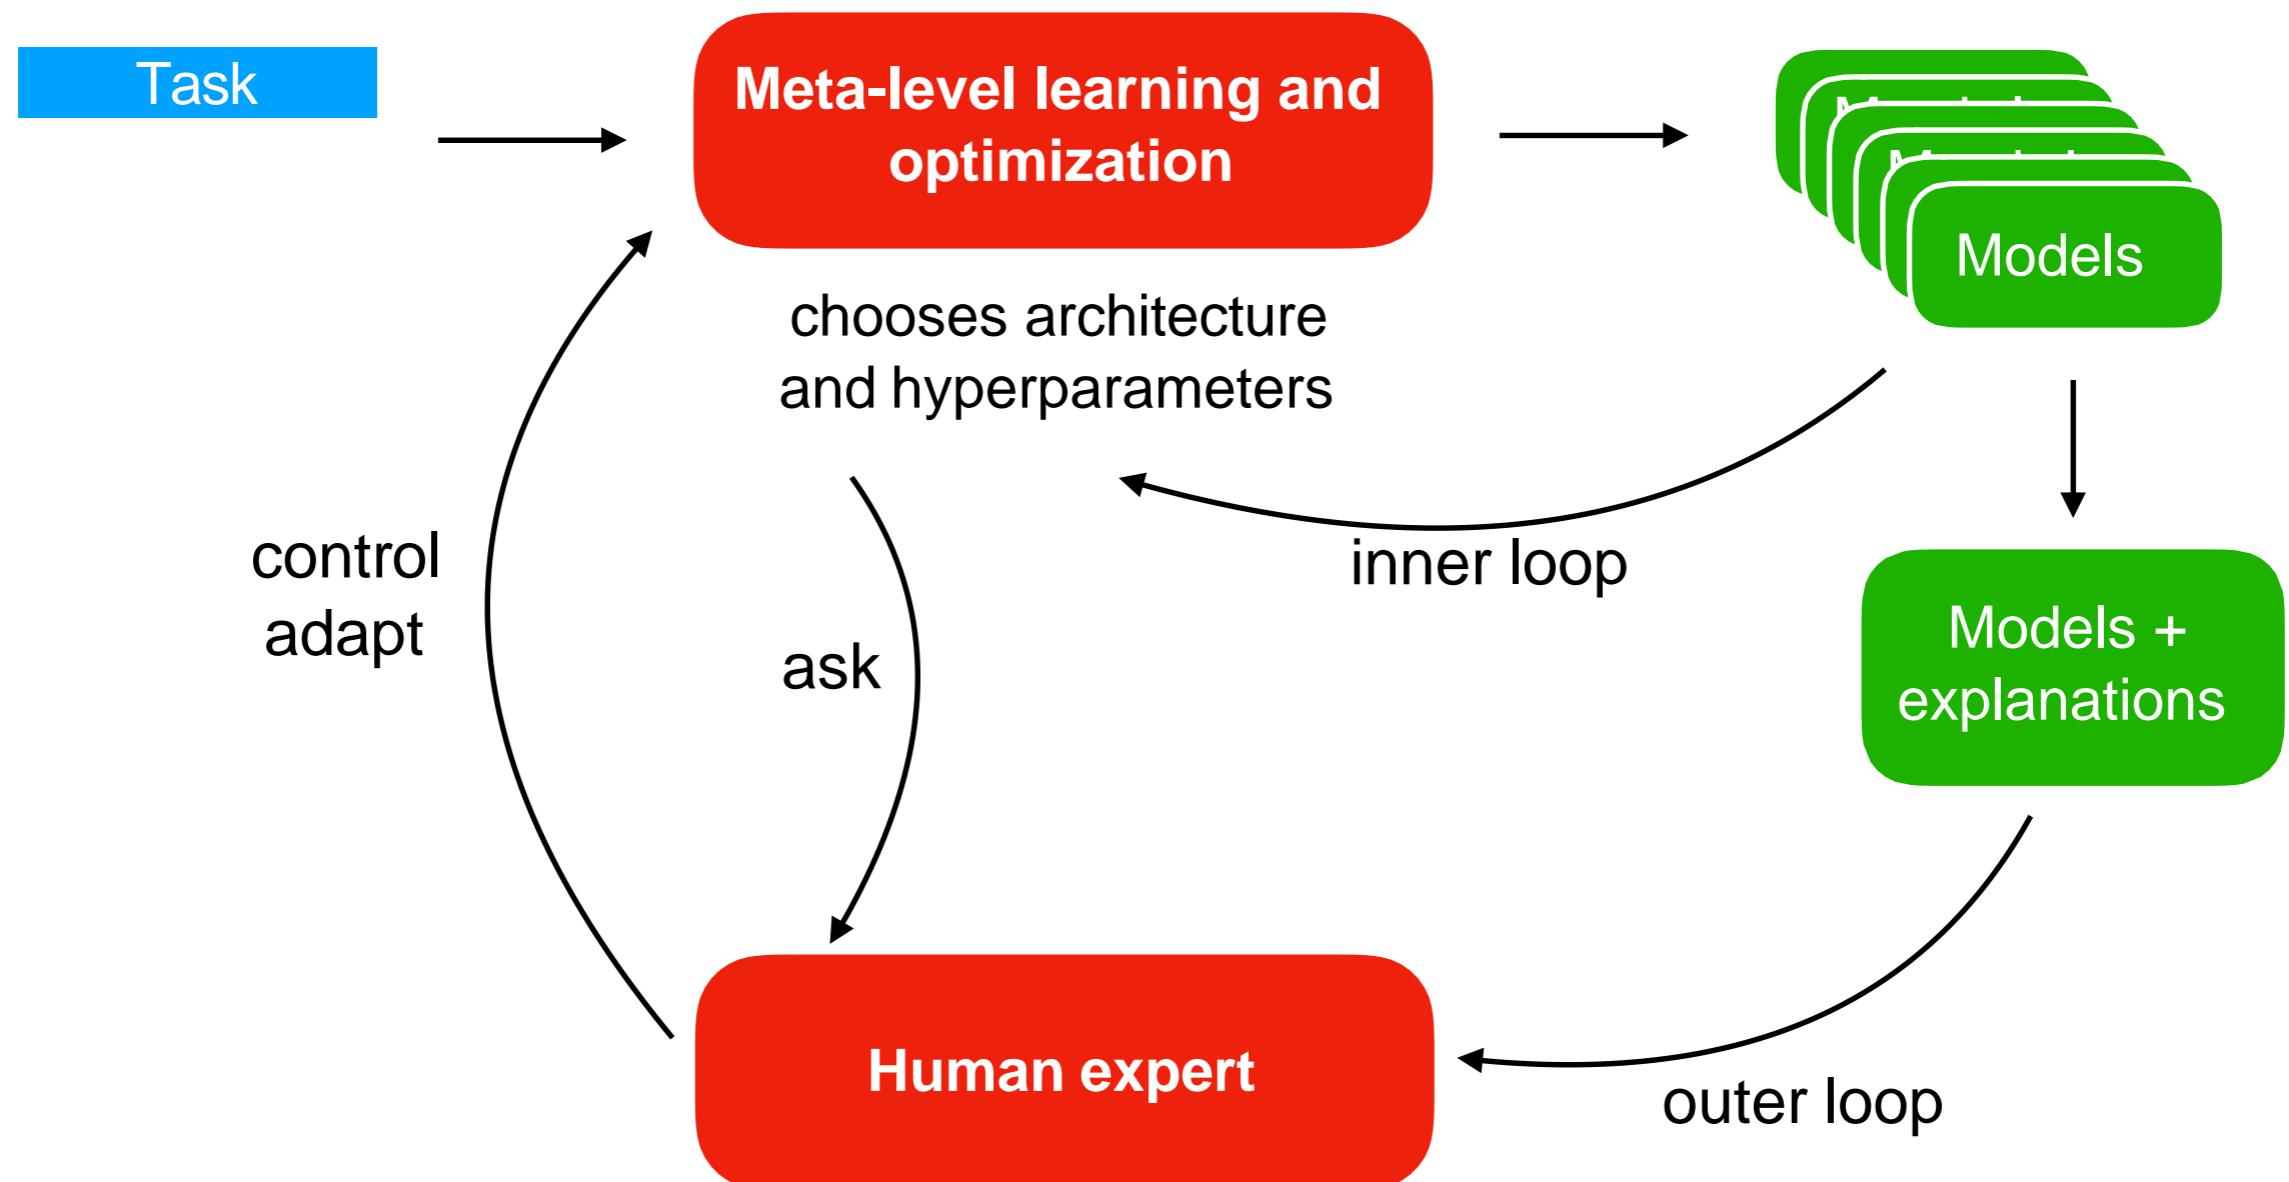
Replace manual model building by automation





Human-in-the-loop AutoML (semi-AutoML)

Domain knowledge and human expertise are very valuable
e.g. unknown unknowns, preference learning,...



Overview

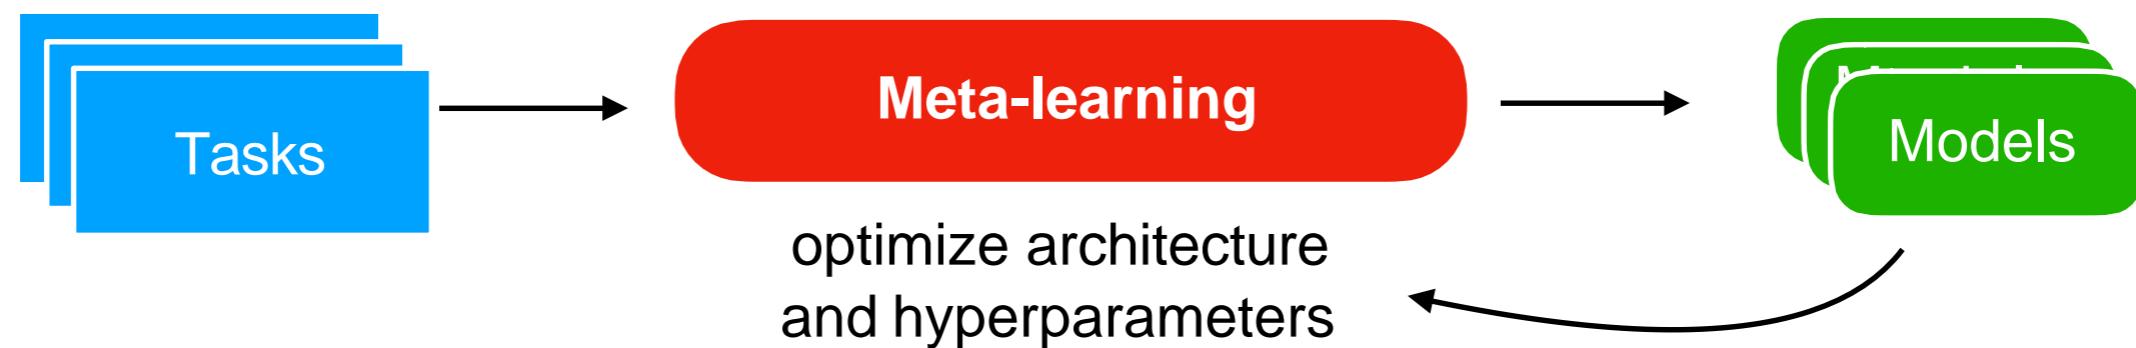
part AutoML introduction, promising optimization techniques

1



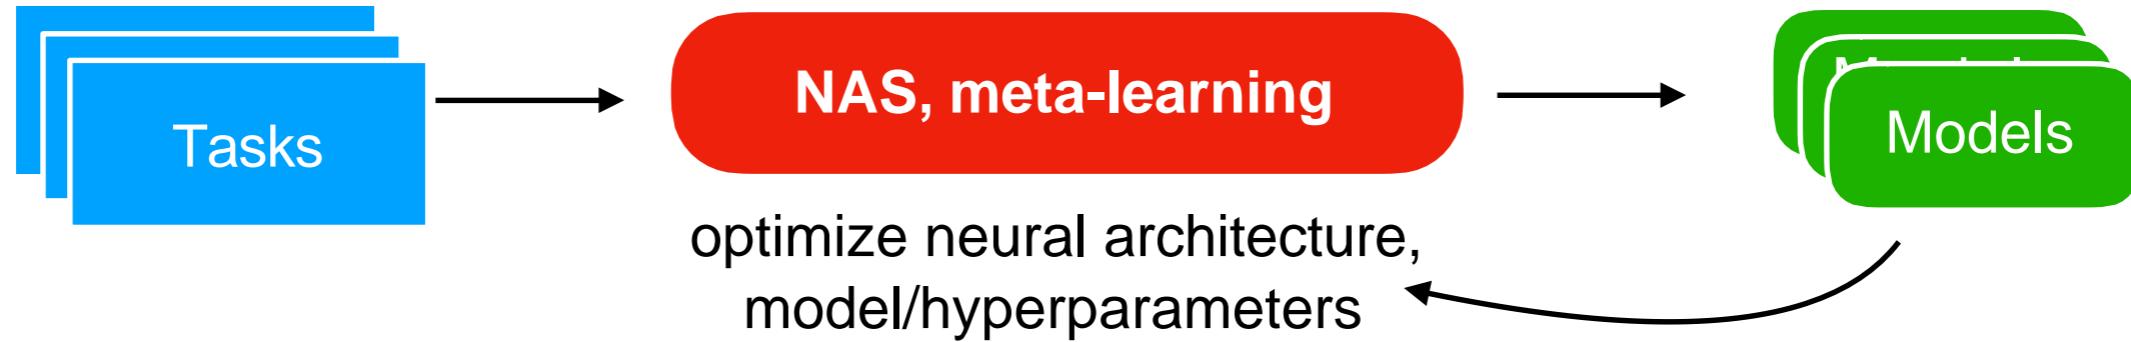
part Meta-learning

2



part Neural architecture search, meta-learning on neural nets

3



Overview

part AutoML introduction, optimization techniques

1



1. Problem definition
2. (Pipeline) architecture search
3. Optimization techniques
4. Performance improvements
5. Benchmarks

AutoML Definition

Let:

$\{A^{(1)}, A^{(2)}, \dots, A^{(n)}\}$ be a set of *algorithms* (operators)

$\Lambda^{(i)} = \lambda_1 \times \lambda_2 \times \dots \times \lambda_m$ be the *hyperparameter space* for $A^{(i)}$

d be the space of possible *architectures* of one or more algorithms

$\Lambda_d = \Lambda^{(1)} \times \Lambda^{(2)} \times \dots \times \Lambda^{(m)}$ its combined *configuration space*

$\lambda \in \Lambda_d$ a specific *configuration* (of architecture and hyperparameters)

$\mathcal{L}(\lambda, D_{train}, D_{valid})$ the loss of the model created by λ , trained on data D_{train} , and validated on data D_{valid}

Find the configuration that minimizes the expected loss on a dataset fij :

$$\lambda^* = \operatorname{argmin}_{\lambda \in \Lambda_d} \mathbb{E}_{(D_{train}, D_{valid}) \sim fij} \mathcal{L}(\lambda, D_{train}, D_{test})$$

Types of hyperparameters

- Continuous (e.g. learning rate, SVM_C,...)
- Integer (e.g. number of hidden units, number of boosting iterations,...)
- Categorical
 - e.g. choice of algorithm (SVM, RandomForest, Neural Net,...)
 - e.g. choose of operator (Convolution, MaxPooling, DropOut,...)
 - e.g. activation function (ReLU, Leaky ReLU, tanh,...)
- Conditional
 - e.g. SVM kernel if SVM is selected, kernel width if RBF kernel is selected
 - e.g. Convolution kernel size if Convolution layer is selected

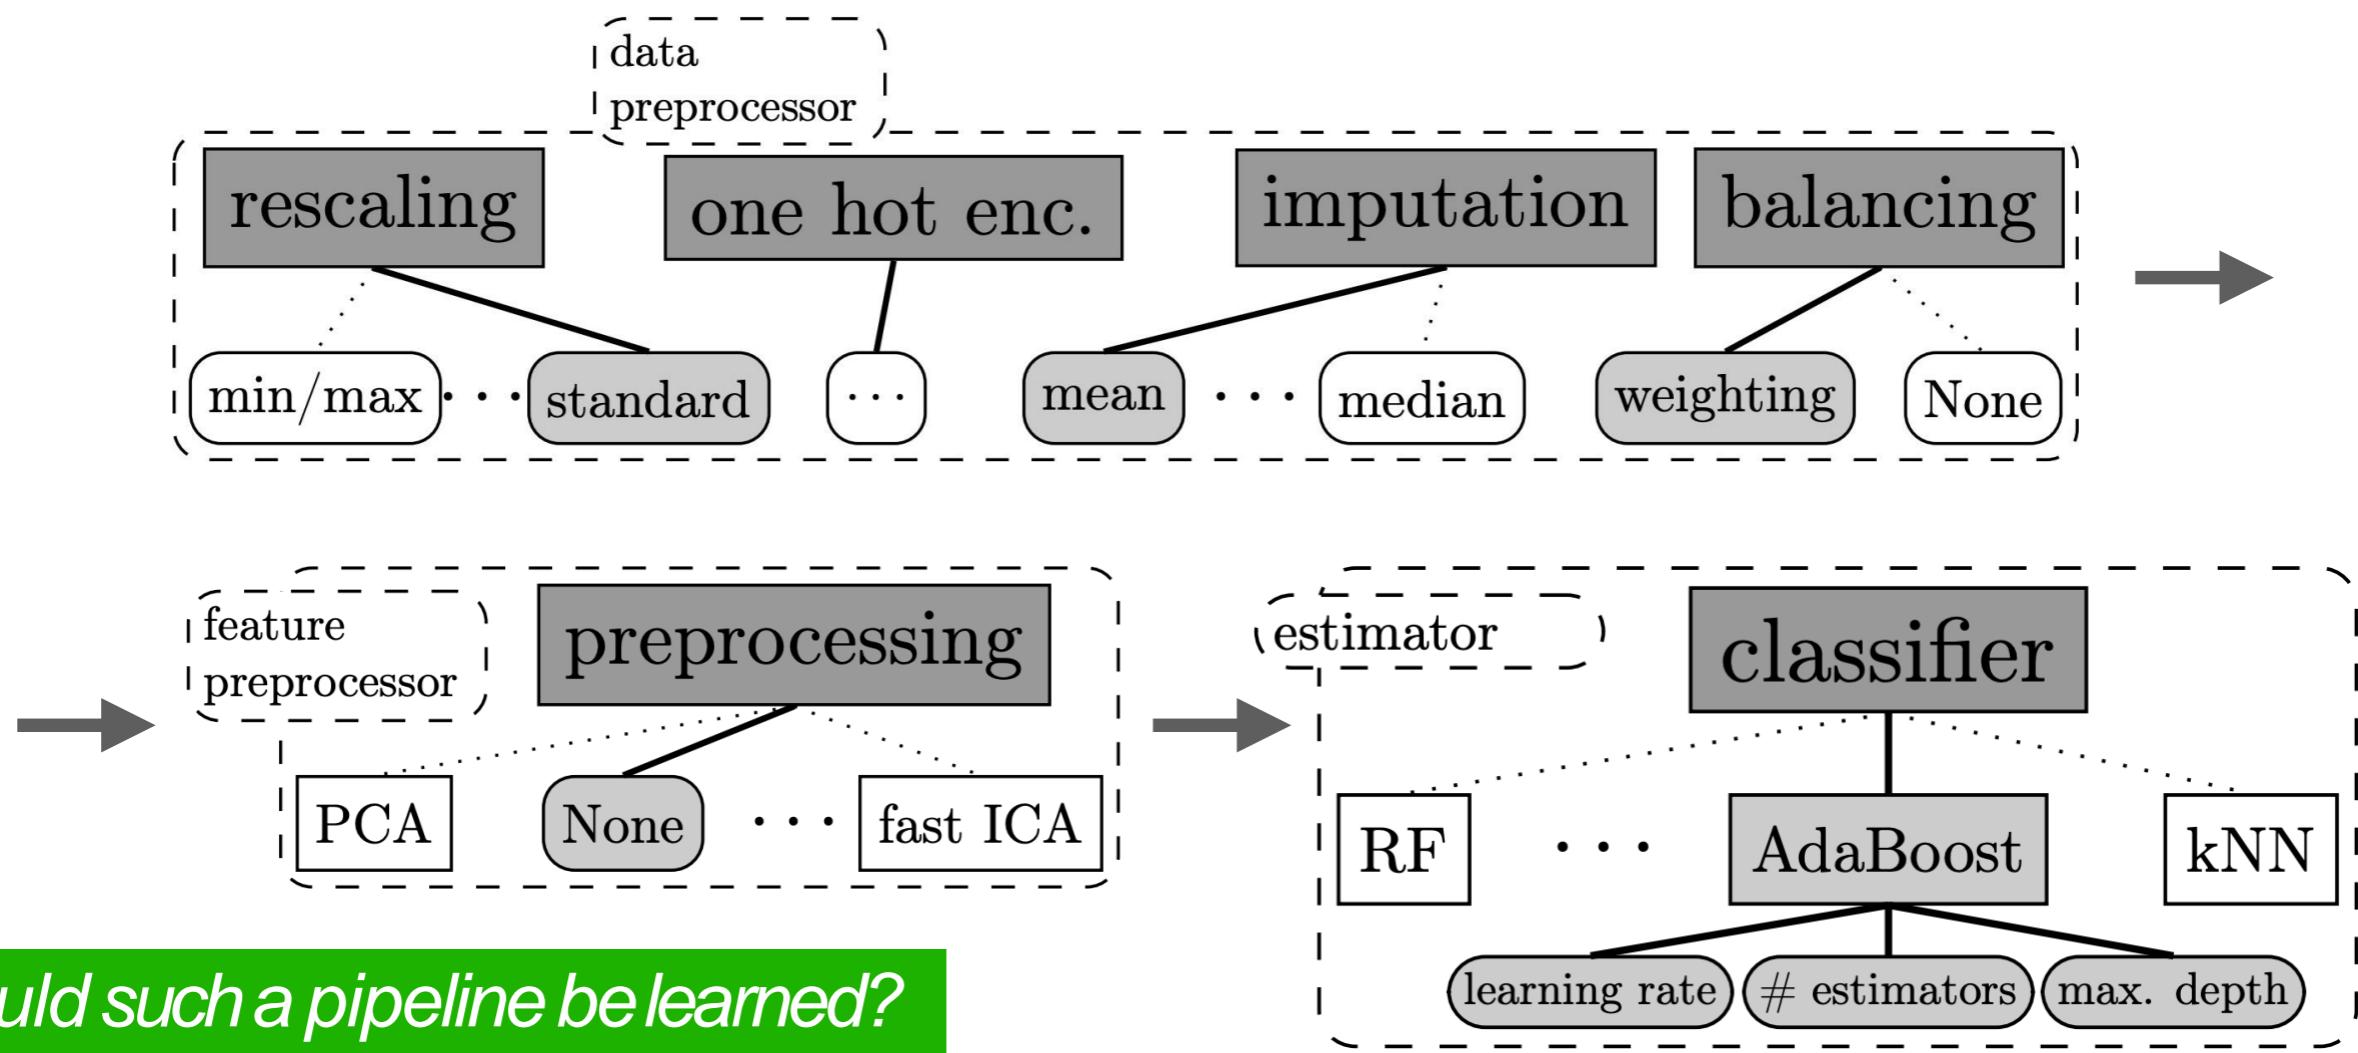
Architecture vs hyperparameters

- We can identify two subproblems:
 - **Architecture search**: search space of all possible architectures
 - **Pipelines**: Fixed predefined pipeline, grammars, genetic programming, planning, Monte-Carlo Tree Search
 - **Neural architectures**: See part 3
 - **Hyperparameter optimization**: optimize remaining hyperparameters
 - **Optimization**: grid/random search, Bayesian optimization, evolution, multi-armed bandits, gradient descent (only NNs)
 - **Meta-learning**: See part 2
- Can be solved *consecutively*, *simultaneously* or *interleaved*
- **Compositionality**: breaking down the learning process into smaller reusable tasks makes it easier, more transferable, more robust

Pipeline search: fixed pipelines

- Parameterize best-practice (linear) pipelines
 - Introduce conditional hyperparameters $\lambda_r \in \{A^{(1)}, \dots, A^{(k)}, \emptyset\}$
 - Combined Algorithm Selection and Hyperparameter optimization (CASH):

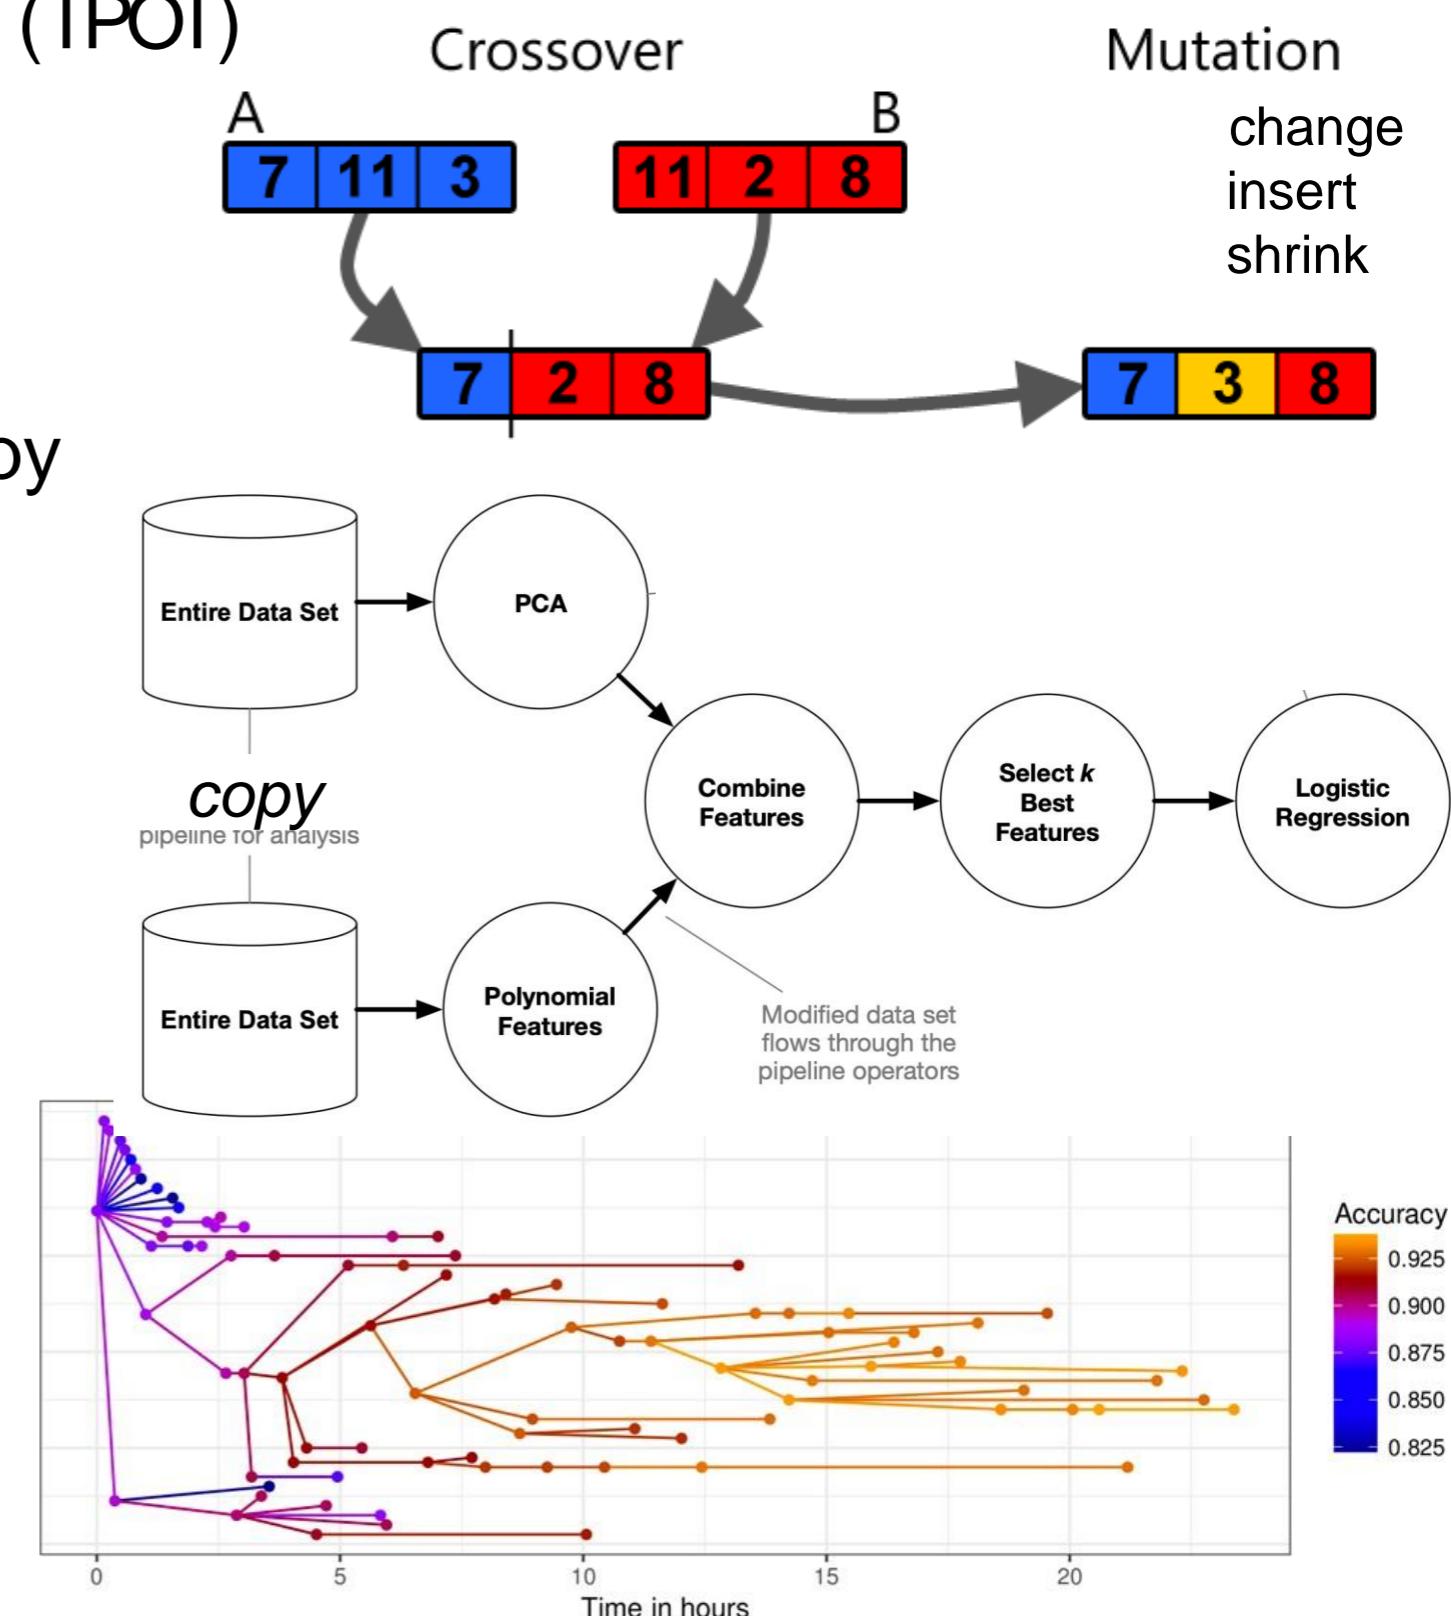
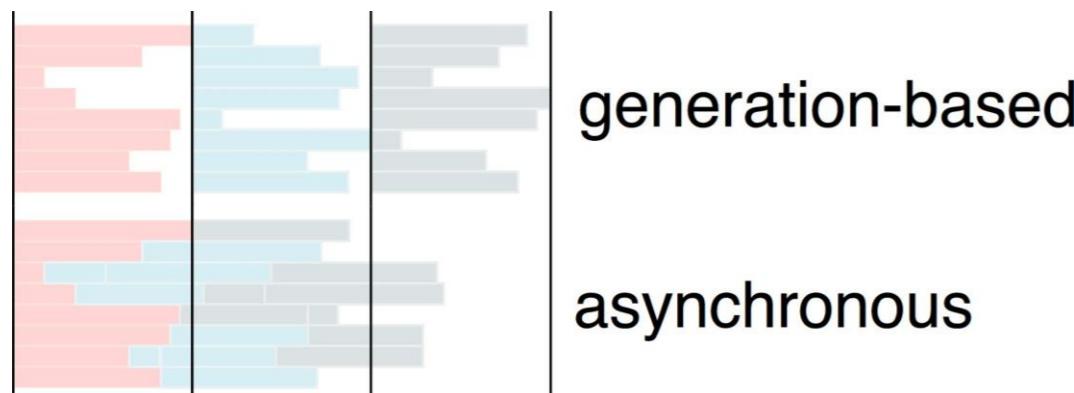
$$\Lambda_d = \Lambda^{(1)} \times \Lambda^{(2)} \times \dots \times \Lambda^{(m)} \times \lambda_r$$



Pipeline search: genetic programming

- Tree-based pipeline optimization (TPOT)
 - Start with random pipelines, best of every generation will cross-over or mutate
 - GP primitives include data copy and feature joins: trees
 - Multi-objective optimization: accurate but short
 - Easy to parallelize
- Asynchronous evolution (GAMA)

[Gijsbers, Vanschoren 2019](#)



Pipeline search: TPOT demo

```
from tpot import TPOTClassifier
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split

digits = load_digits()
X_train, X_test, y_train, y_test = train_test_split(digits.data, digits.target,
                                                    train_size=0.75, test_size=0.25)

tpot = TPOTClassifier(generations=5, population_size=50, verbosity=2, n_jobs=-1)
tpot.fit(X_train, y_train)

Optimization Progress:  0% |  0/300 [00:00<?, ?pipeline/s]

print(tpot.score(X_test, y_test))
```

Pipeline search: genetic programming

- Grammar-based genetic programming (RECIPE) [\[de Sa et al. 2017\]](#)
 - Encodes background knowledge, avoids invalid pipelines
 - Cross-over and mutation respect the grammar

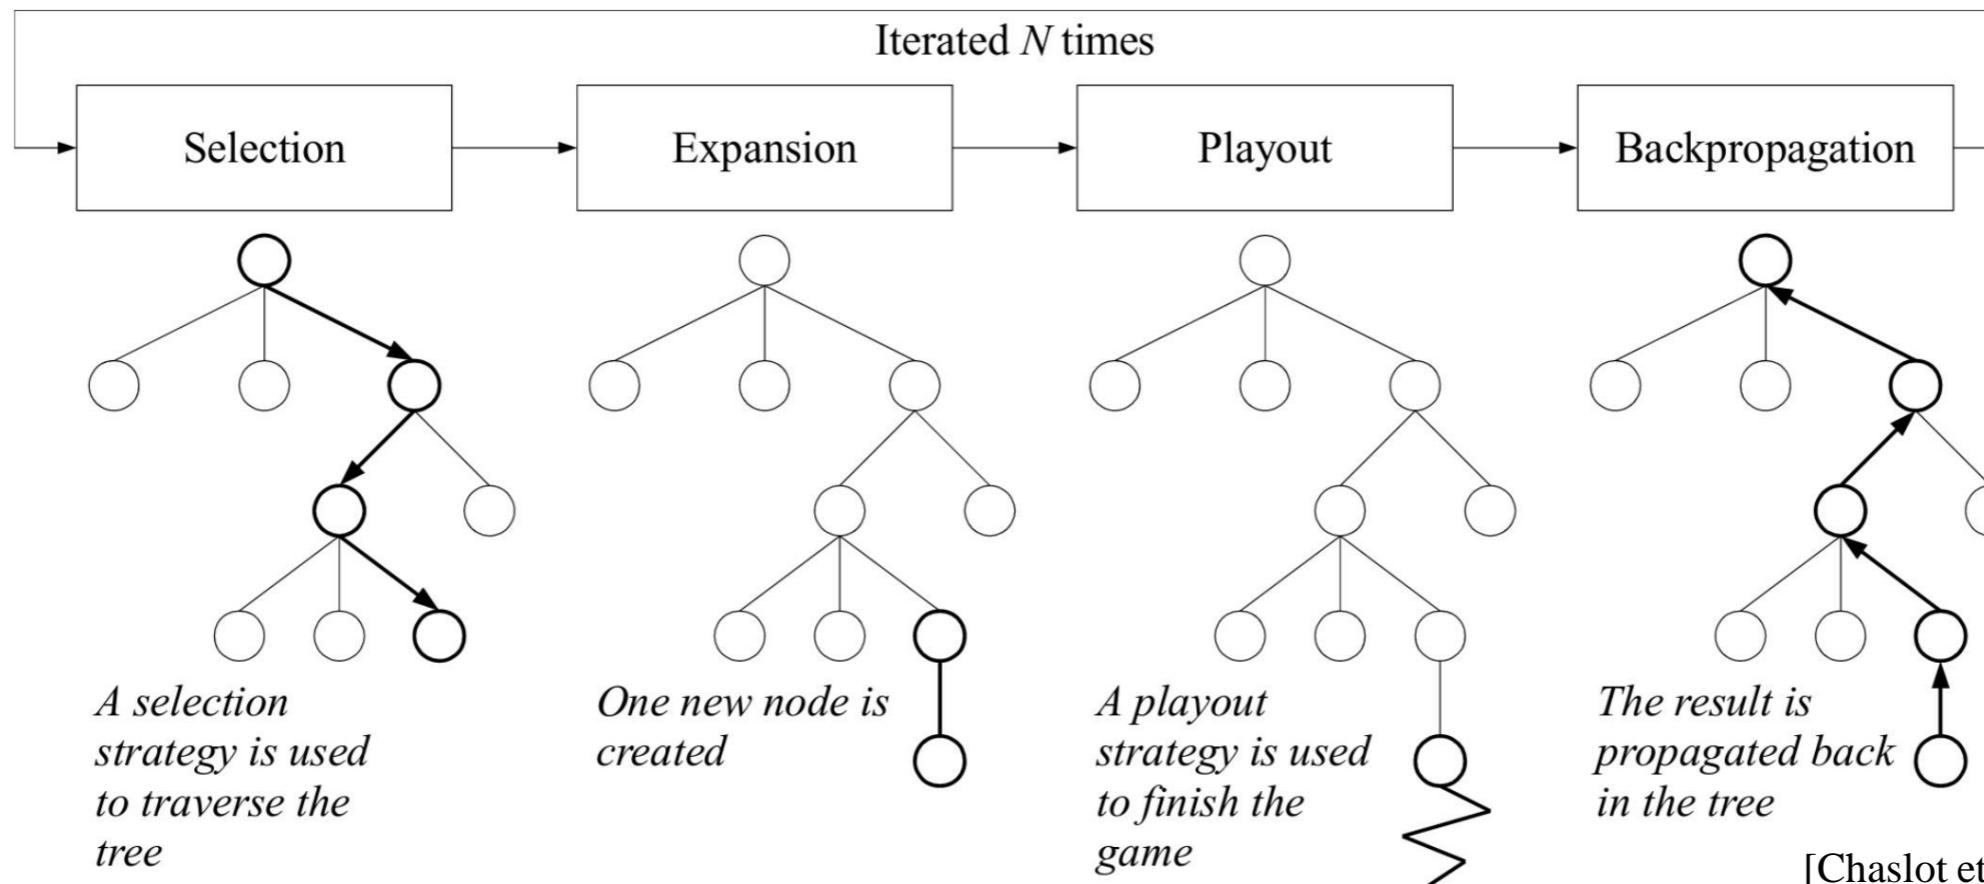
The diagram shows a grammar for pipeline search. Red annotations with arrows point to specific parts of the grammar rules:

- A red arrow labeled "production rule" points to the start symbol <Start>.
- A red arrow labeled "optional" points to the square brackets [] used for optional components like <Pre-processing>.
- A red arrow labeled "non-terminal" points to the non-terminal symbols like <Algorithm>, <Imputation>, etc.
- A red arrow labeled "terminal" points to the terminal symbols like SelectKBest, PCA, Euclidian, etc.

```
<Start> ::= [<Pre-processing>] <Algorithm>
<Pre-processing> ::= [<Imputation>] <DimensionalityDefinition>
<Imputation> ::= Mean | Median | Max
<DimensionalityDefinition> ::= <FeatureSelection> [<FeatureConstruction>]
                           [<FeatureSelection>] <FeatureConstruction>
<FeatureSelection> ::= <Supervised> | <Unsupervised>
<Supervised> ::= SelectKBest <K> <score> | VarianceThreshold | [...]
<score> ::= f-classification | chi2
<K> ::= 1 | 2 | 3 | [...] | NumberOfFeatures - 1
<perc> ::= 1 | 2 | 3 | [...] | 99
<Unsupervised> ::= PCA | FeatureAgglomeration <affinity> | [...]
<affinity> ::= Euclidian | L1 | L2 | Manhattan | Cosine
<FeatureConstruction> ::= PolynomialFeatures
<Algorithm> ::= <NaiveBayes> | <Trees> | [...]
<NaiveBayes> ::= GaussianNB | MultinomialNB | BernoulliNB
```

Pipeline search: Monte Carlo Tree Search

- Use MCTS to search for optimal pipelines
- Optimize the structure and hyperparameters simultaneously by building a surrogate model to predict configuration performance
 - Bayesian surrogate model: MOSAIC [\[Rakotoarison et al. 2019\]](#)
 - Neural network: AlphaD3M [\[Drori et al. 2018\]](#)



AlphaD3M Example



```
In [2]: from d3m_interface import AutoML
```

Generating pipelines using AlphaD3M

```
In [4]: train_dataset_path = '/home/ubuntu/datasets/seed_datasets_current/JIDO_SOHR_Articles_1061/TRAIN'
test_dataset_path = '/home/ubuntu/datasets/seed_datasets_current/JIDO_SOHR_Articles_1061/TEST'
score_dataset_path = '/home/ubuntu/datasets/seed_datasets_current/JIDO_SOHR_Articles_1061/SCORE'
output_path = '/home/ubuntu/output/'
tmp_path = '/home/ubuntu/tmp'
```

```
In [5]: automl = AutoML(output_path, 'AlphaD3M')
pipelines = automl.search_pipelines(train_dataset_path, time_bound=5)
```

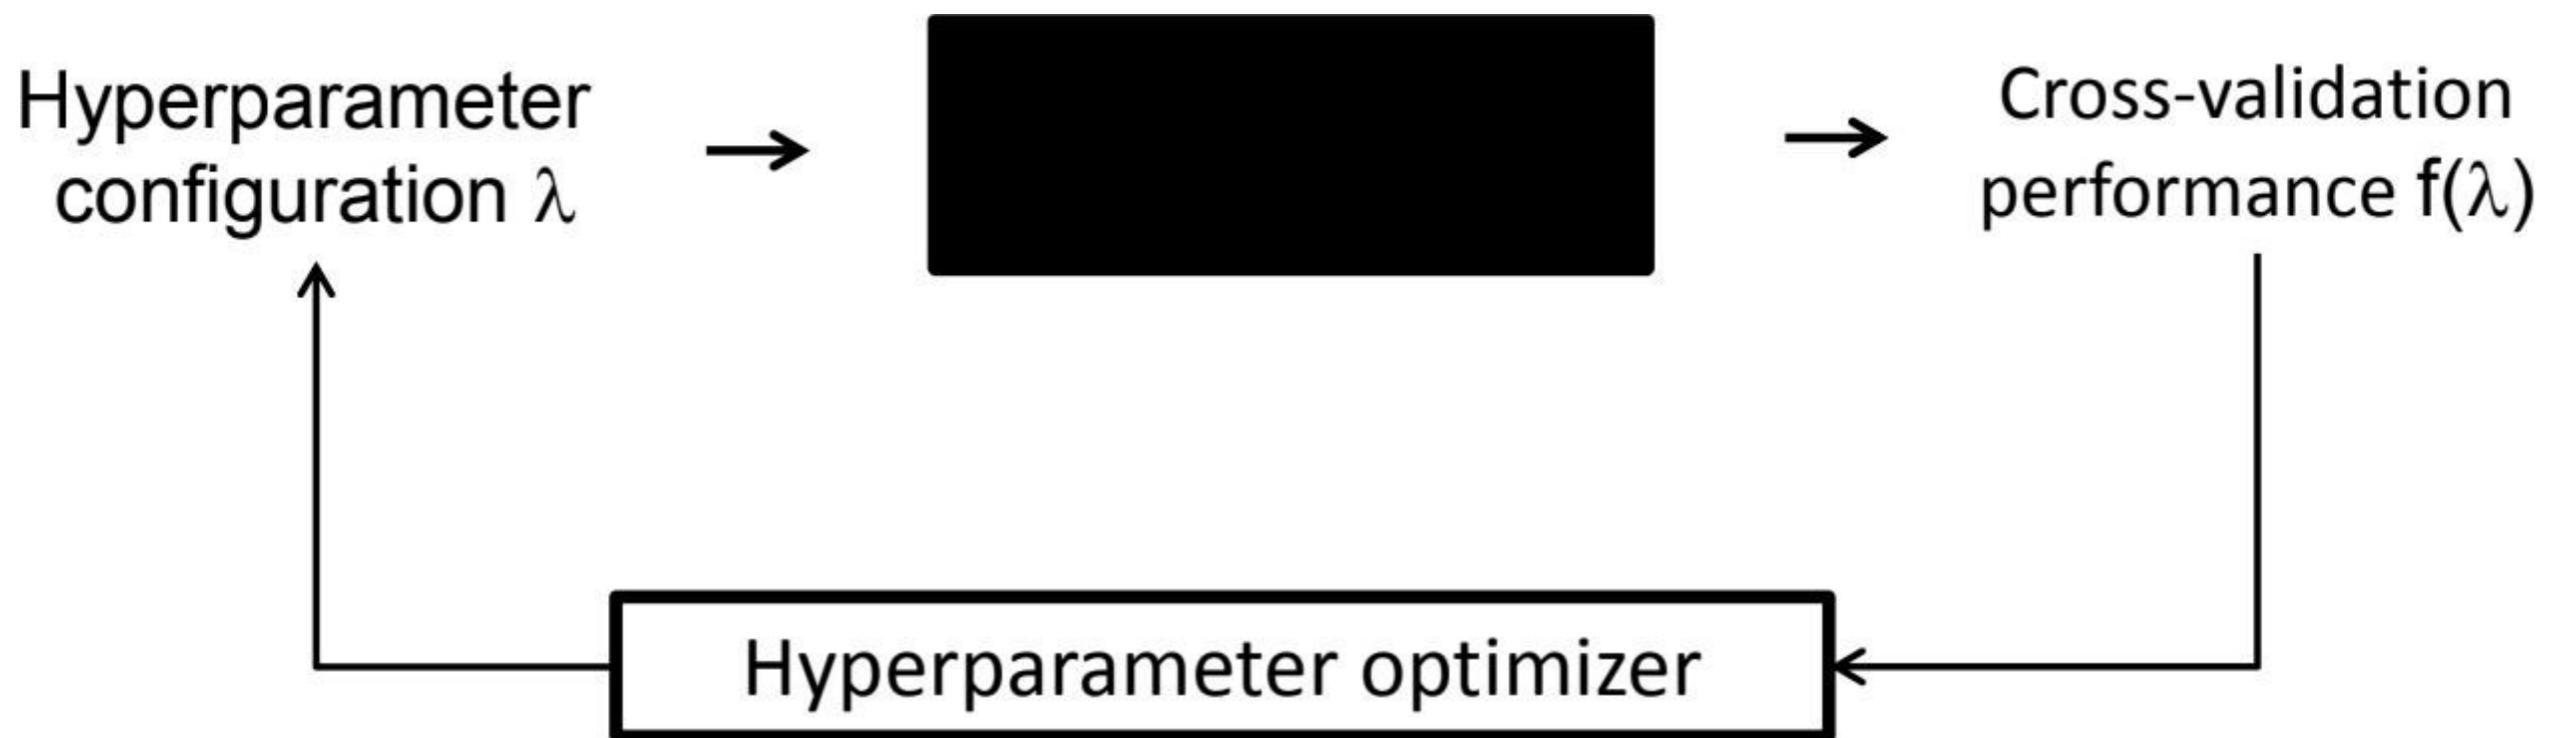
```
automl_red.leaderboard
```

ranking		id	summary	f1
0	1	c53f5b6e-872f-4ae0-9366-6b7cff85b5de	text_reader.common, add_semantic_types.common,...	0.83410
1	2	75f86ebc-37f1-446a-9baa-24ef474fd215	add_semantic_types.common, imputer.sklearn, en...	0.05021
2	3	bbfd3b0-6086-4a8a-8e9f-708ed586edef	add_semantic_types.common, imputer.sklearn, en...	0.04255
3	4	bb02b1a1-93fa-4df4-98d5-ee5574c65917	add_semantic_types.common, imputer.sklearn, en...	0.02542
4	5	343507f5-a408-4b8e-a1cc-91cde7135623	add_semantic_types.common, imputer.sklearn, en...	0.00000
5	6	f37da102-0663-4e9d-86b0-9e1de8488434	add_semantic_types.common, imputer.sklearn, en...	0.00000

Hyperparameter optimization (HPO)

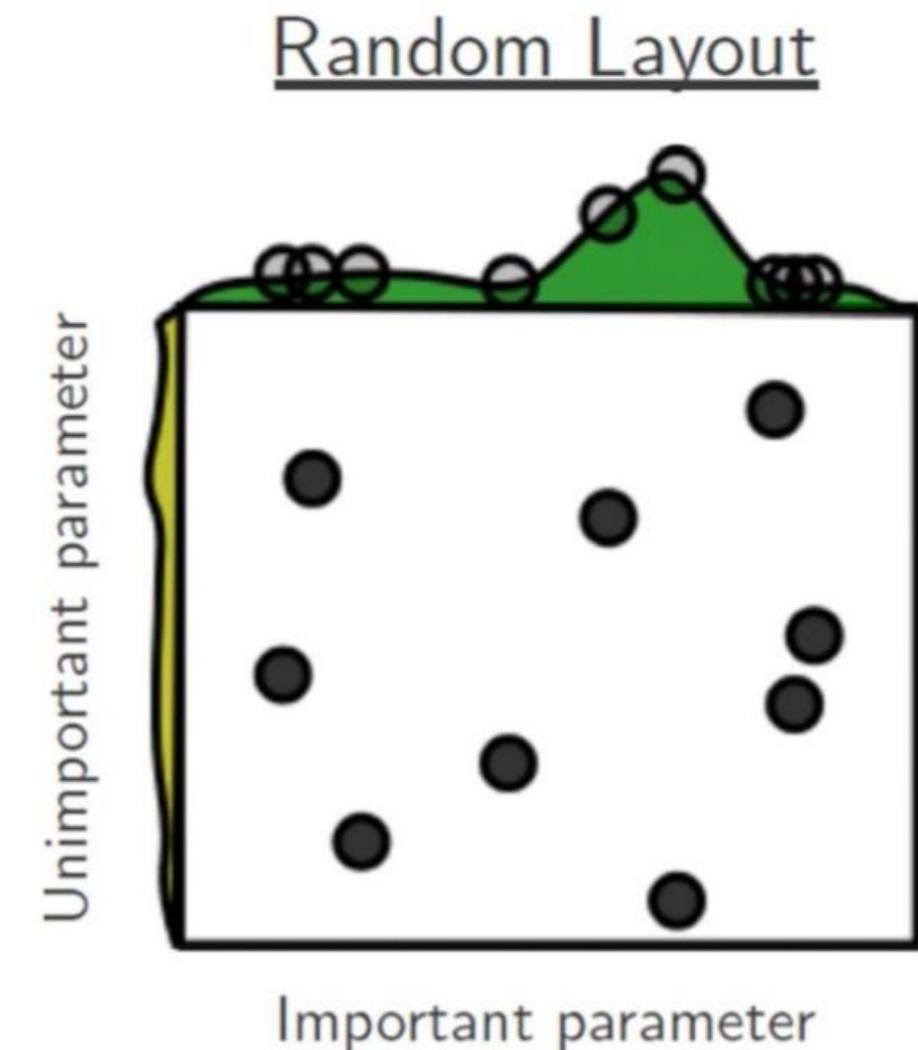
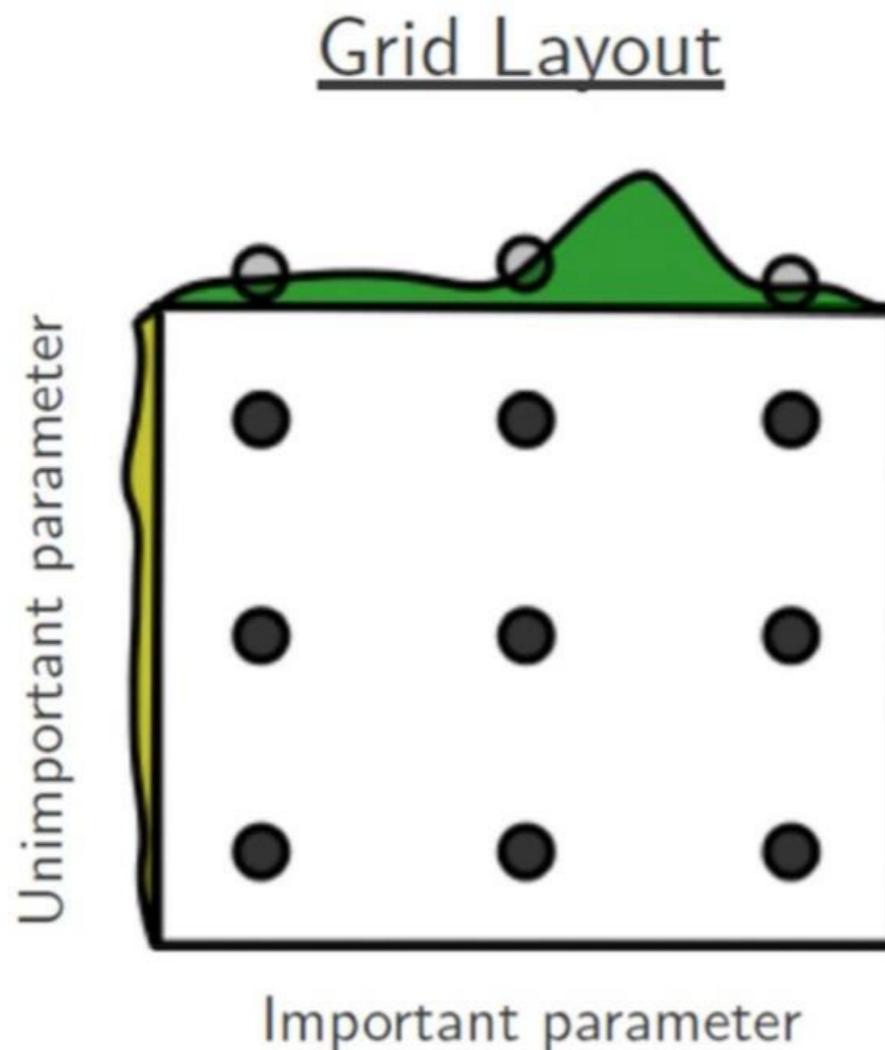
Black box optimization: expensive for machine learning.

Sample efficiency is important!



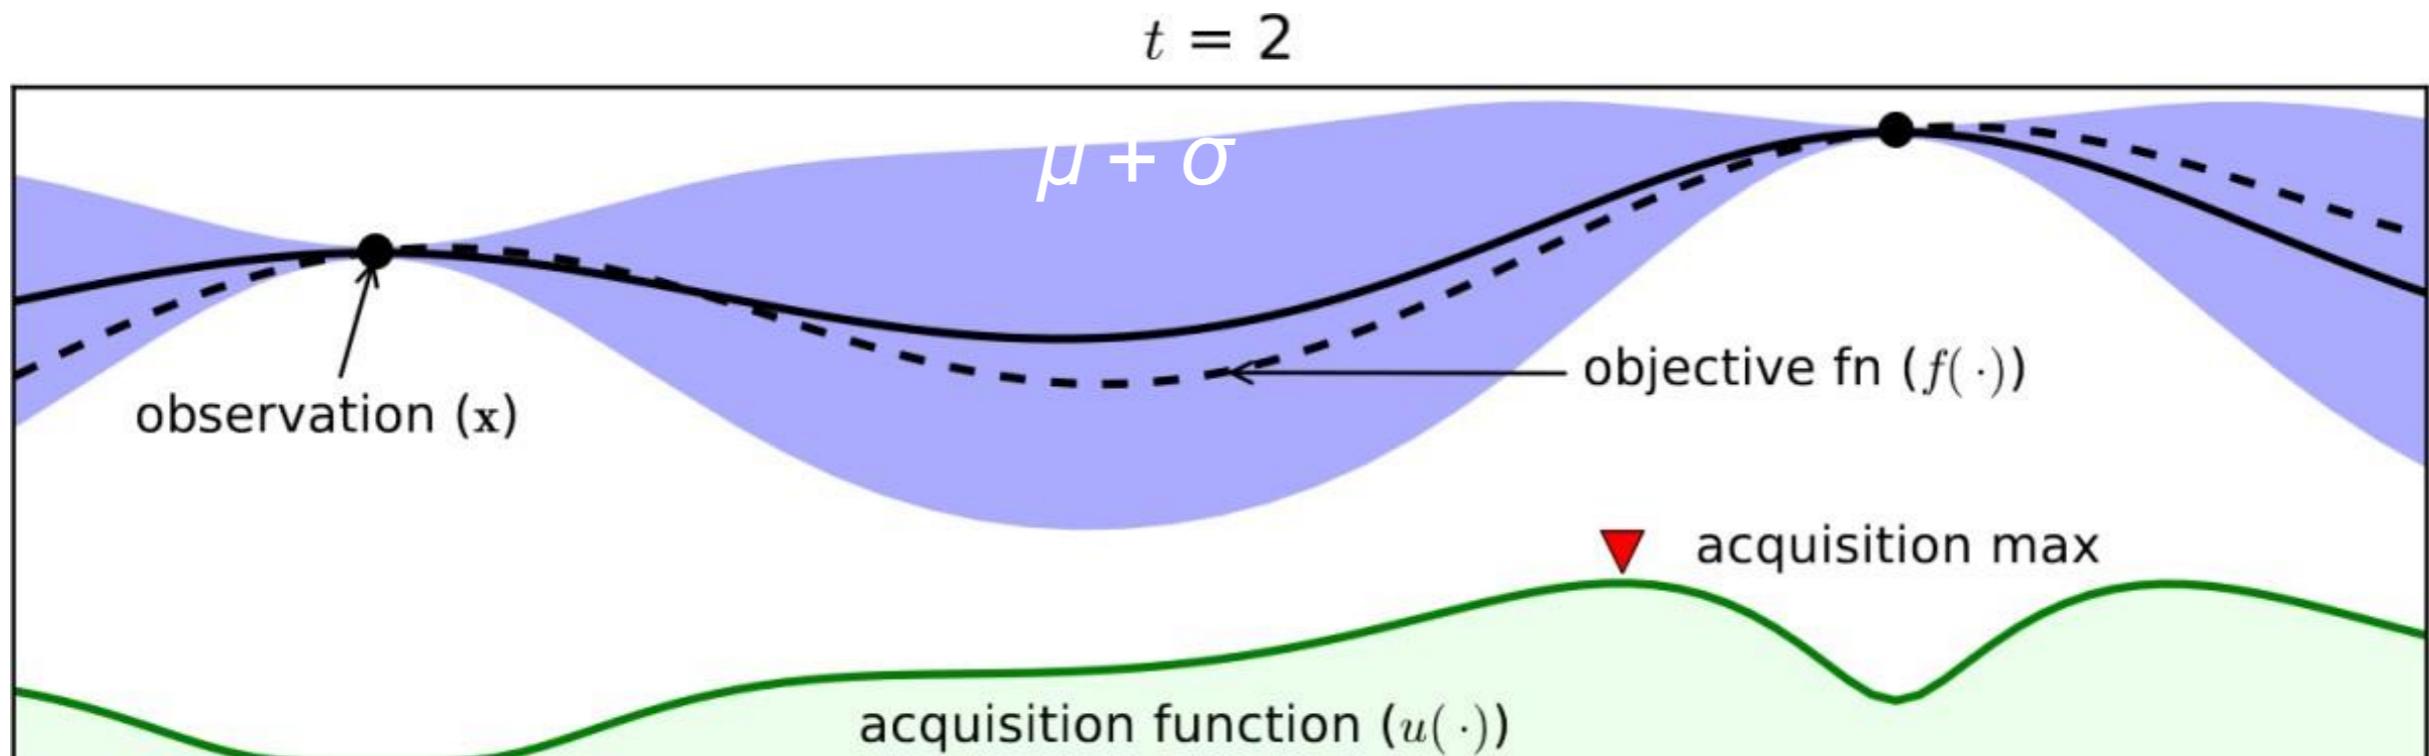
HPO: grid and random search

- Random search handles unimportant dimensions better
- Easily parallelizable, but uninformed (no learning)



HPO: Bayesian Optimization

- Start with a few (random) hyperparameter configurations
- Build a **surrogate model** to predict how well other configurations will work: mean and standard deviation (blue band)
 - Any probabilistic regression model: e.g. Gaussian processes
- To avoid a greedy search, use an **acquisition function** to trade off exploration and exploitation, e.g. Expected Improvement (EI)
- Sample for the best configuration under that function

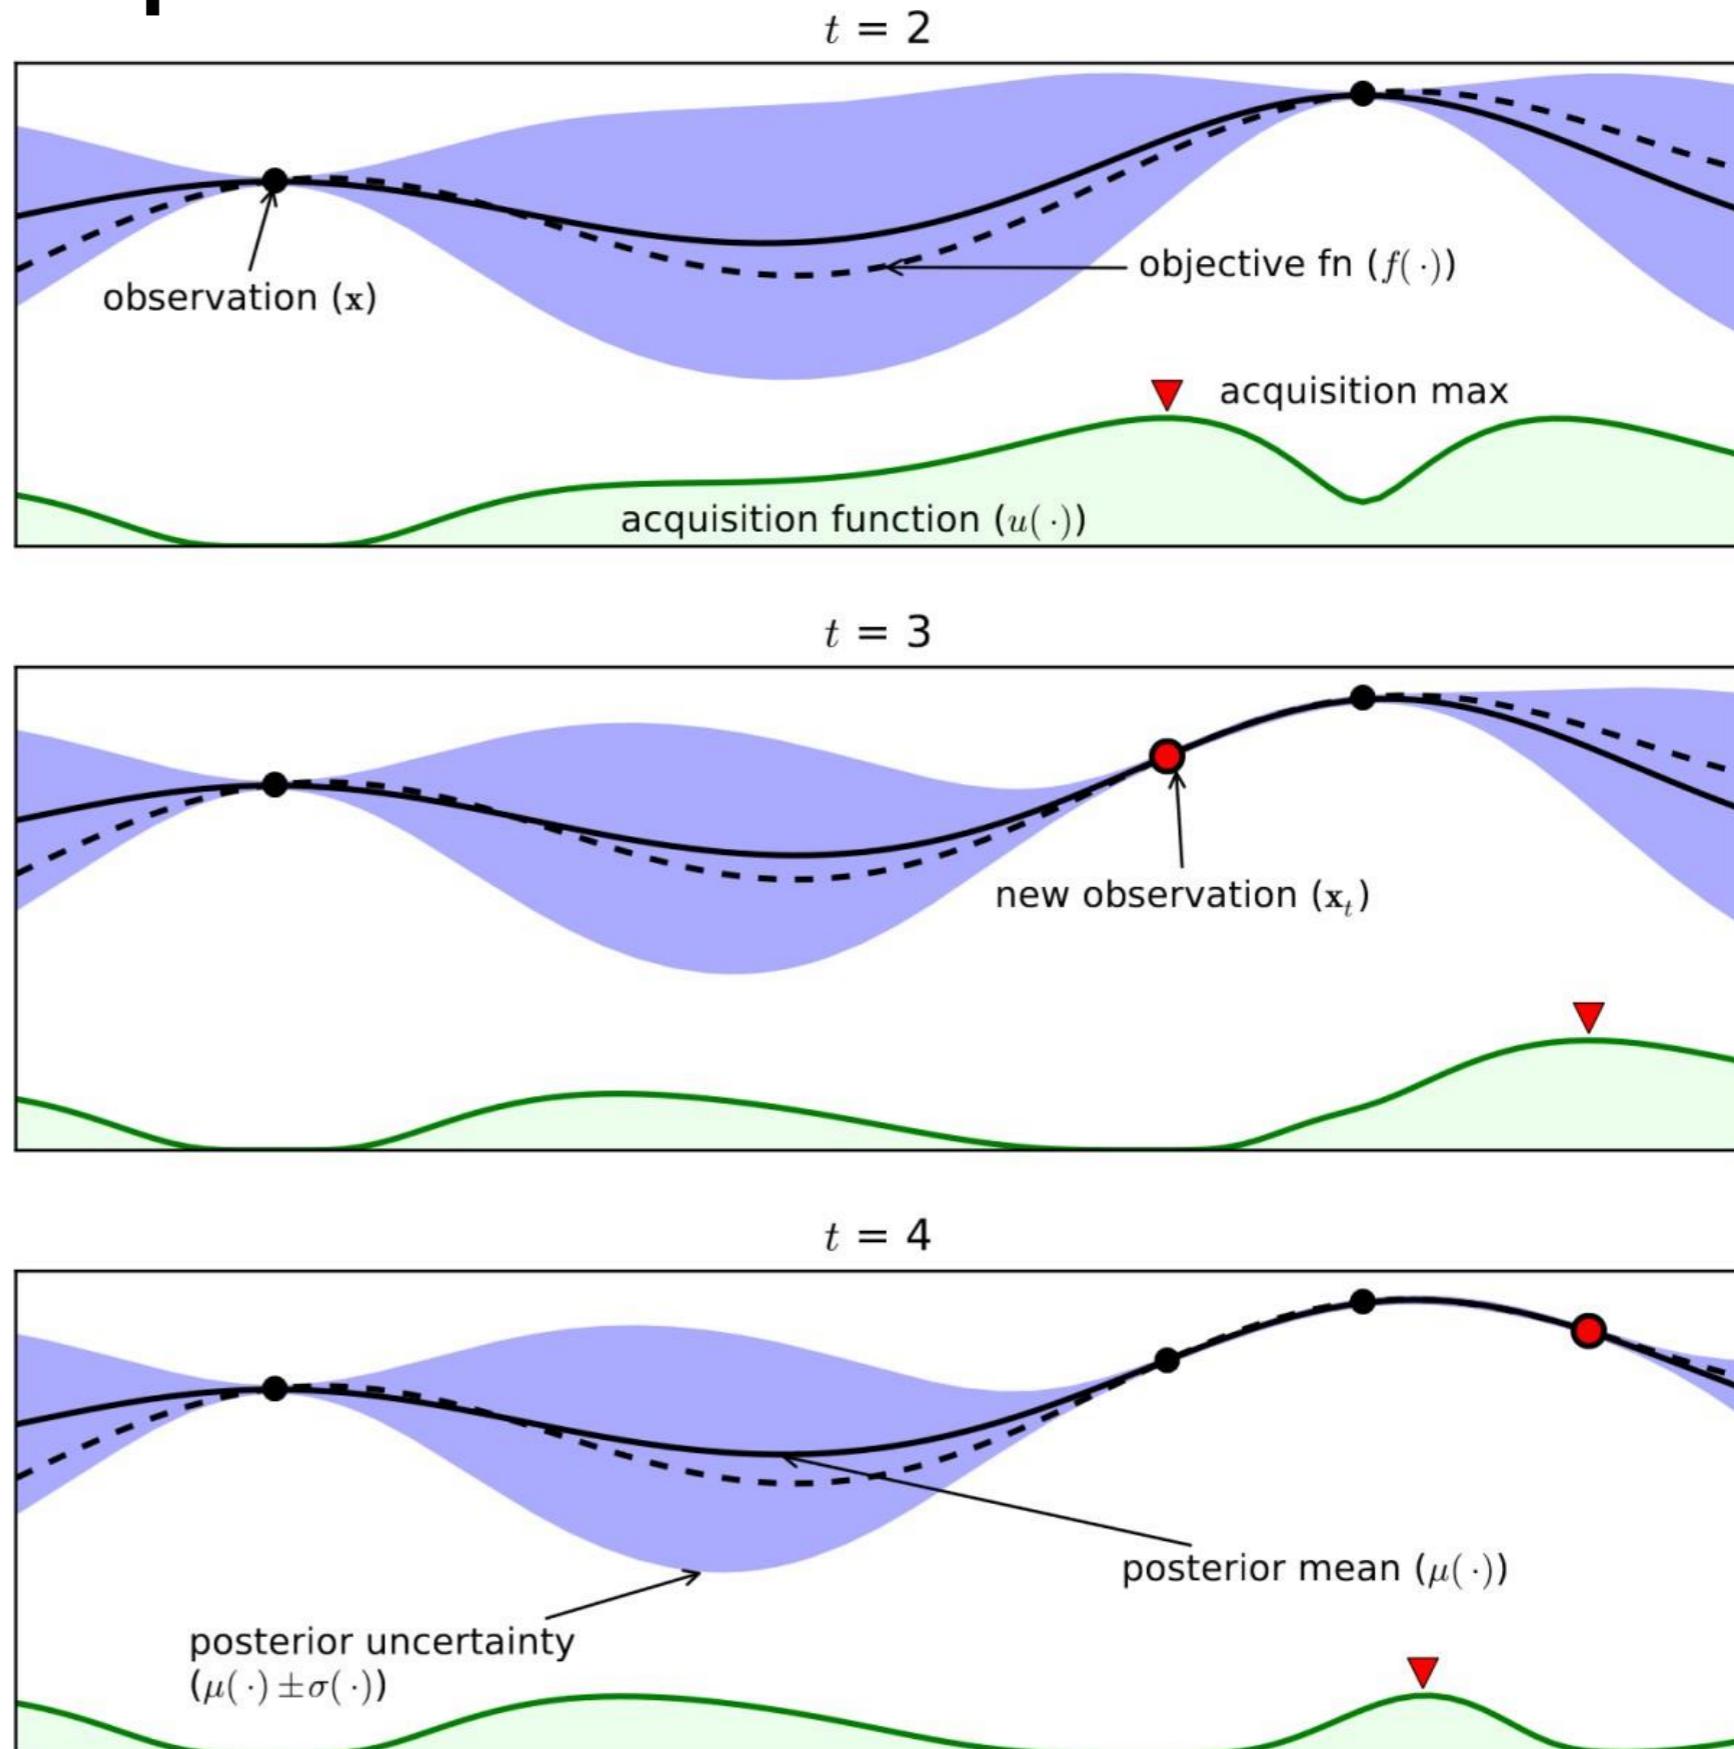


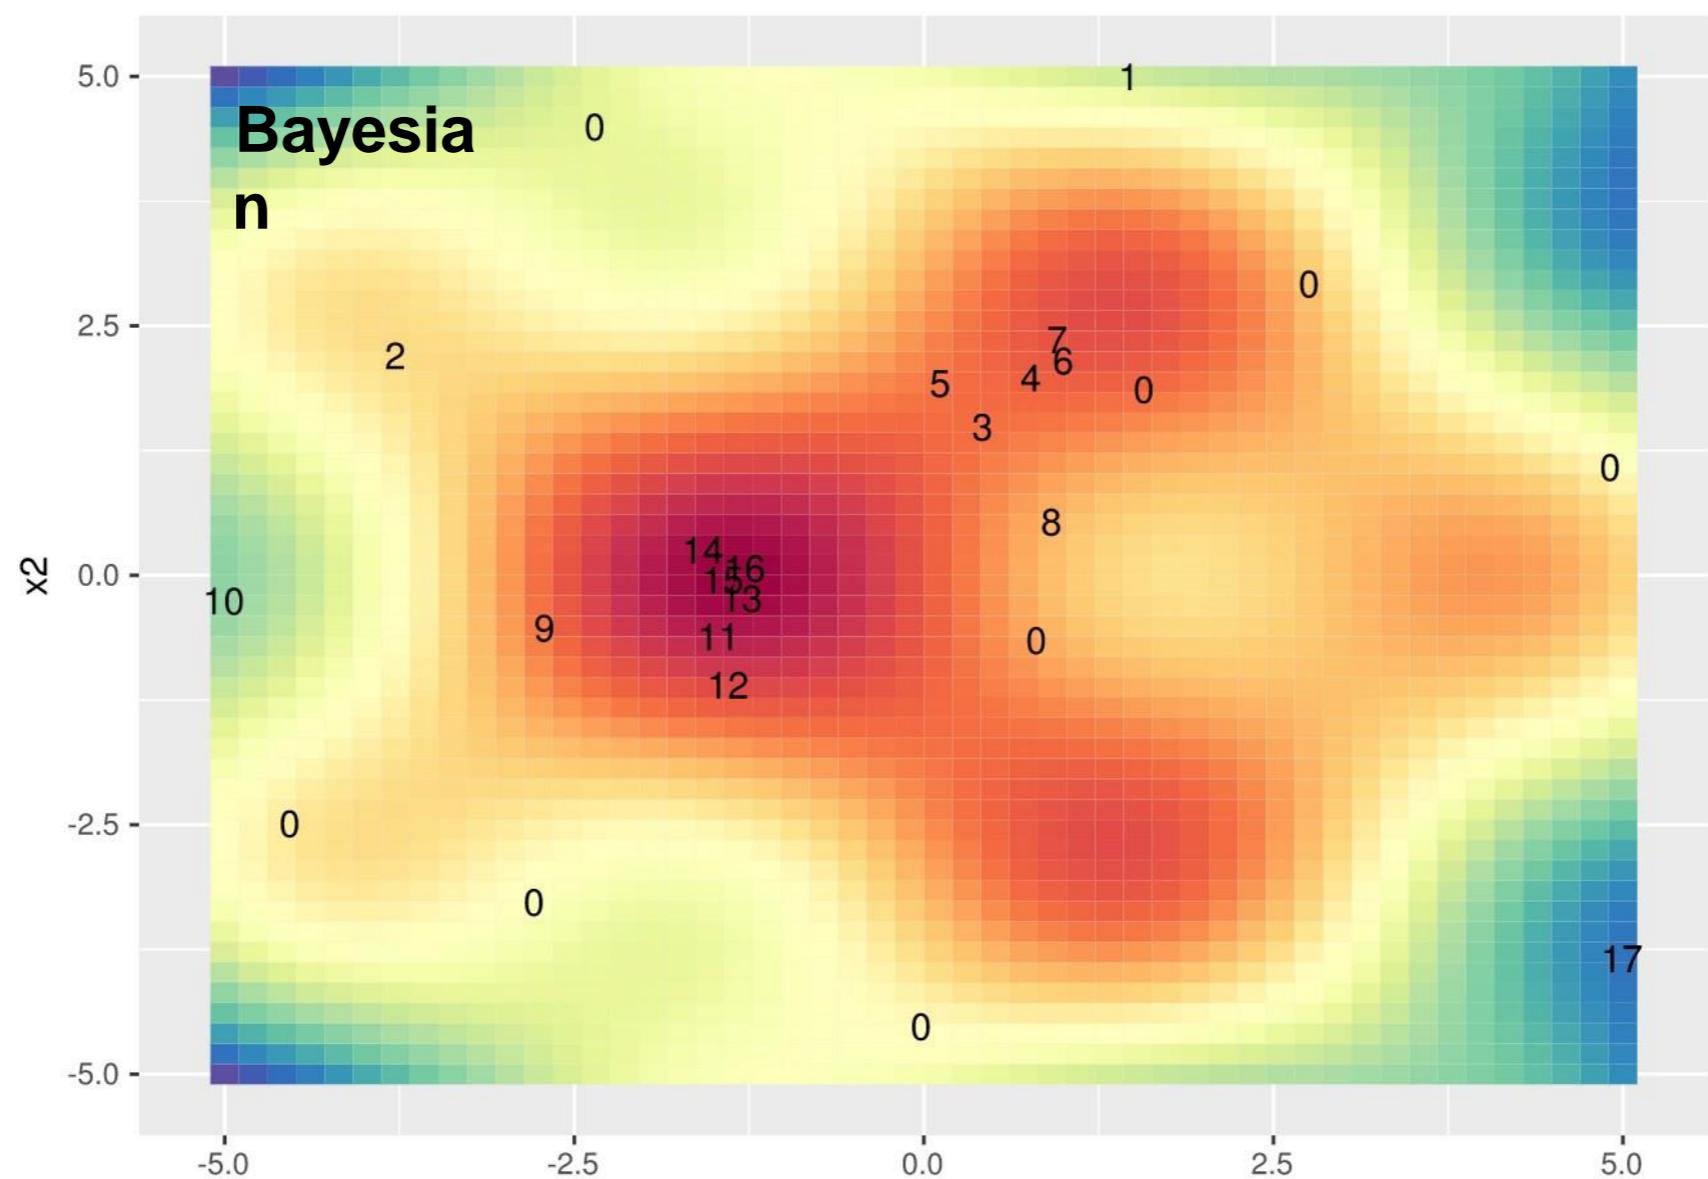
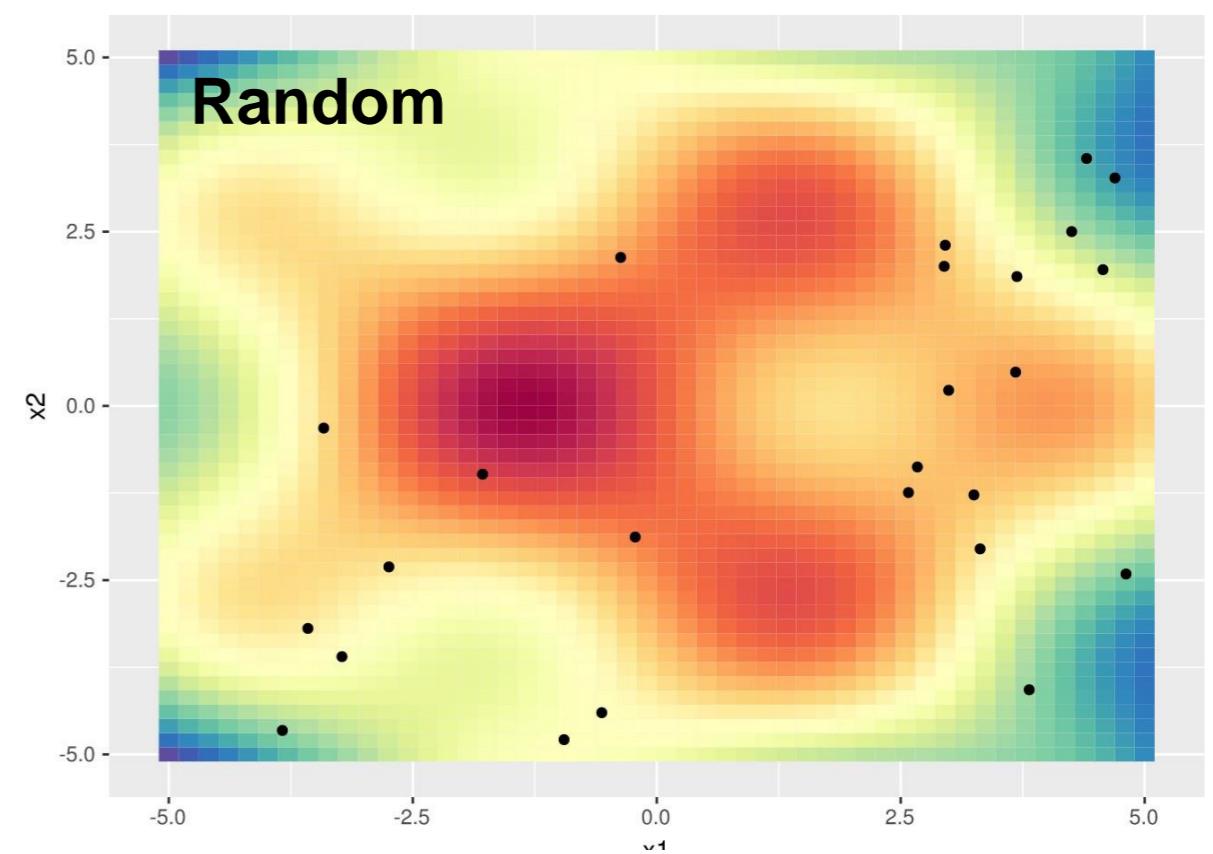
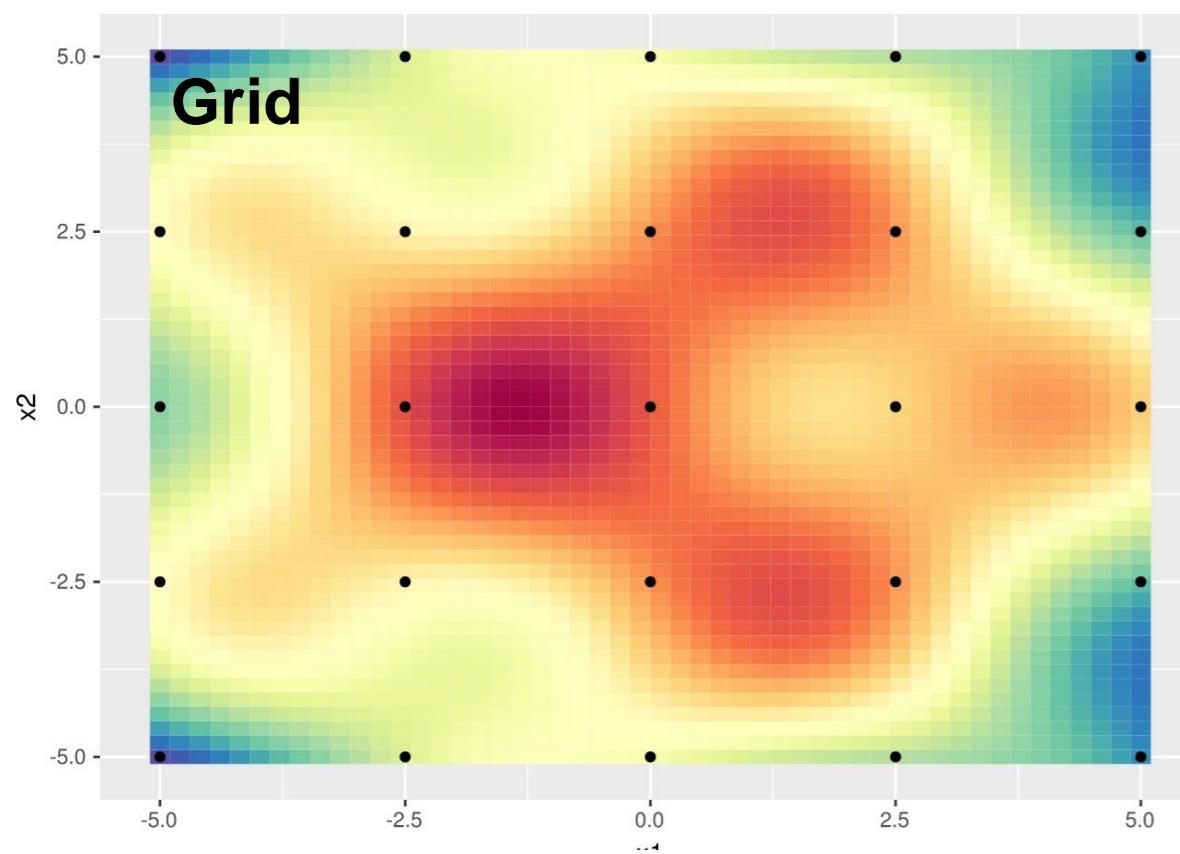
HPO: Bayesian Optimization

- Repeat
- Stopping criterion:
 - Fixed budget (time, evaluations)
 - Min. distance between configs
 - Threshold for acquisition function
 - Still overfits easily
- Convergence results

[Srinivas et al. 2010](#), [Freitas et al. 2012](#), [Kawaguchi et al. 2016](#)

- Good for non-convex, noisy objectives
- Used in AlphaGo





Meta-learning

Meta-learning can drastically speed up the architecture and hyperparameter search, in combination with the optimization techniques we just discussed

- Learn which hyperparameters are really important
- Learn which hyperparameter values should be tried first
- Learn which architectures will most likely work
- Learn how to clean and annotate data
- Learn which feature representations to use
- ...

Overview

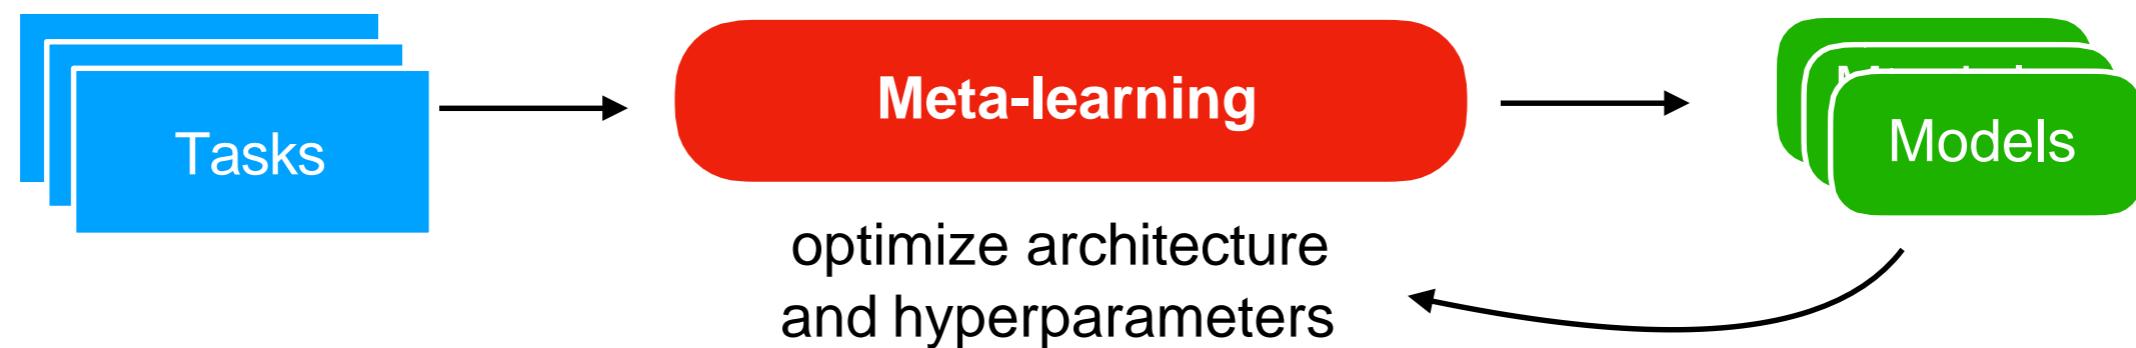
part AutoML introduction, promising optimization techniques

1



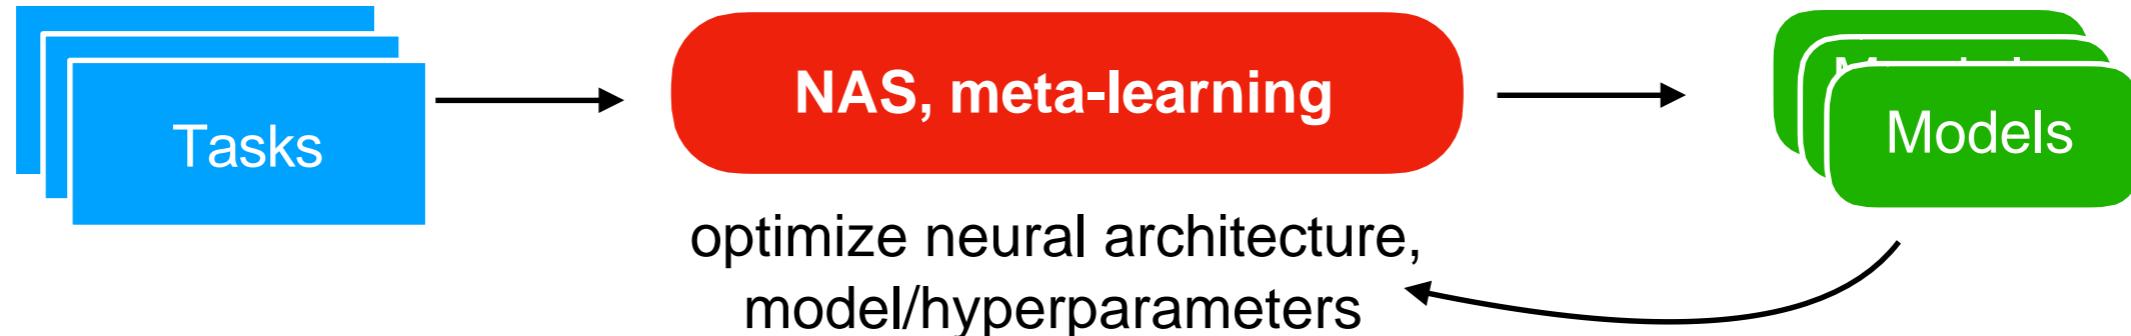
part Meta-learning

2



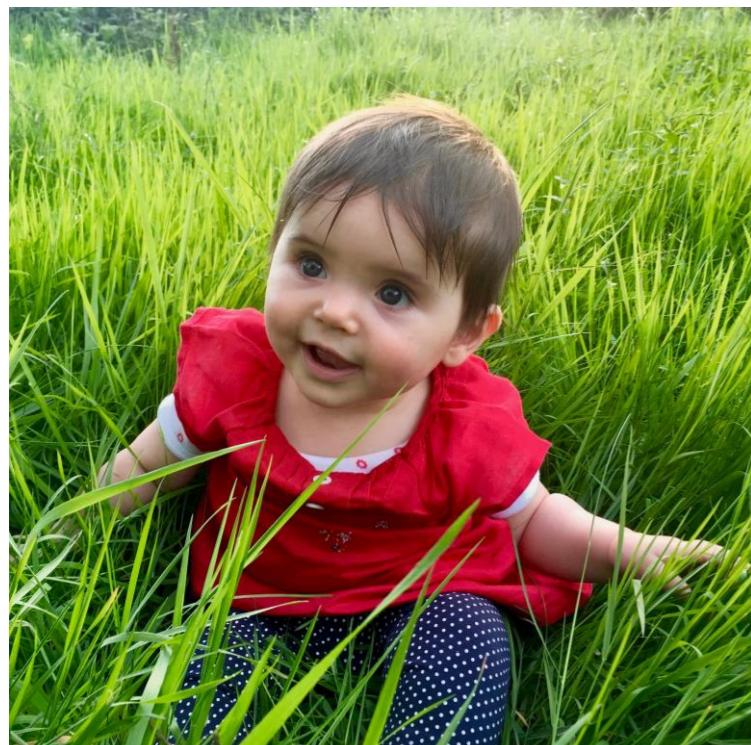
part Neural architecture search, meta-learning on neural nets

3



Learning is a never-ending process

Humans don't learn from scratch



Learning is a never-ending process

Learning humans also seek/create related tasks

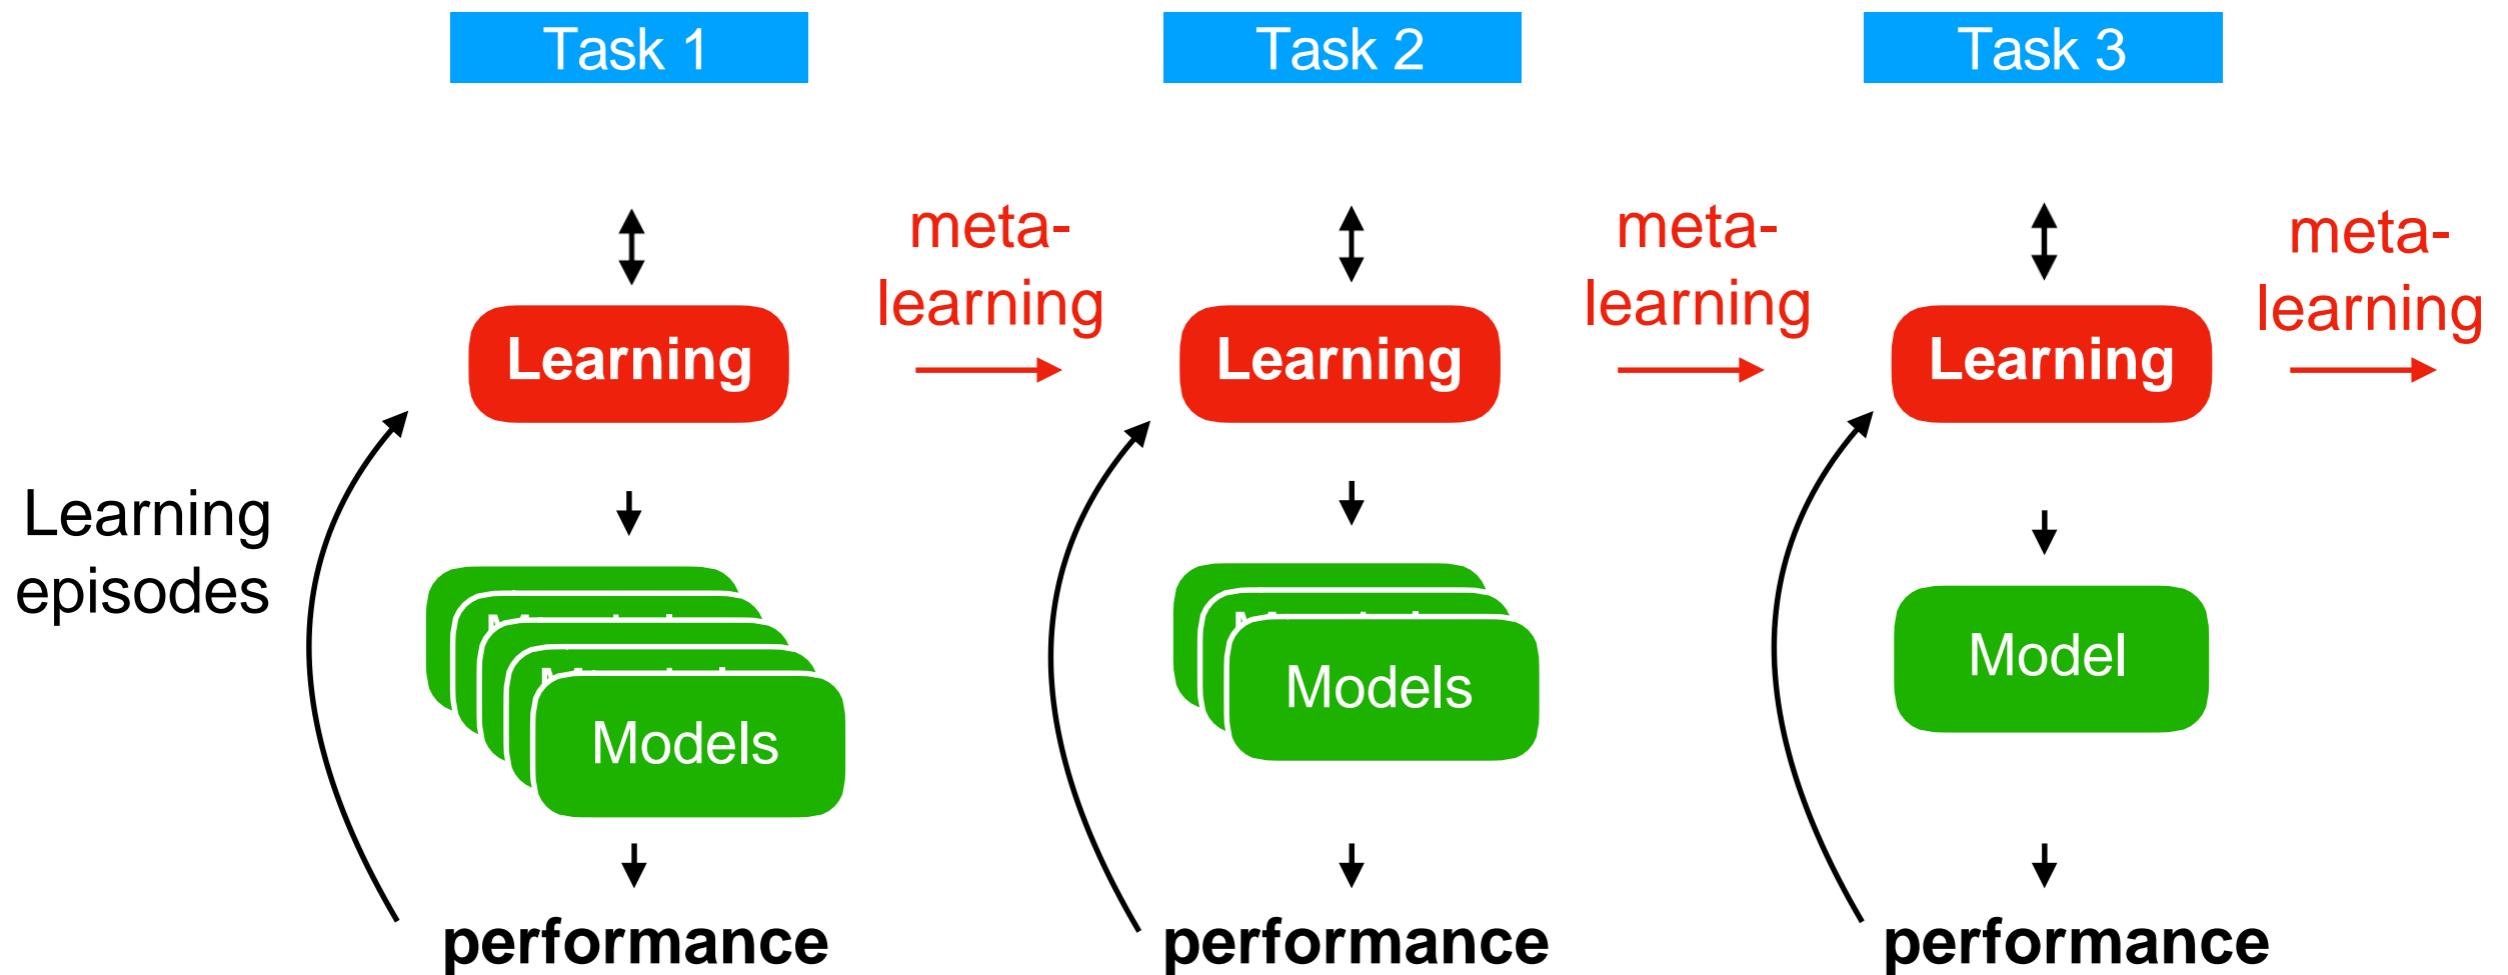


E.g. find many similar puzzles, solve them in different ways,...

Learning is a never-ending process

Humans learn across tasks

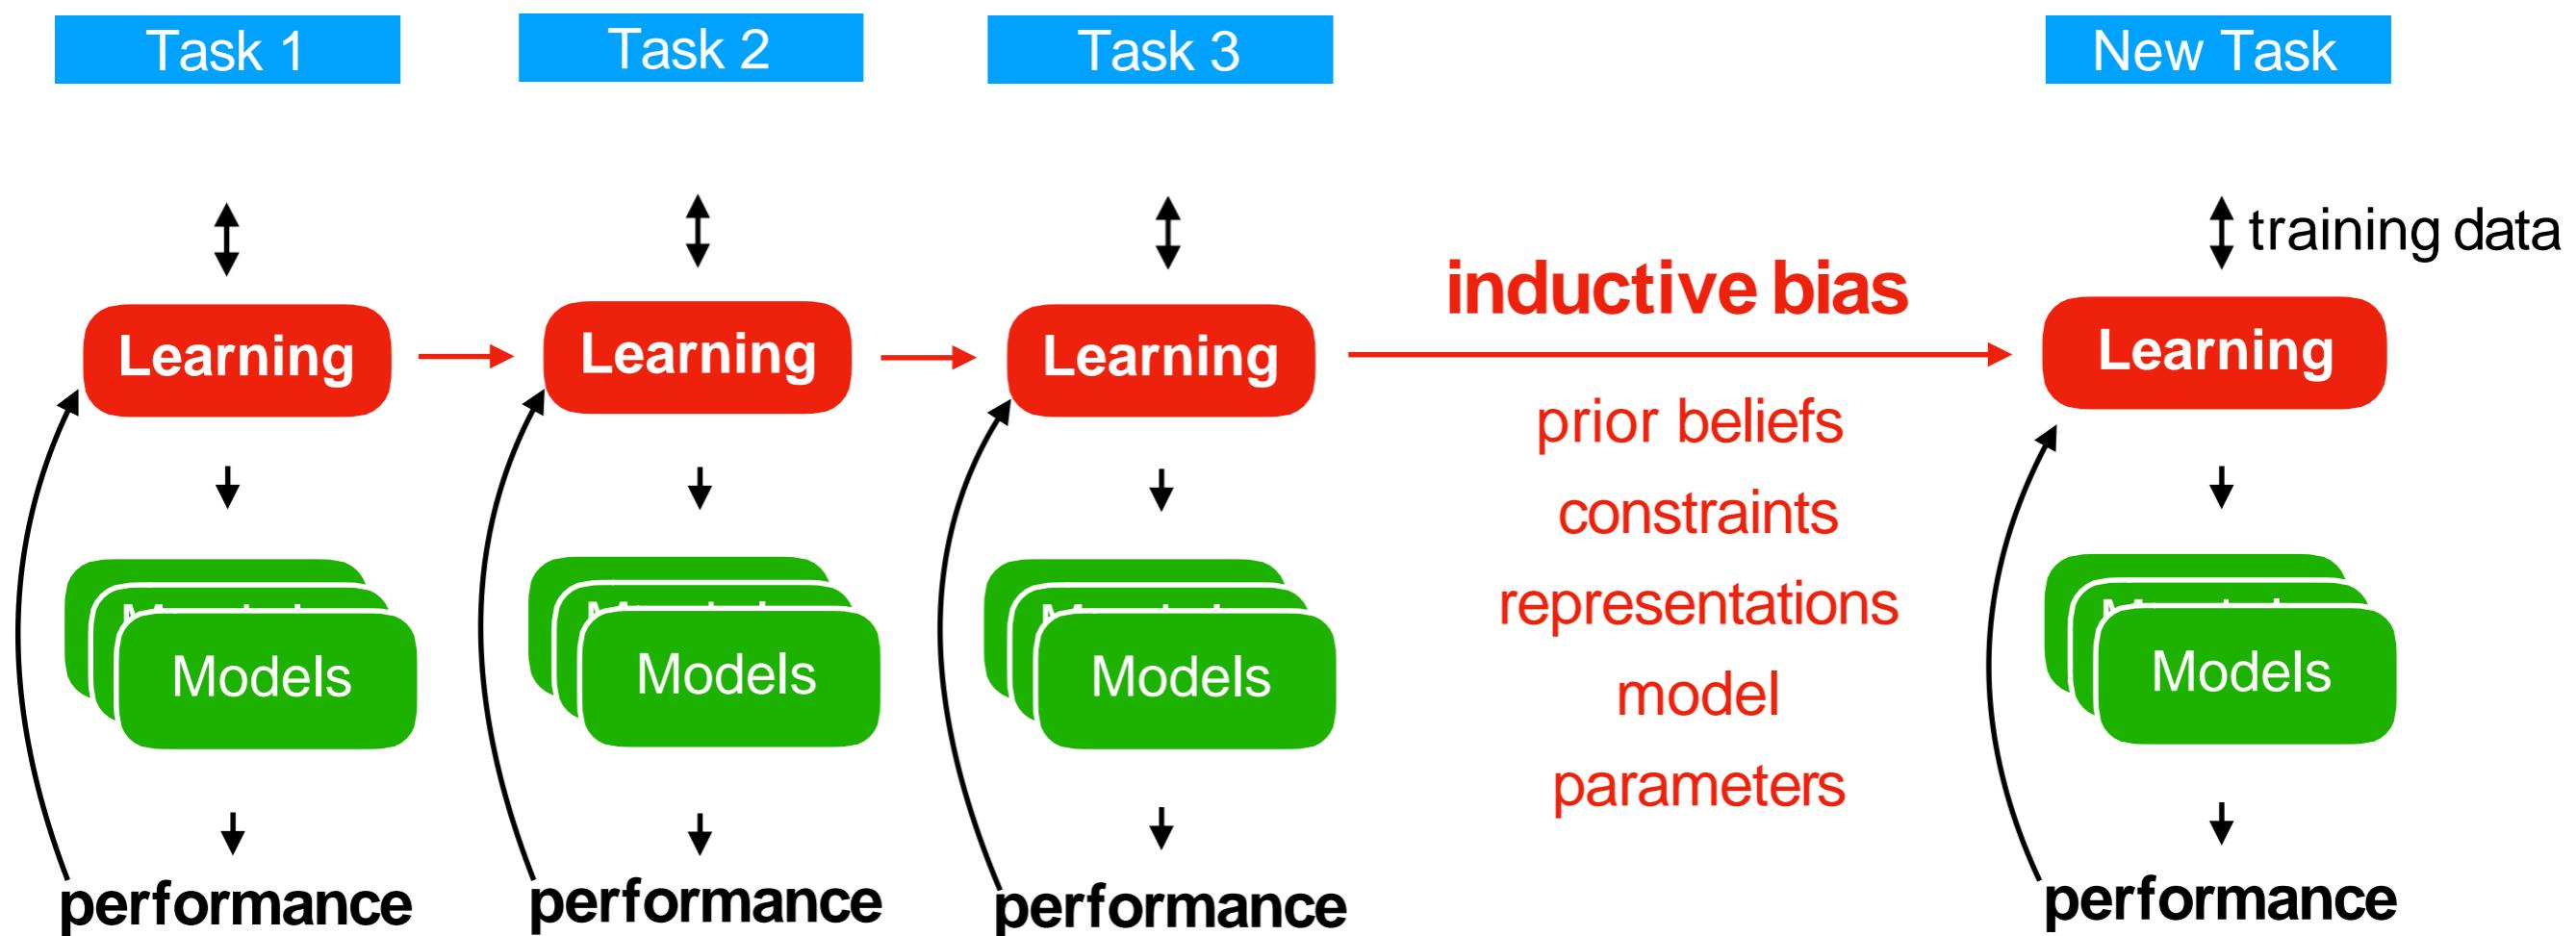
Why? Requires less trial-and-error, less data



Learning to learn

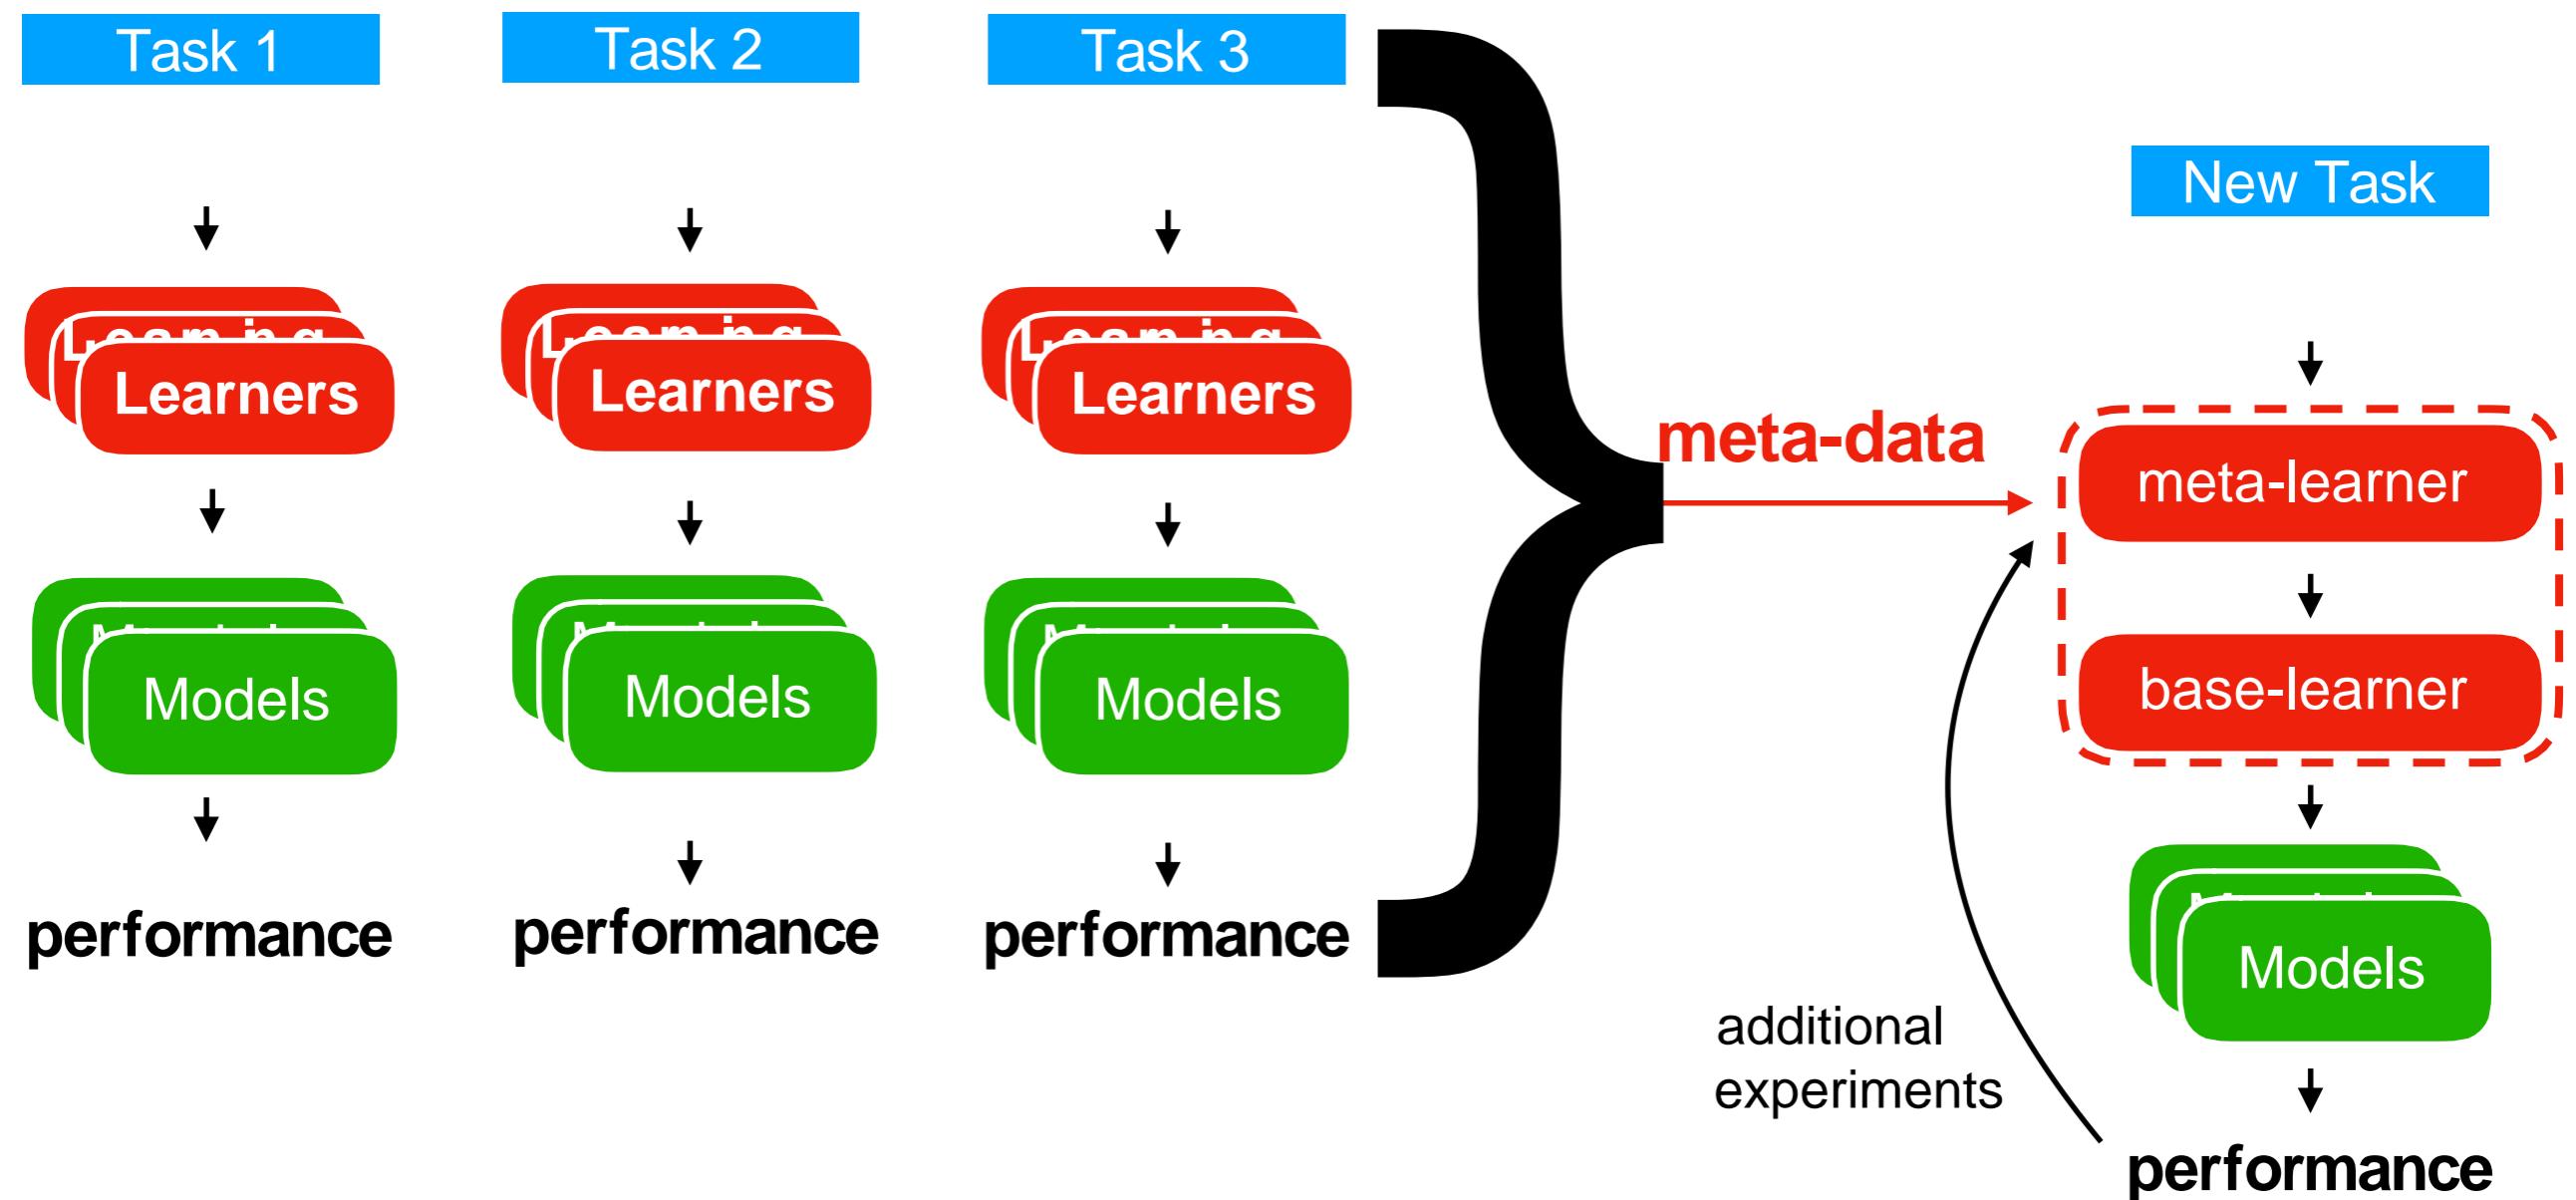
If prior tasks are *similar*, we can **transfer** prior knowledge to new tasks

Inductive bias: assumptions added to the training data to learn effectively



Meta-learning

Meta-learner *learns* a (base-)learning algorithm, based on *meta-data*



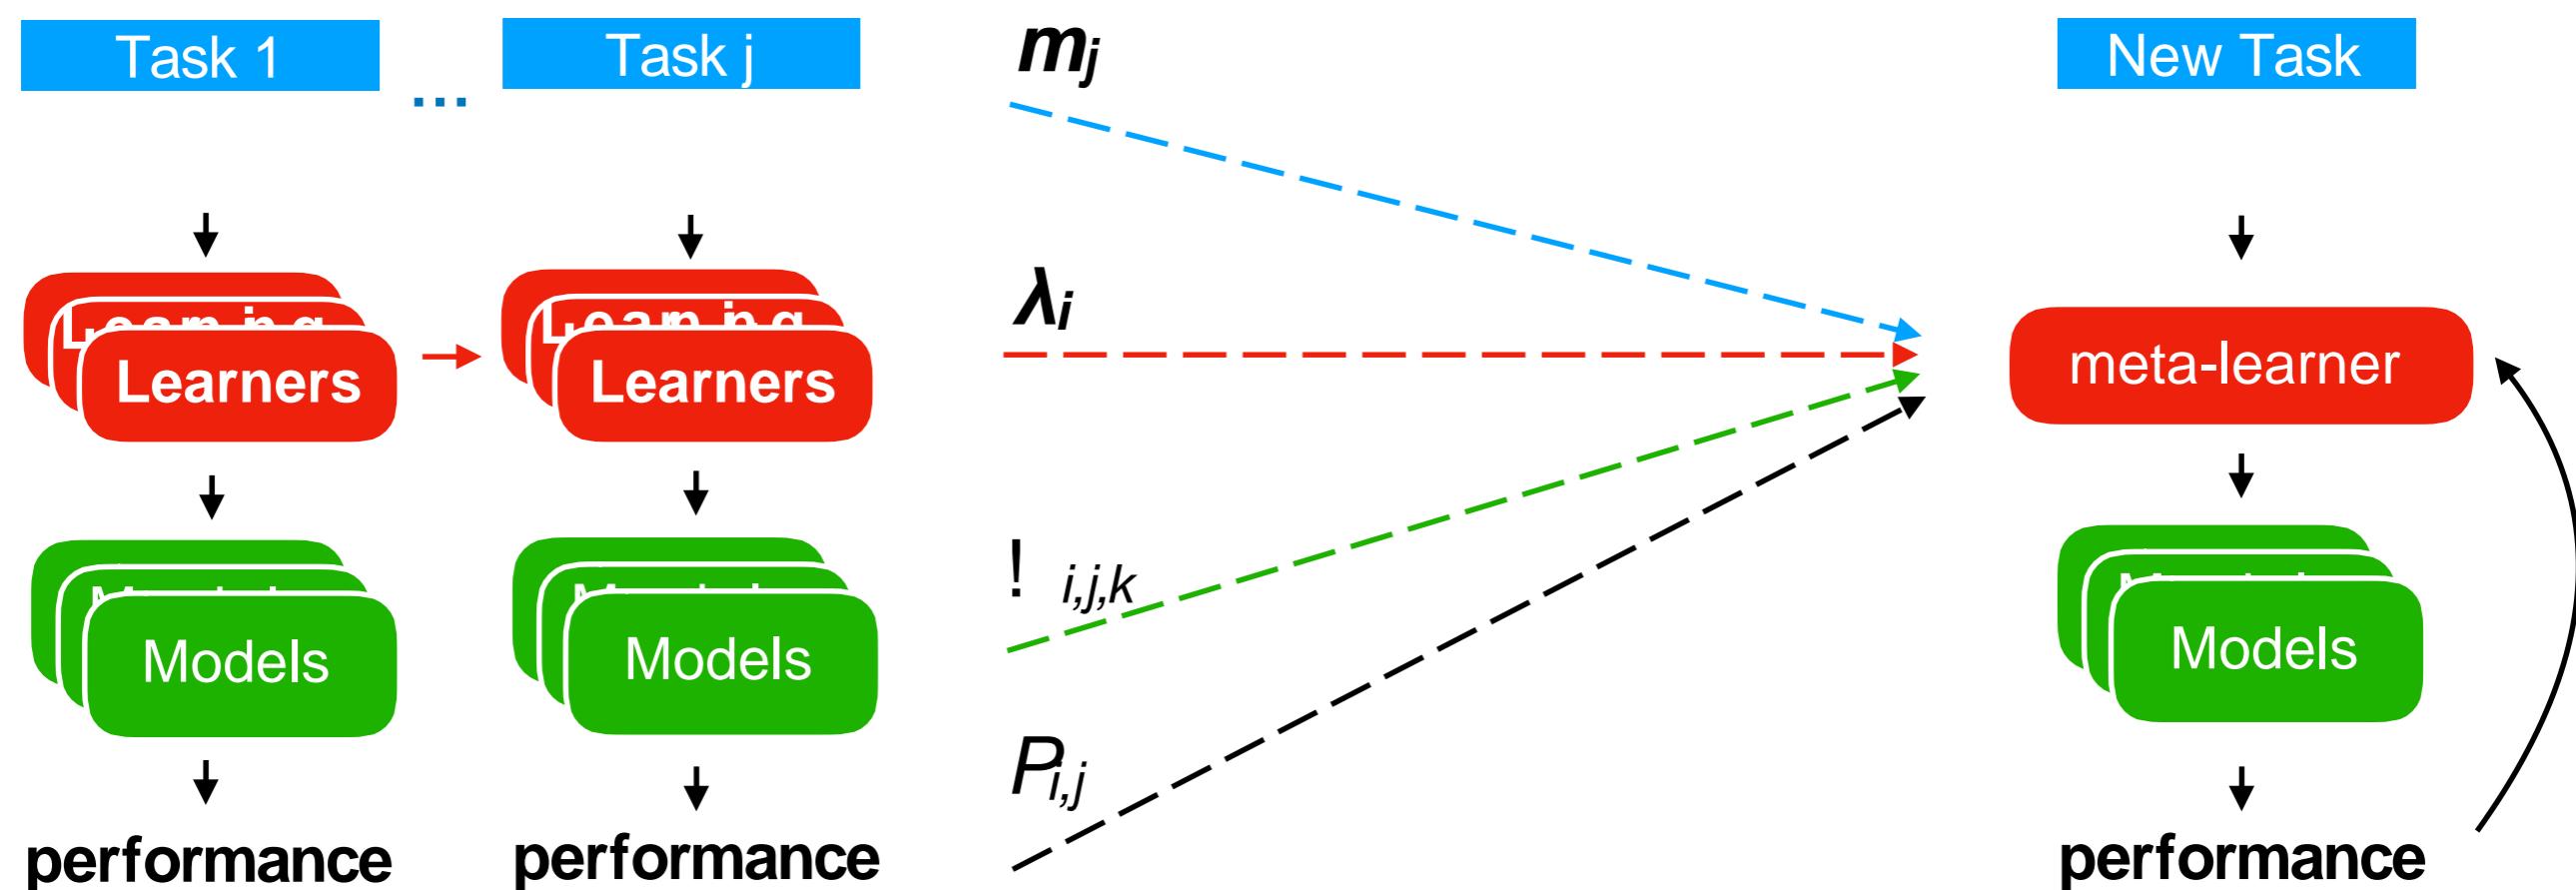
Meta-data

m_j Description of task j (meta-features)

λ_i Configuration i (architecture + hyperparameters)

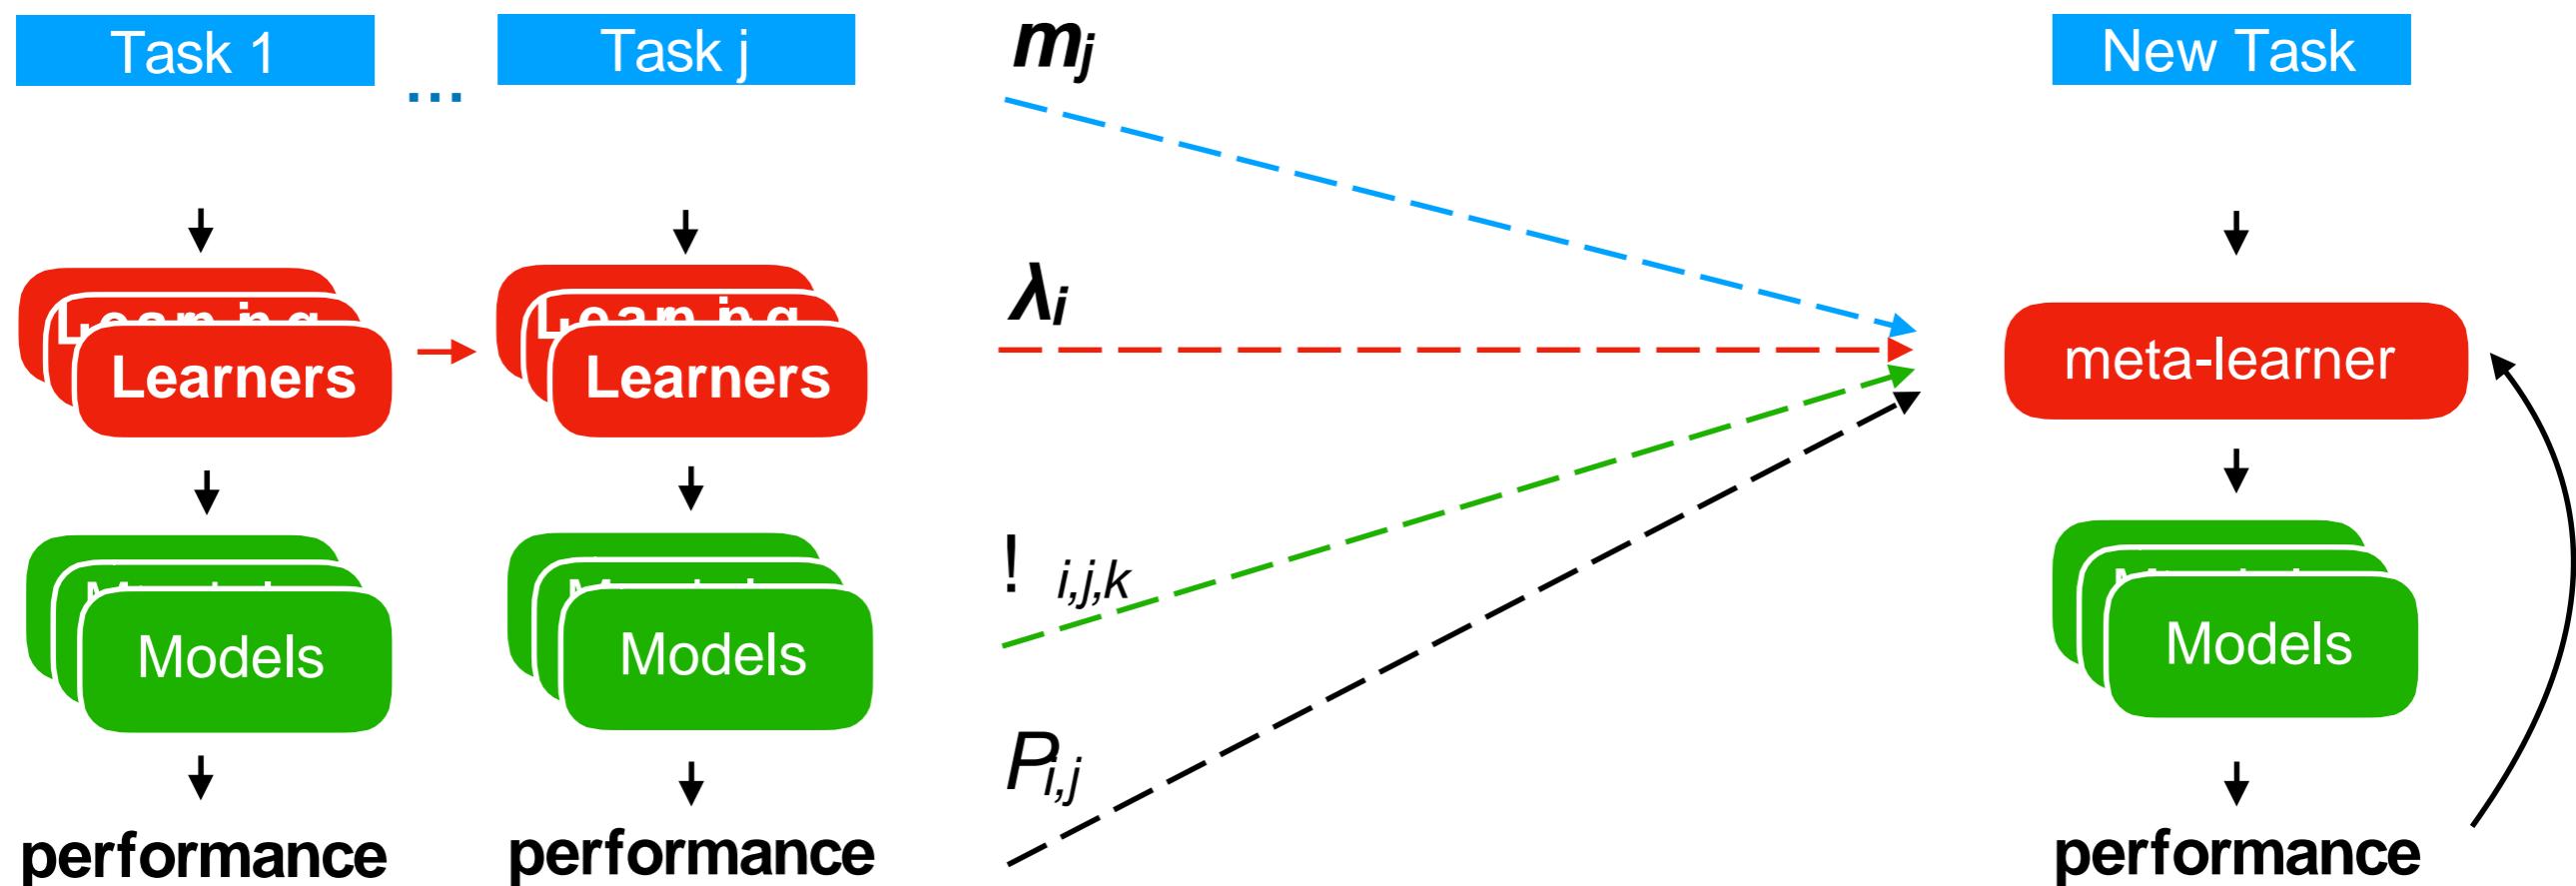
$!_{i,j,k}$ Model parameters (e.g. weights)

P_{ij} Performance estimate of model built with λ_i on task j



How?

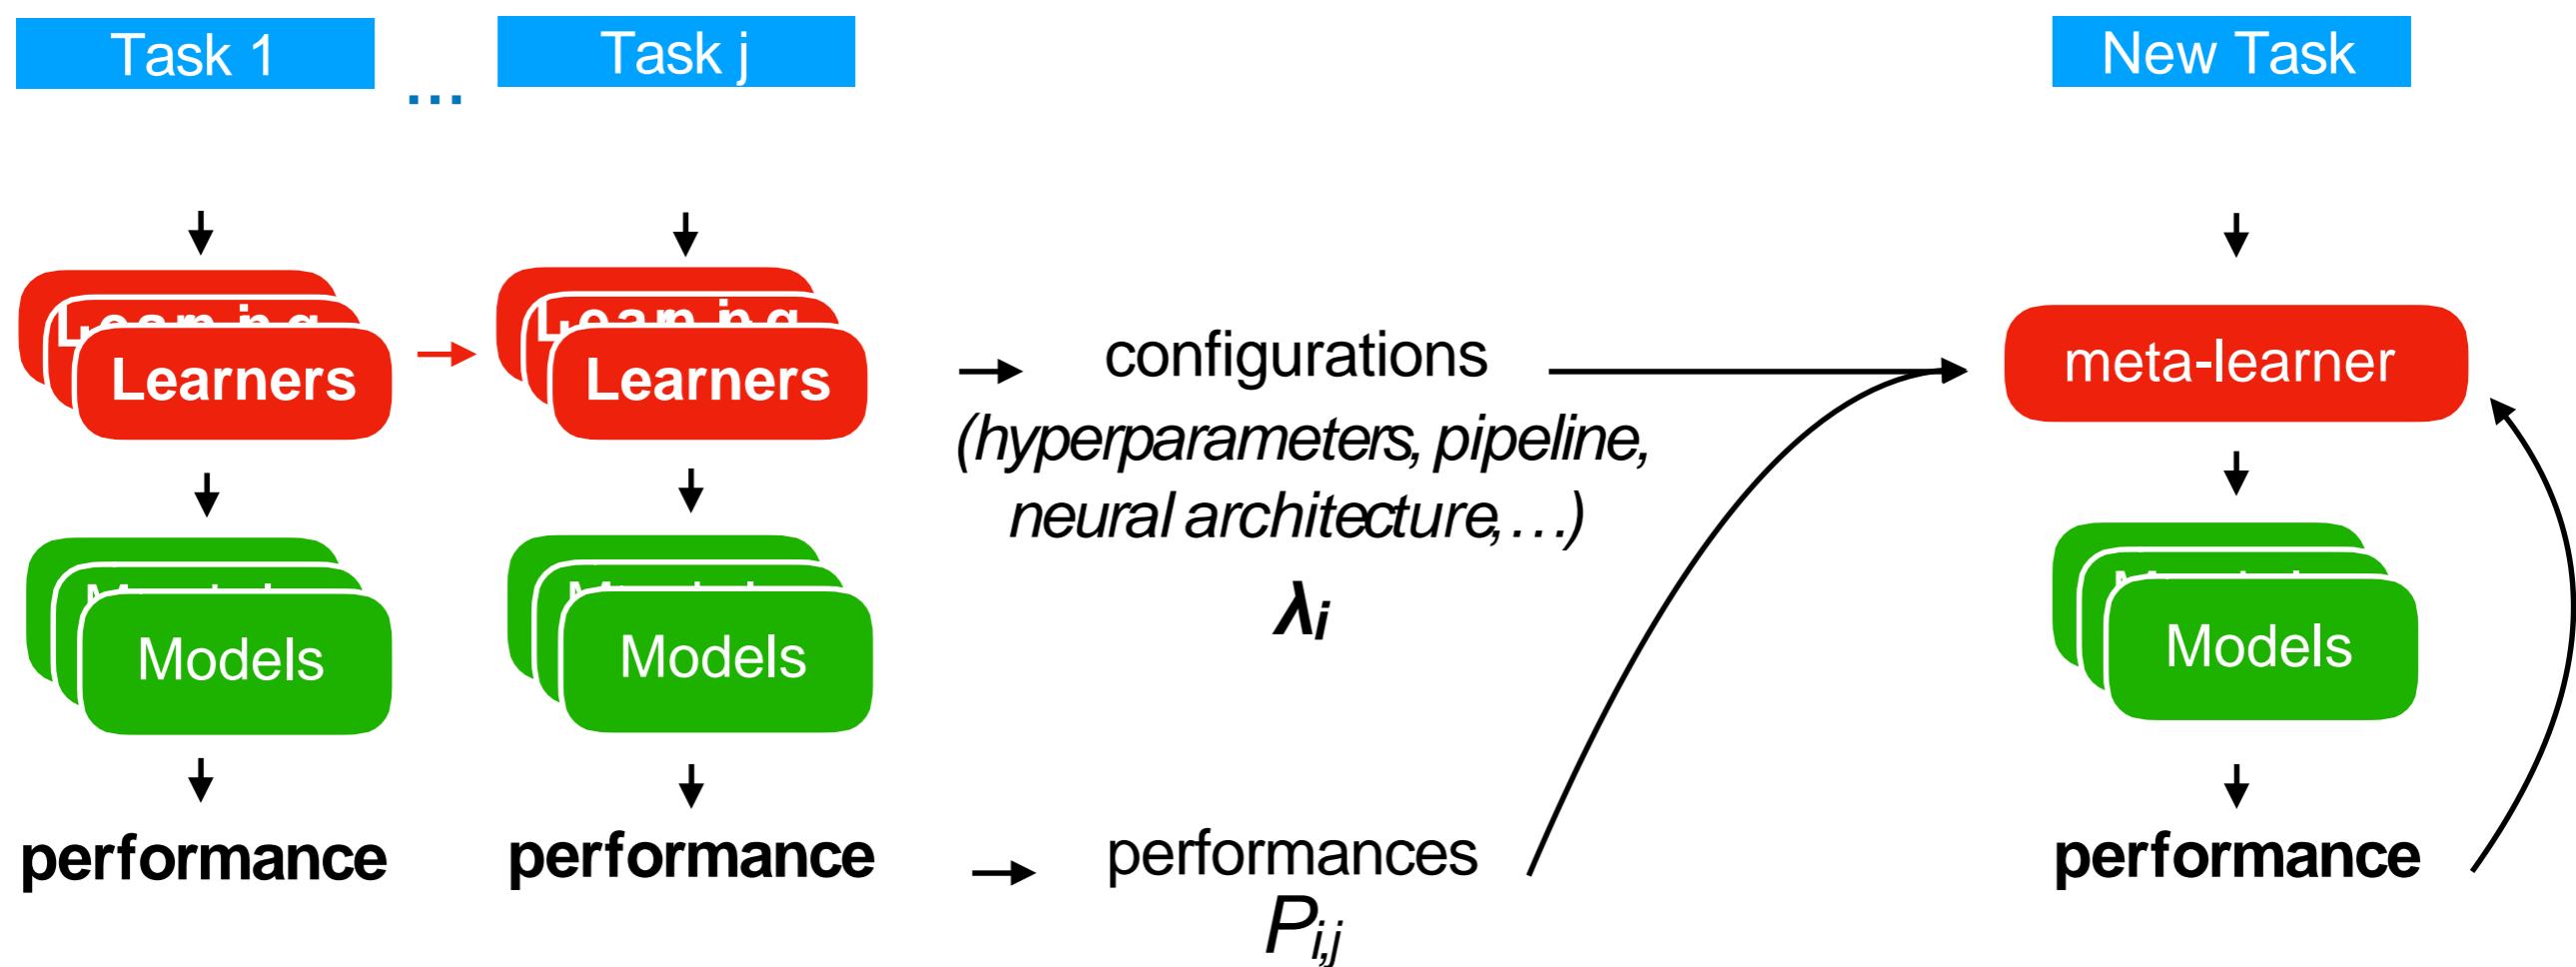
1. Learn from observing learning algorithms (for *any* task) $\lambda_i \ P_{ij}$
2. Learn what may likely work (for *somewhat similar* tasks) $m_j \ \lambda_i \ P_{ij}$
3. Learn from previous models (for *very similar* tasks) $!_{i,j,k} (m_j) \lambda_i \ P_{ij}$



1. Learn from observing learning algorithms

Define, store and use meta-data:

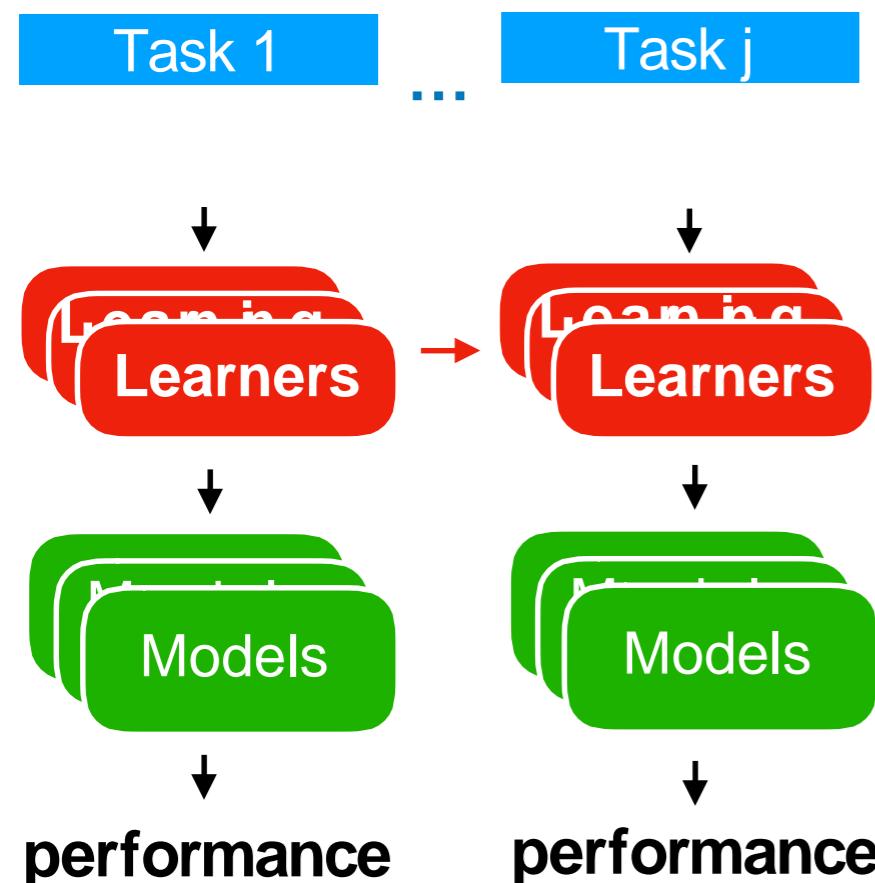
- **configurations**: settings that uniquely define the model
- **performance** (e.g. accuracy) on specific tasks



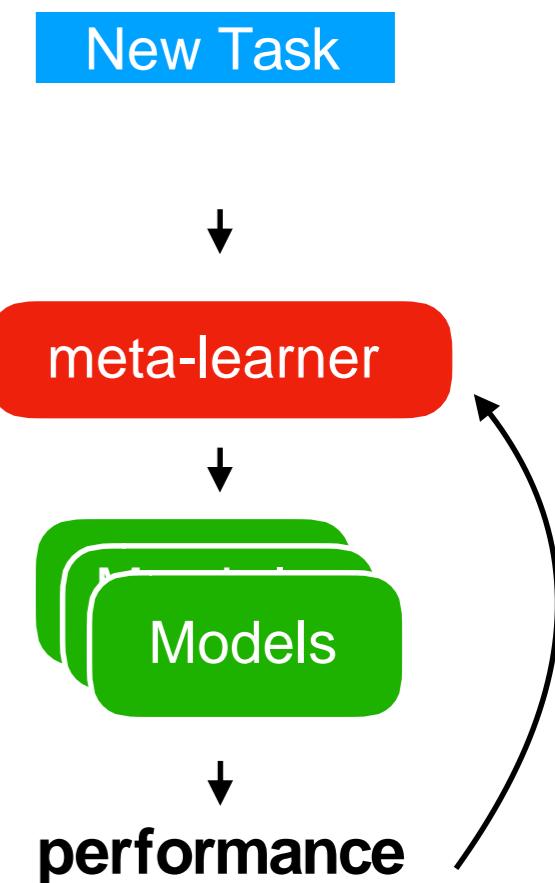
1. Learn from observing learning algorithms

Define, store and use meta-data:

- **configurations**: settings that uniquely define the model
- **performance** (e.g. accuracy) on specific tasks

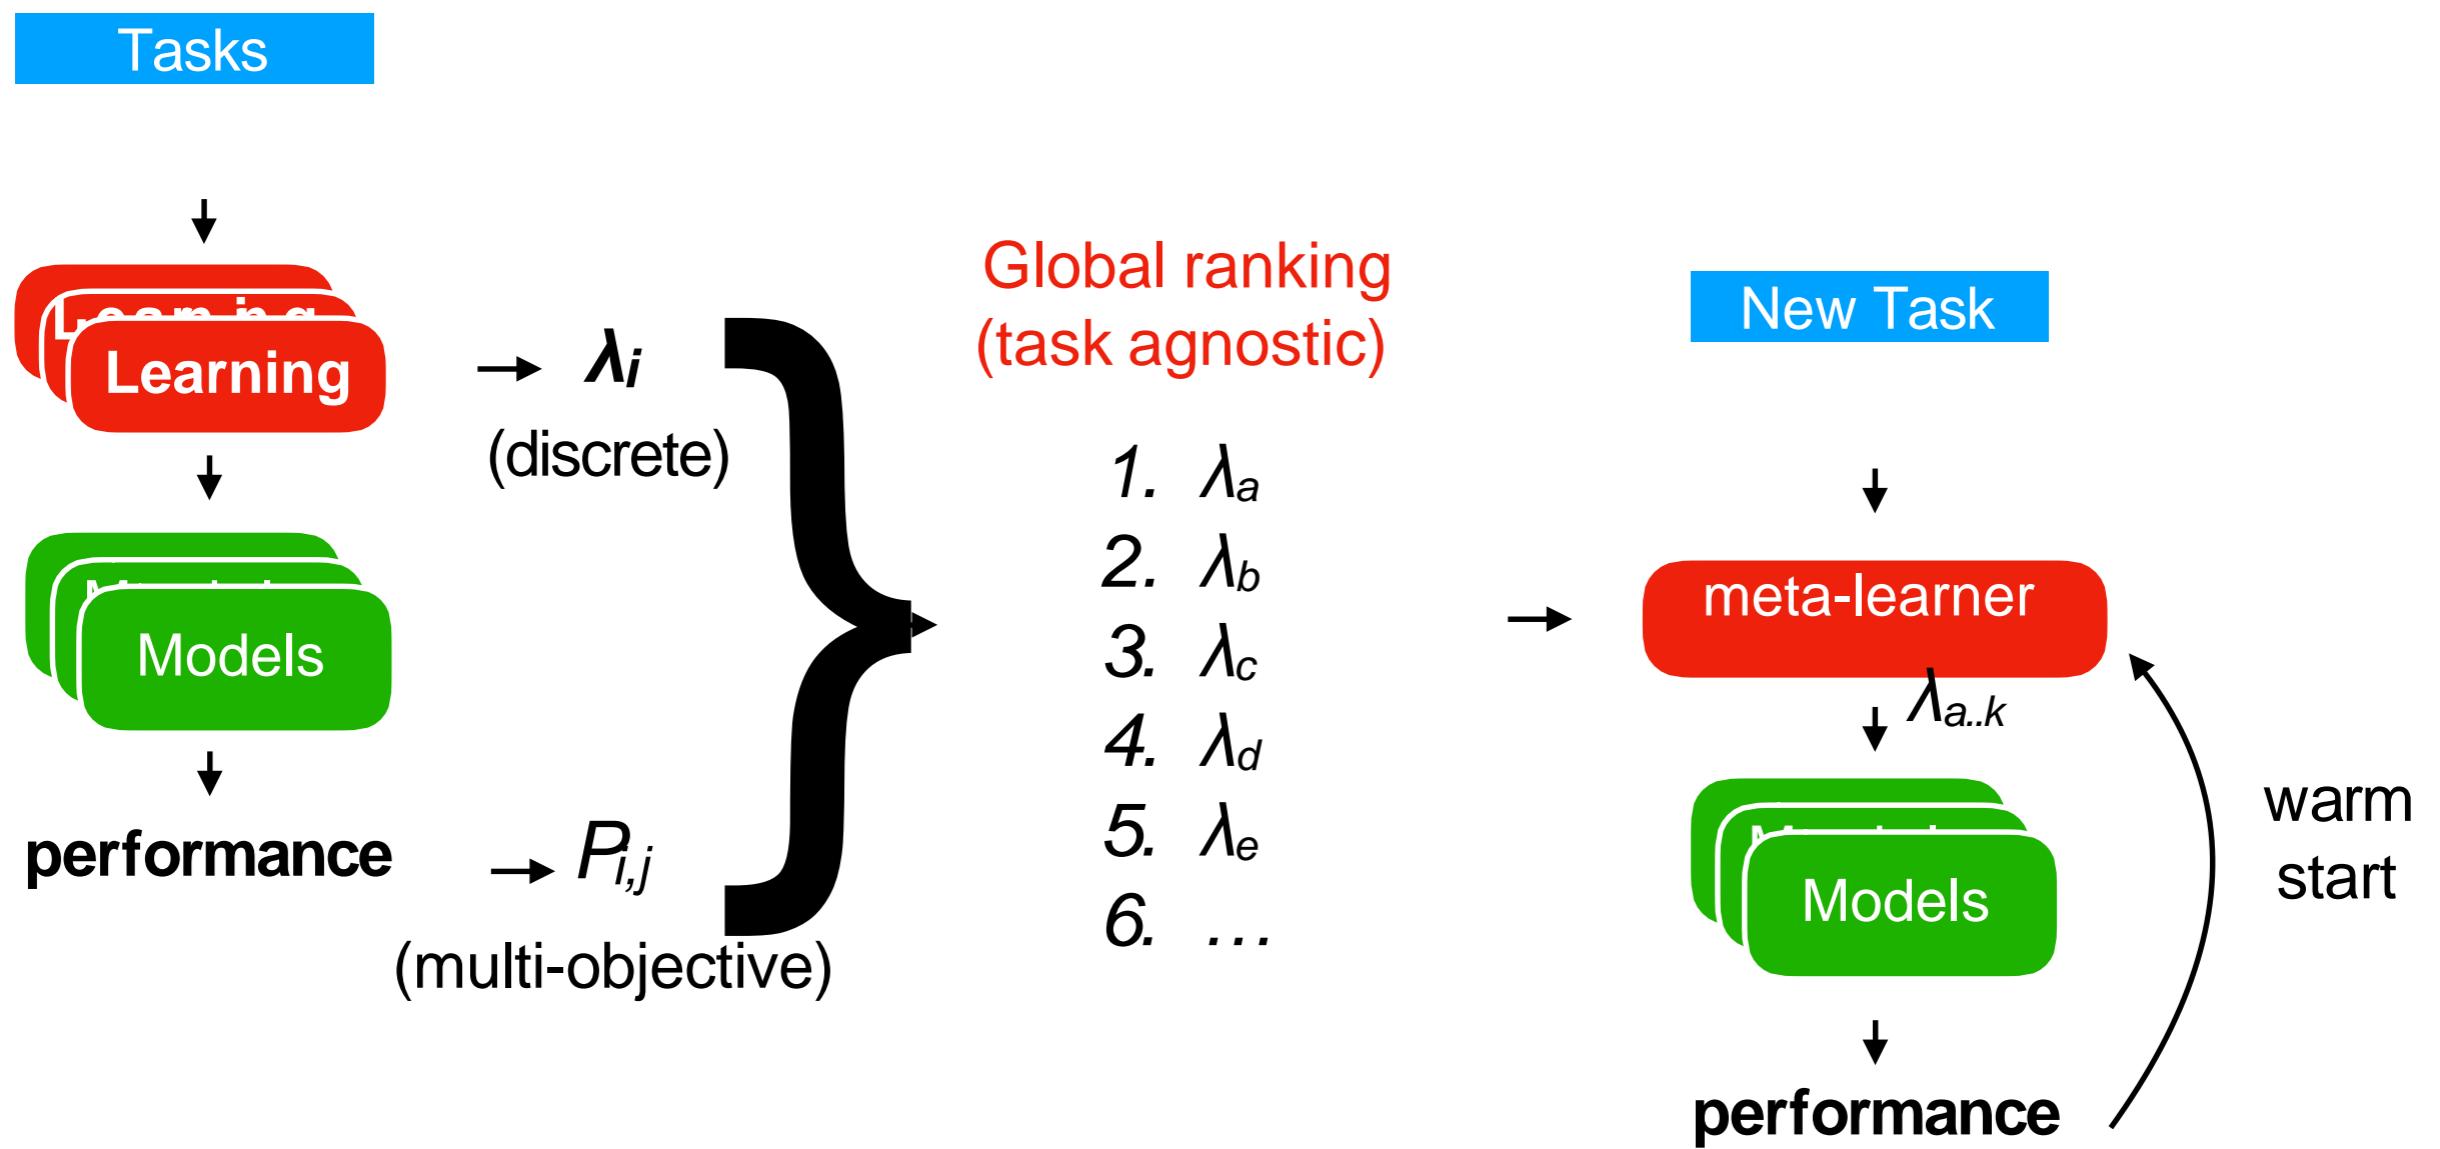


task	model config	λ	score P
0	alg=SVM, C=0.1		0,876
0	alg>NN, lr=0.9		0,921
0	alg=RF, mtry=0.1		0,936
1	alg=SVM, C=0.2		0,674
1	alg>NN, lr=0.5		0,777
1	alg=RF, mtry=0.4		0,791



Rankings

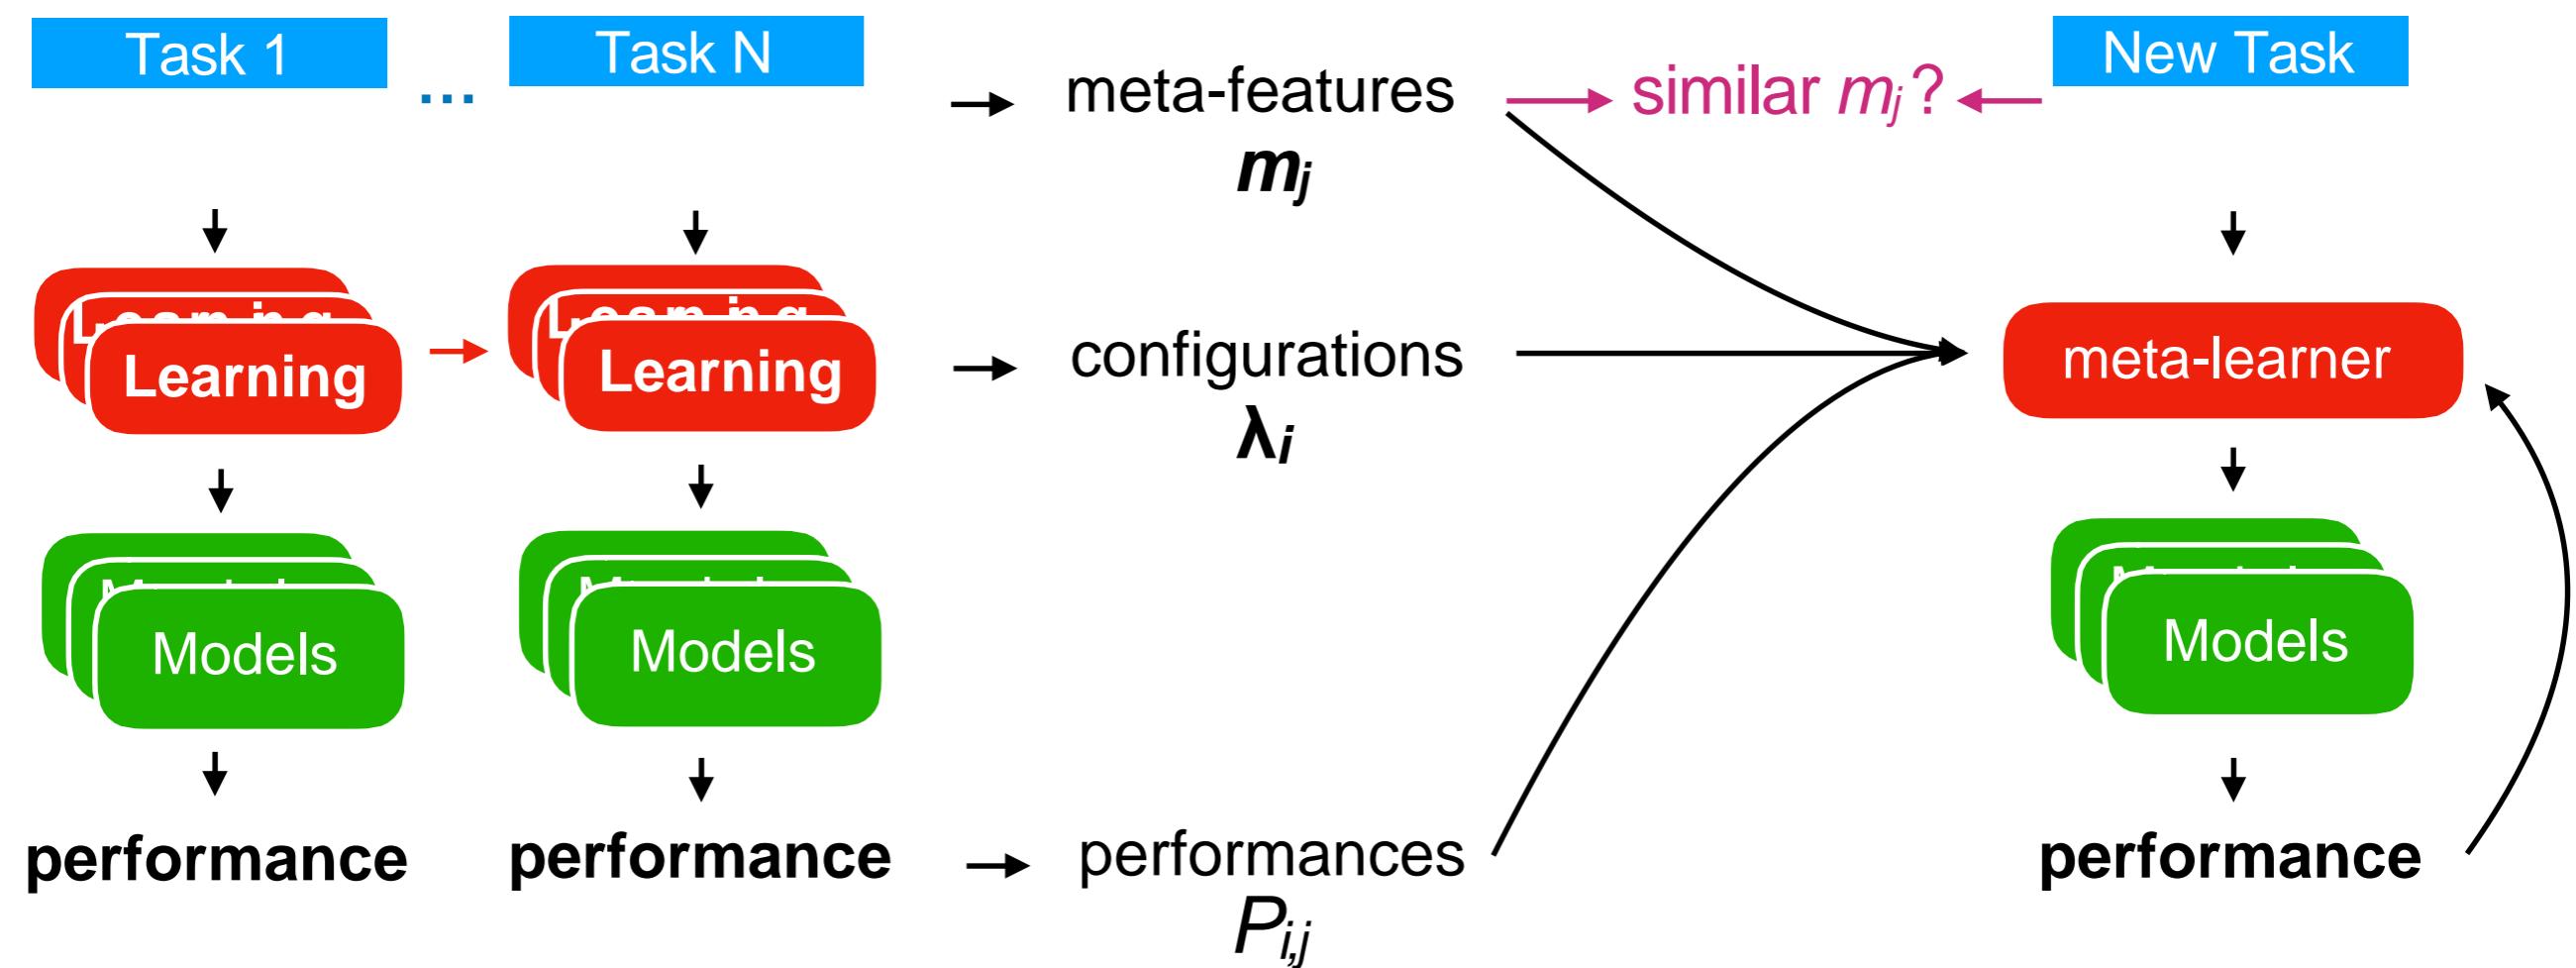
- Build a *global (multi-objective) ranking*, recommend the top-K
- Can be used as a *warm start* for optimization techniques
 - E.g. Bayesian optimization, evolutionary techniques,...



2. Learn what may likely work (for *partially similar* tasks)

Meta-features: measurable properties of the tasks

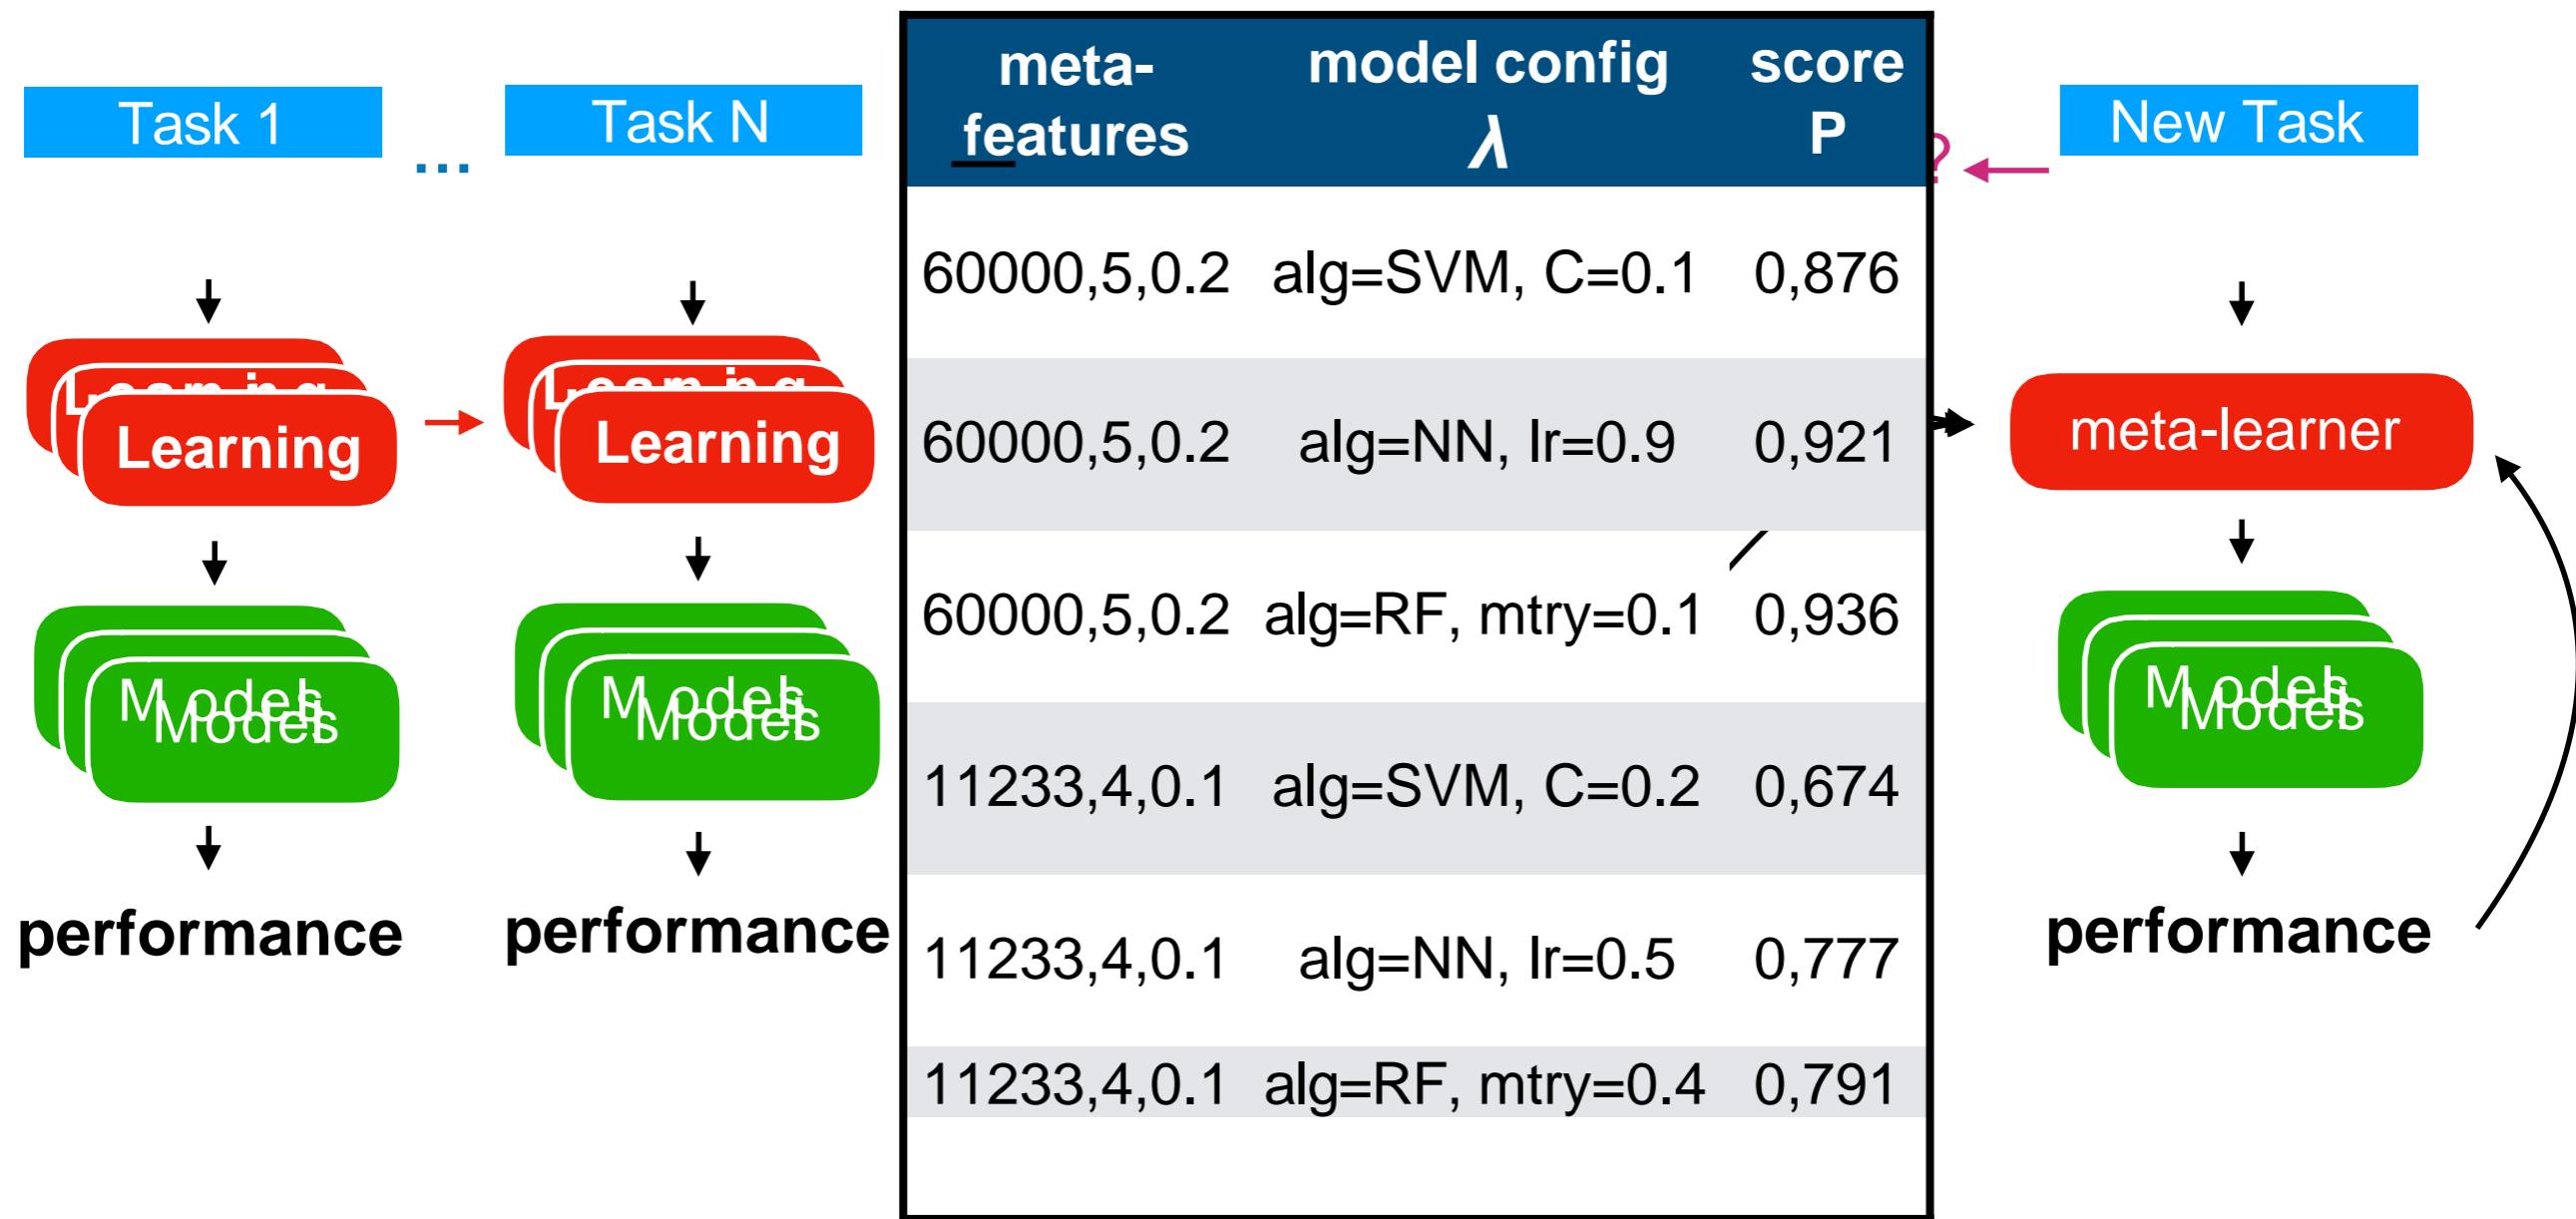
(number of instances and features, class imbalance, feature skewness,...)



2. Learn what may likely work (for *partially similar tasks*)

Meta-features: measurable properties of the tasks

(number of instances and features, class imbalance, feature skewness,...)

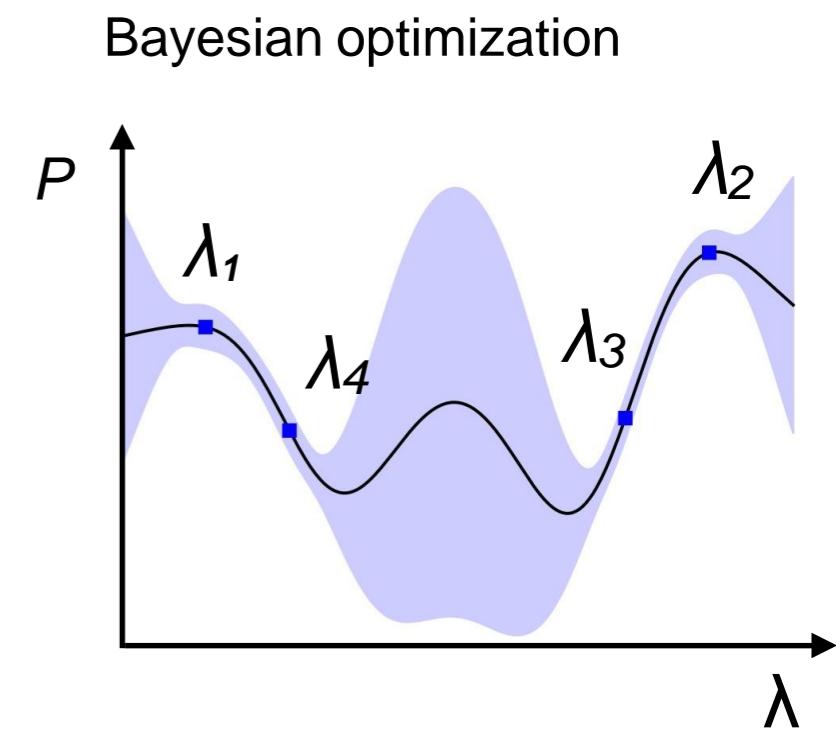
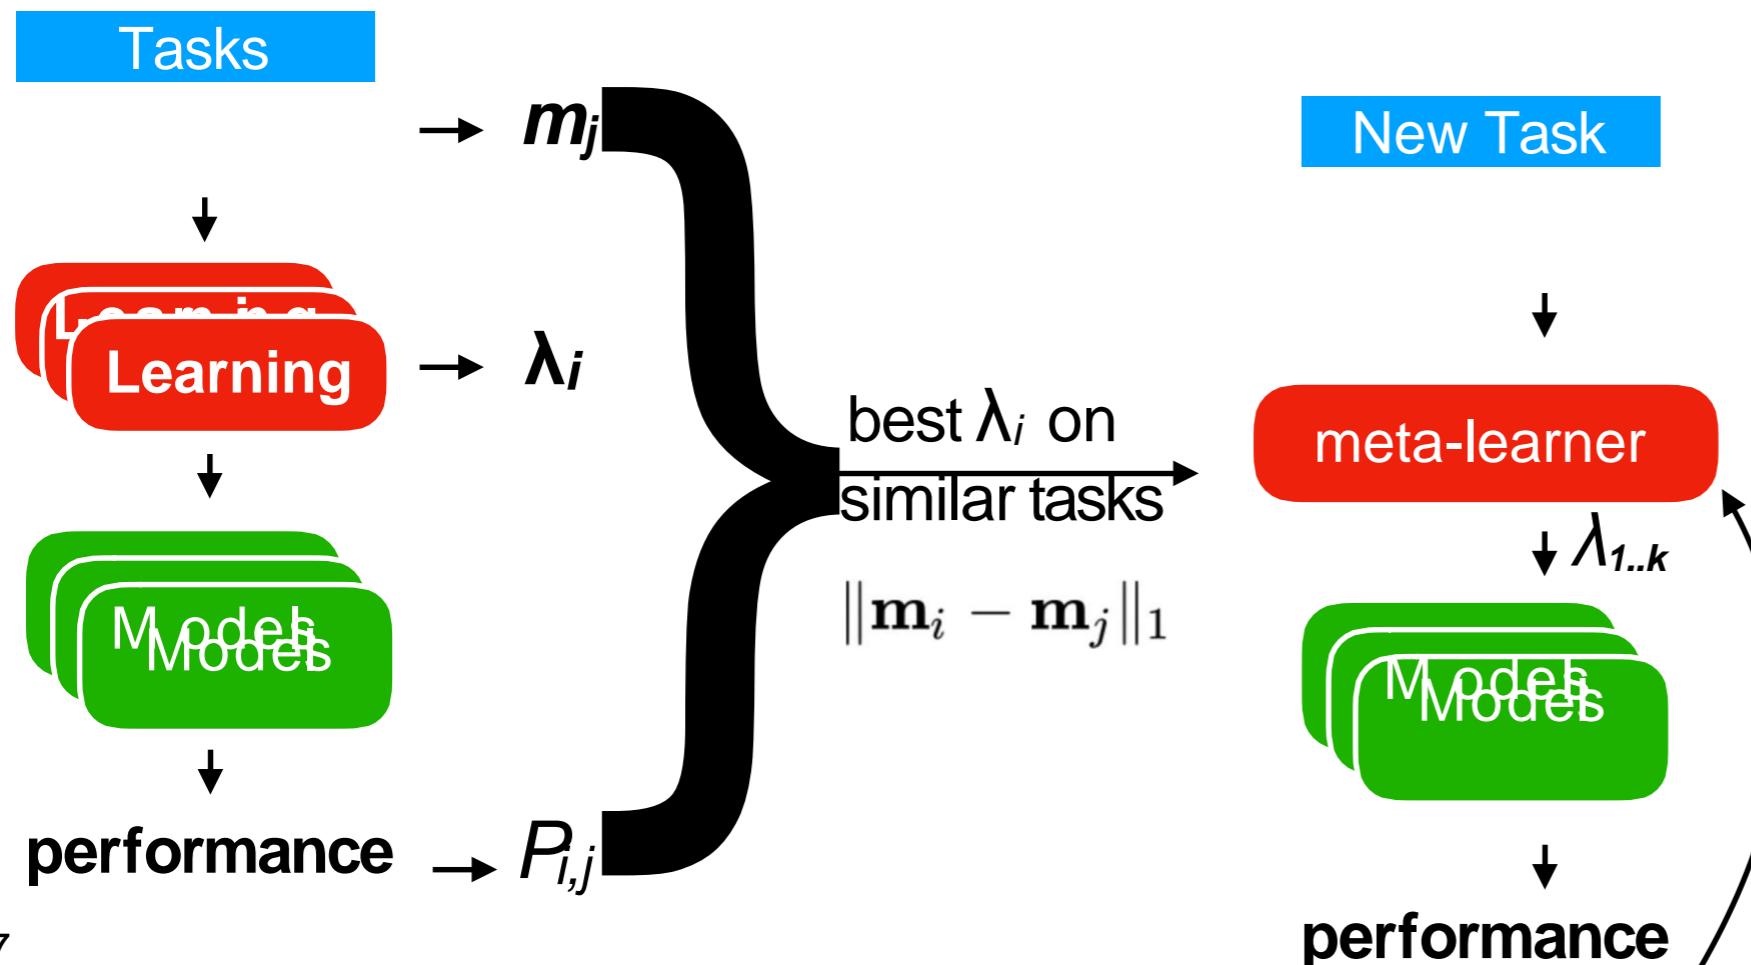


Meta-features

- **Hand-crafted (interpretable) meta-features¹**
 - **Number of** instances, features, classes, missing values, outliers,...
 - **Statistical:** skewness, kurtosis, correlation, covariance, sparsity, variance,...
 - **Information-theoretic:** class entropy, mutual information, noise-signal ratio,...
 - **Model-based:** properties of simple models trained on the task
 - **Landmarkers:** performance of fast algorithms trained on the task
 - Domain specific task properties
- Optimized (recommended) representations exist ²

Warm-starting from similar tasks

- Find k most similar tasks, warm-start search with best λ_i
 - Auto-sklearn: Bayesian optimization (SMAC)
 - Meta-learning yield better models, faster
 - Winner of AutoML Challenges



Auto-Sklearn Example

```
[ ] import autosklearn.classification
import sklearn.model_selection
import sklearn.datasets
import sklearn.metrics

X, y = sklearn.datasets.load_digits(return_X_y=True)

X_train, X_test, y_train, y_test = \
    sklearn.model_selection.train_test_split(X, y, random_state=1)

automl = autosklearn.classification.AutoSklearnClassifier(time_left_for_this_task=60*5) #Auto-sklearn searches pipelines for 5 minutes

automl.fit(X_train, y_train)

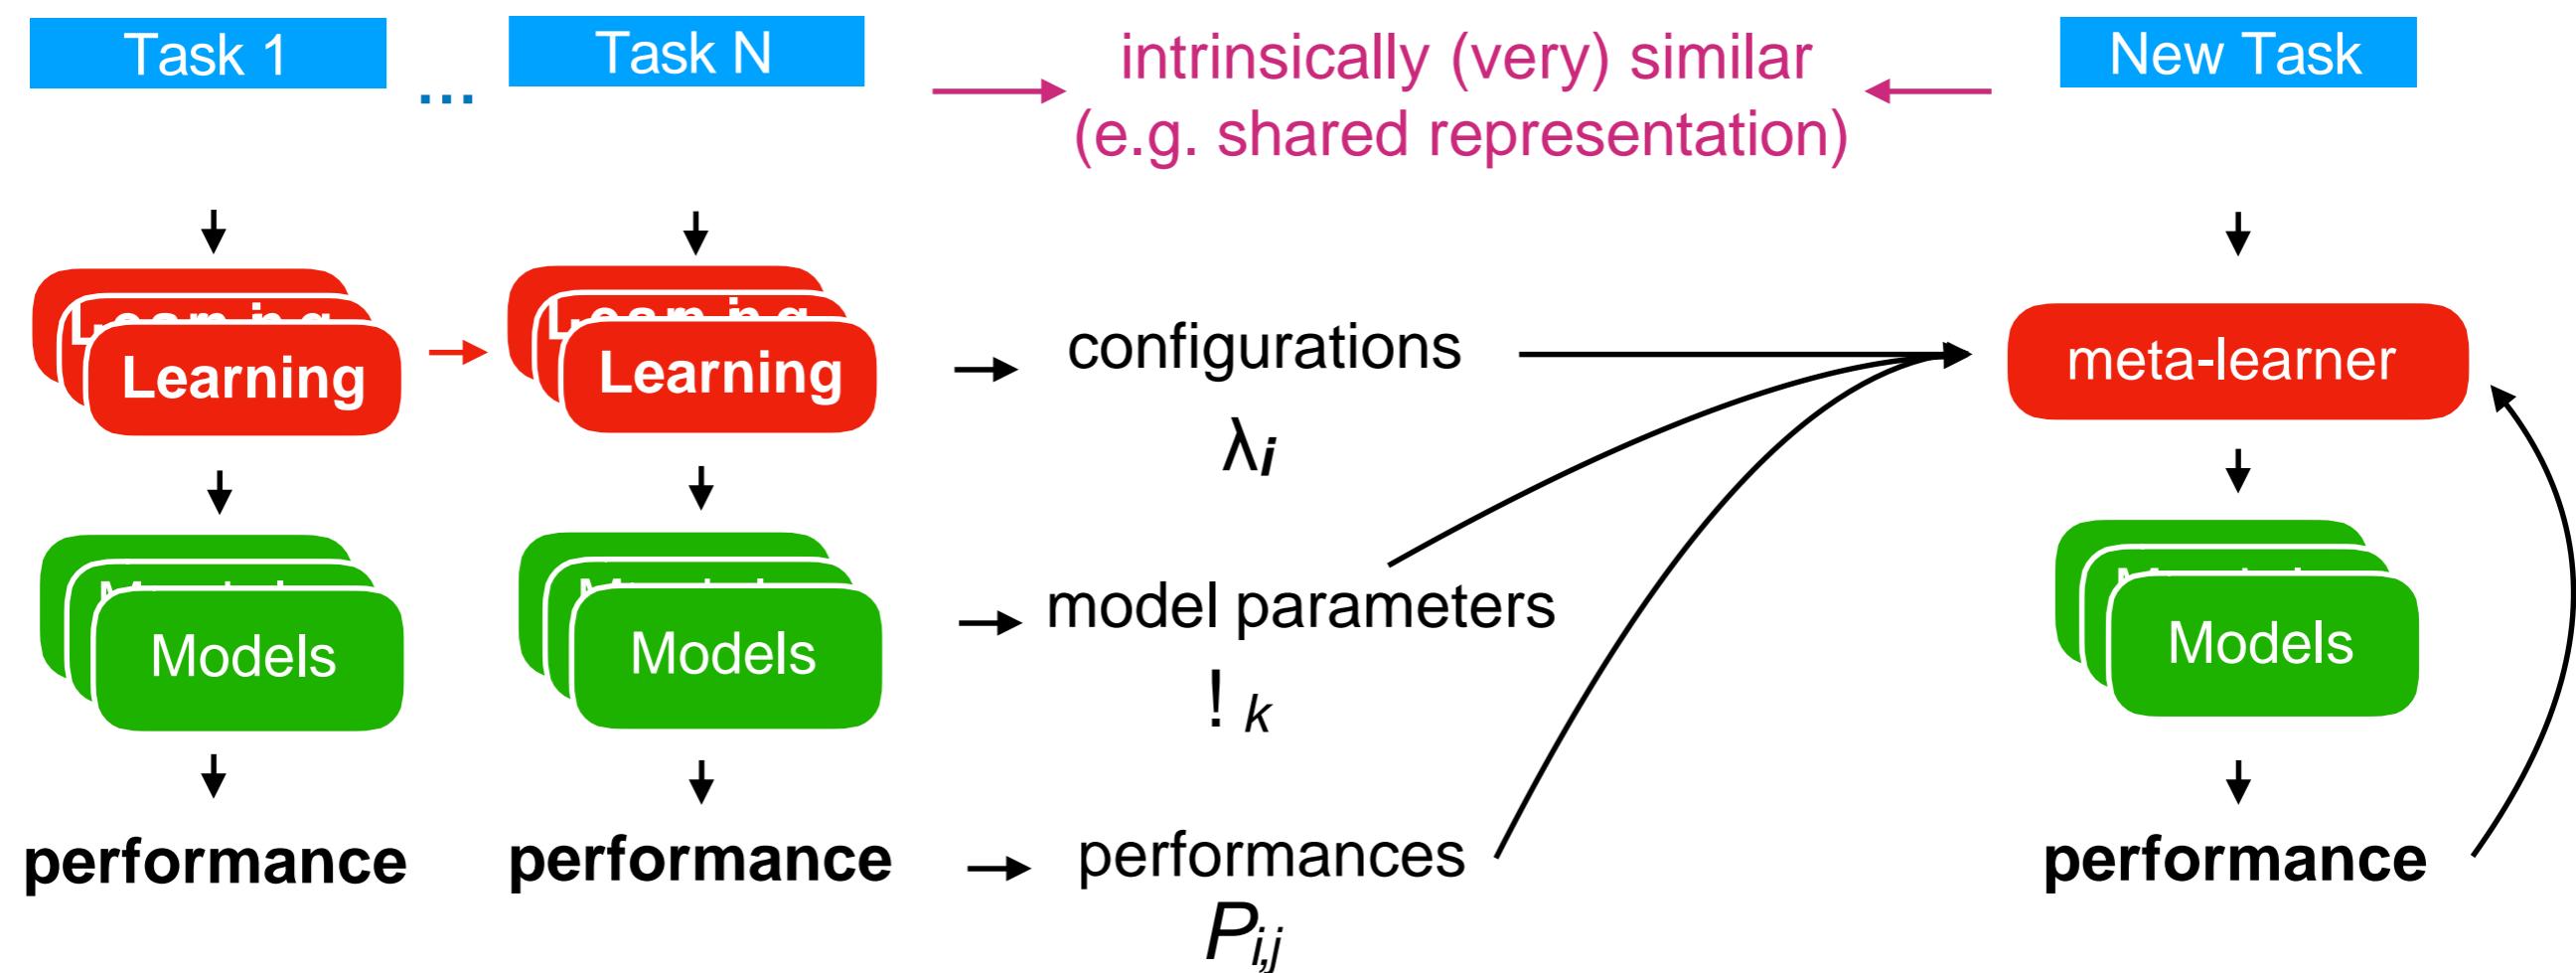
y_hat = automl.predict(X_test)

print("Accuracy score", sklearn.metrics.accuracy_score(y_test, y_hat))
```

3. Learn from previous *models* (for *very similar* tasks)

Models trained on *intrinsically similar* tasks

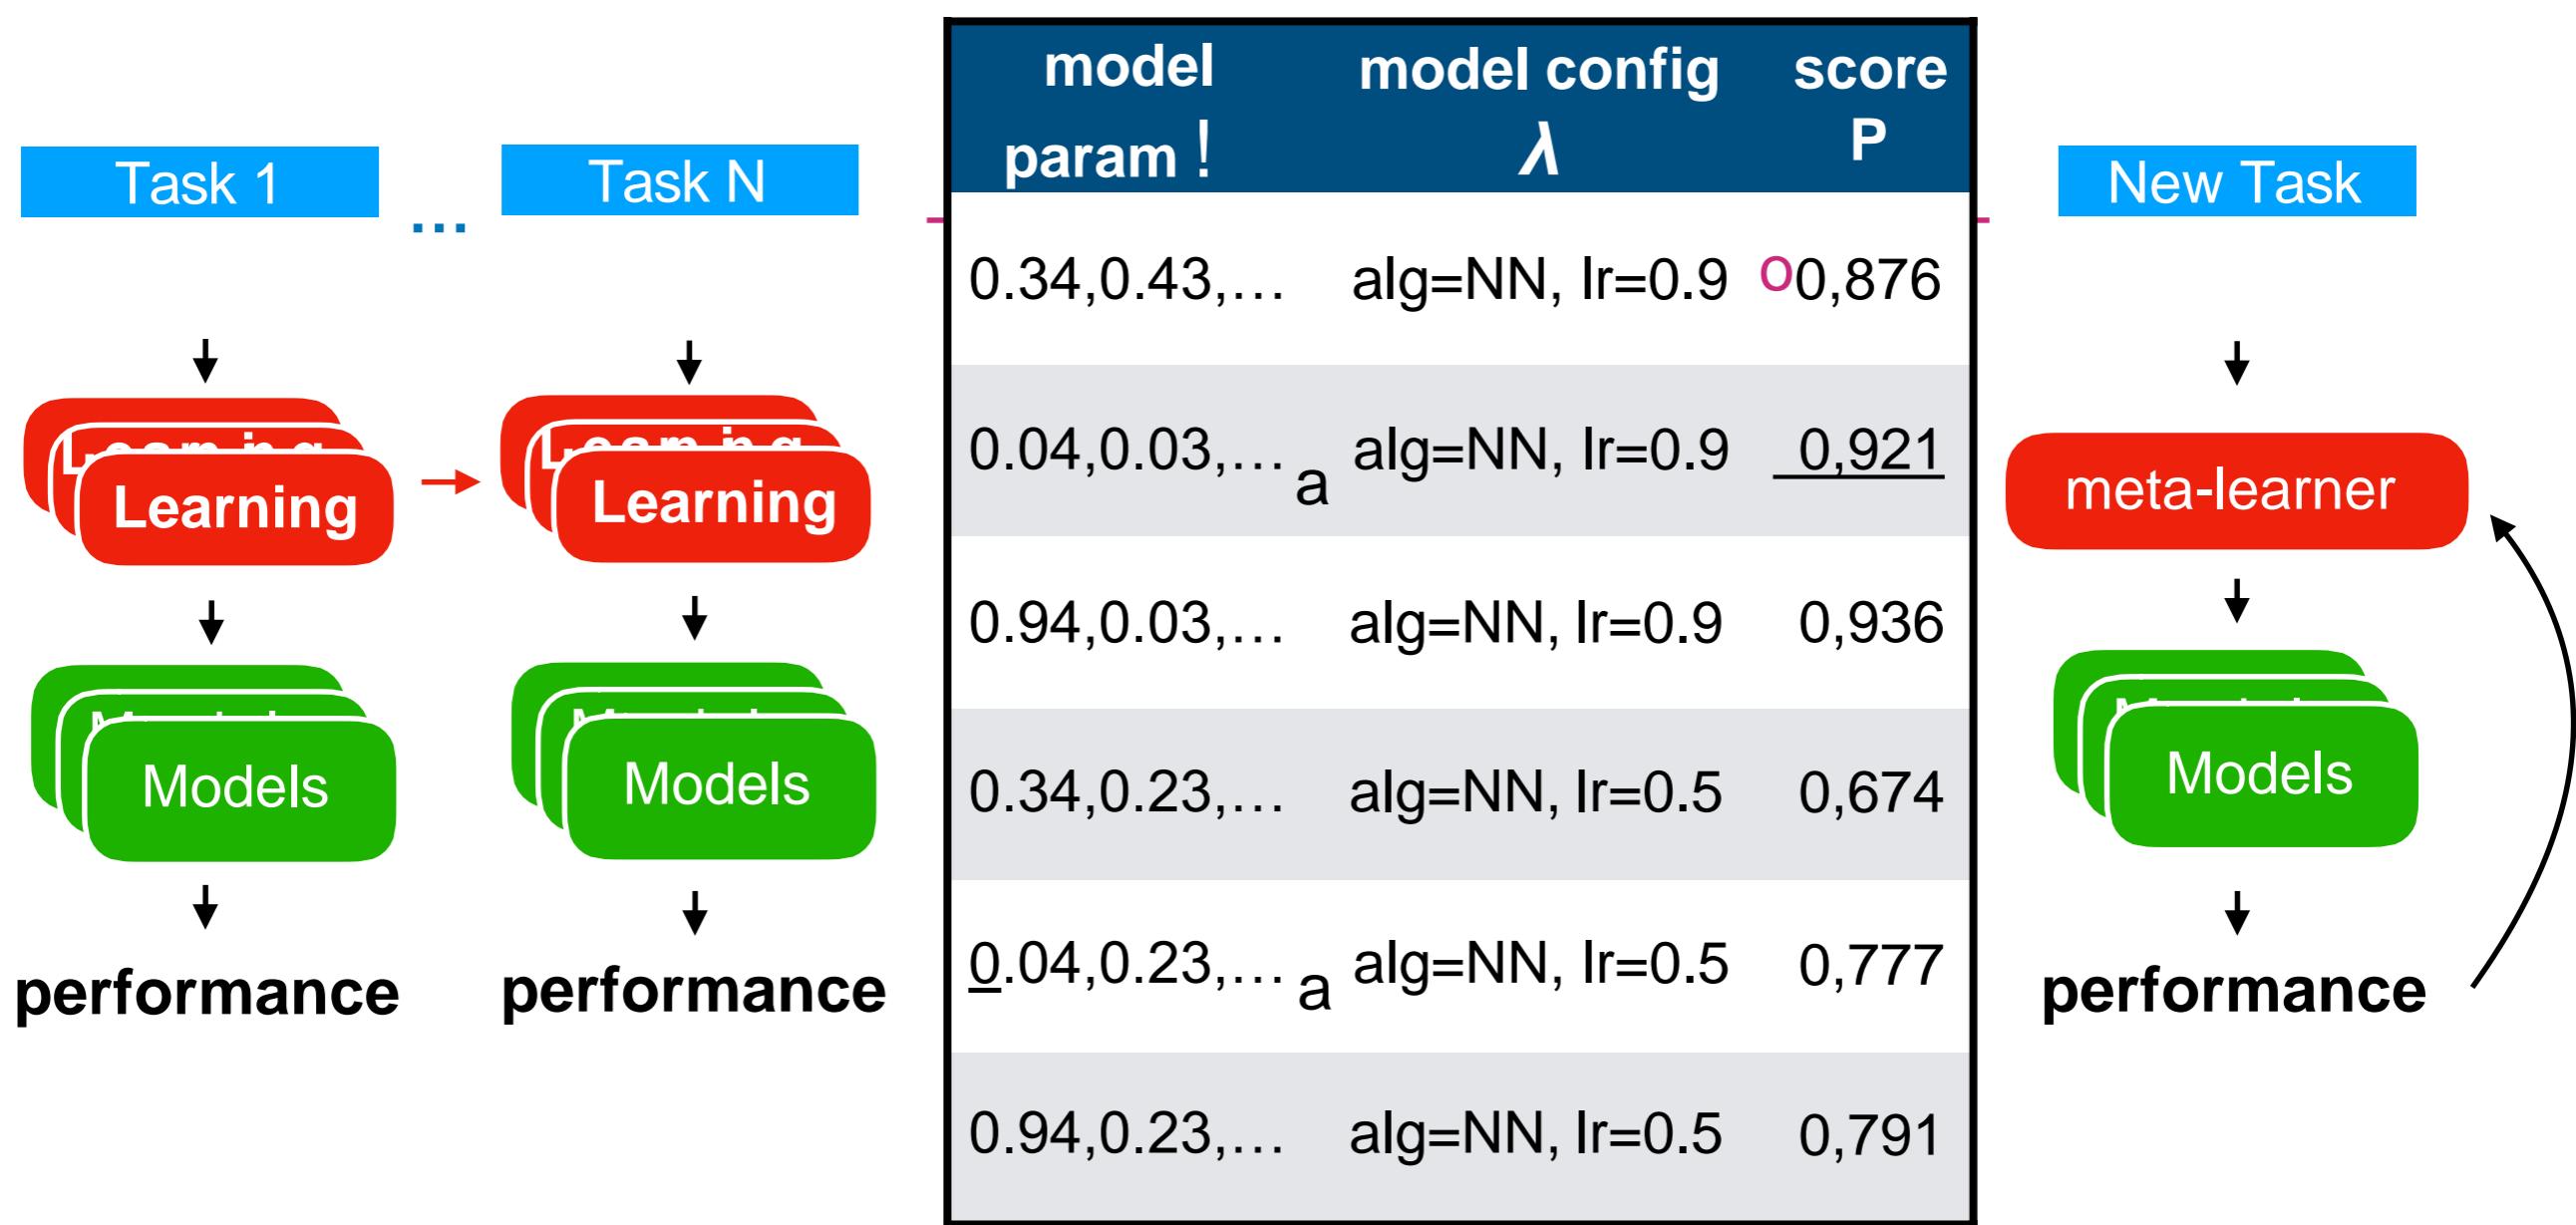
(model parameters, features,...)



3. Learn from previous *models* (for *very similar* tasks)

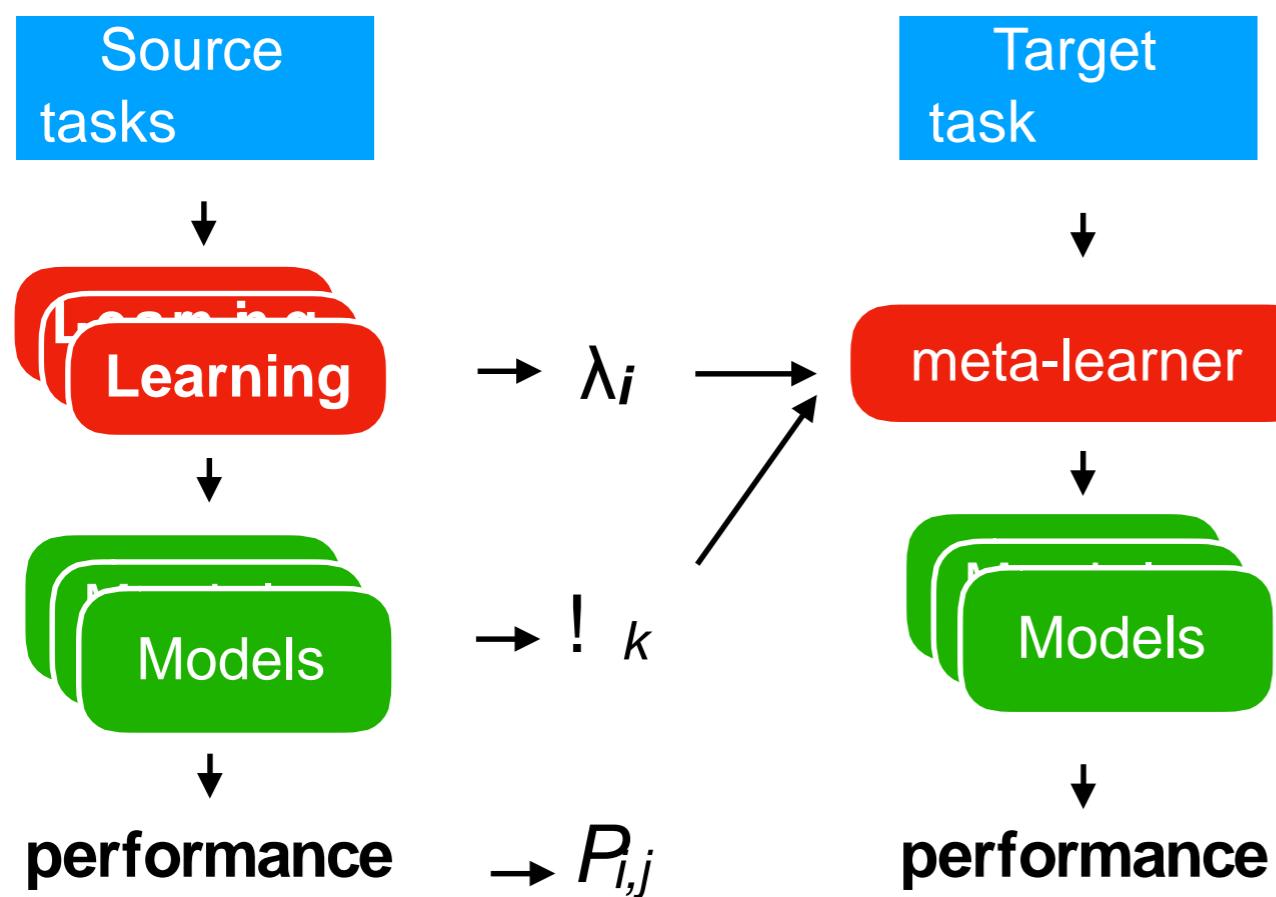
Models trained on *intrinsically similar* tasks

(model parameters, features,...)

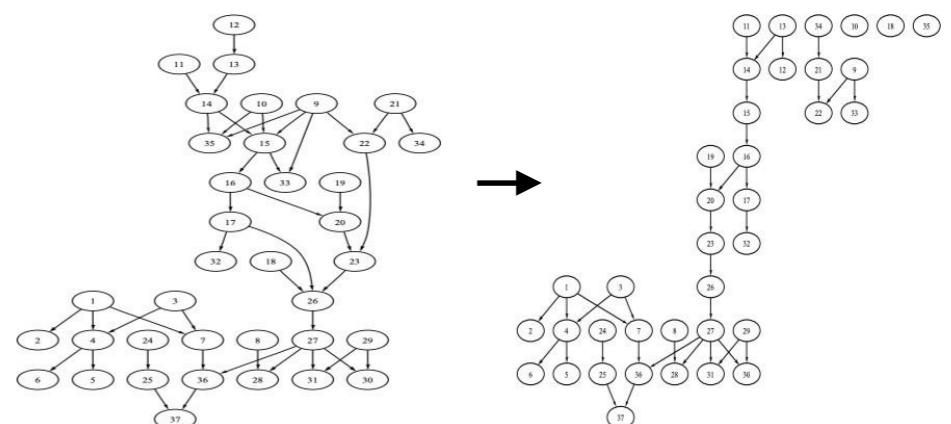


Transfer Learning

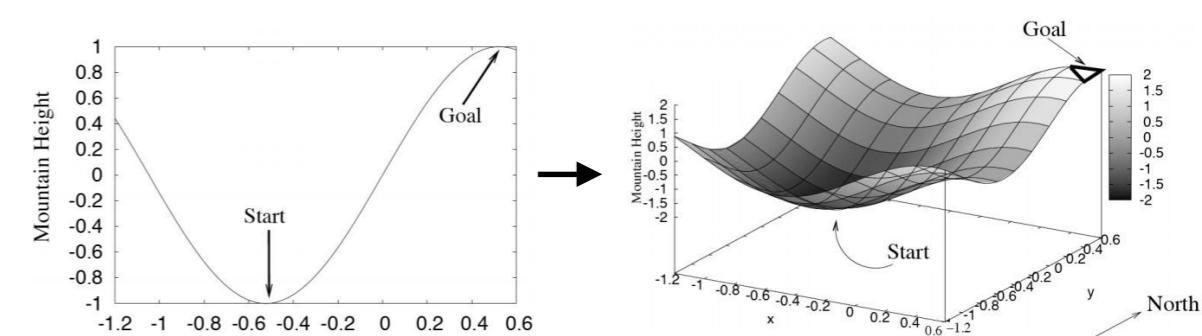
- Select source tasks, transfer trained models to similar target task ¹
- Use as starting point for tuning, or *freeze* certain aspects (e.g. structure)
 - Bayesian networks: start structure search from prior model ²
 - Reinforcement learning: start policy search from prior policy ³



Bayesian Network transfer

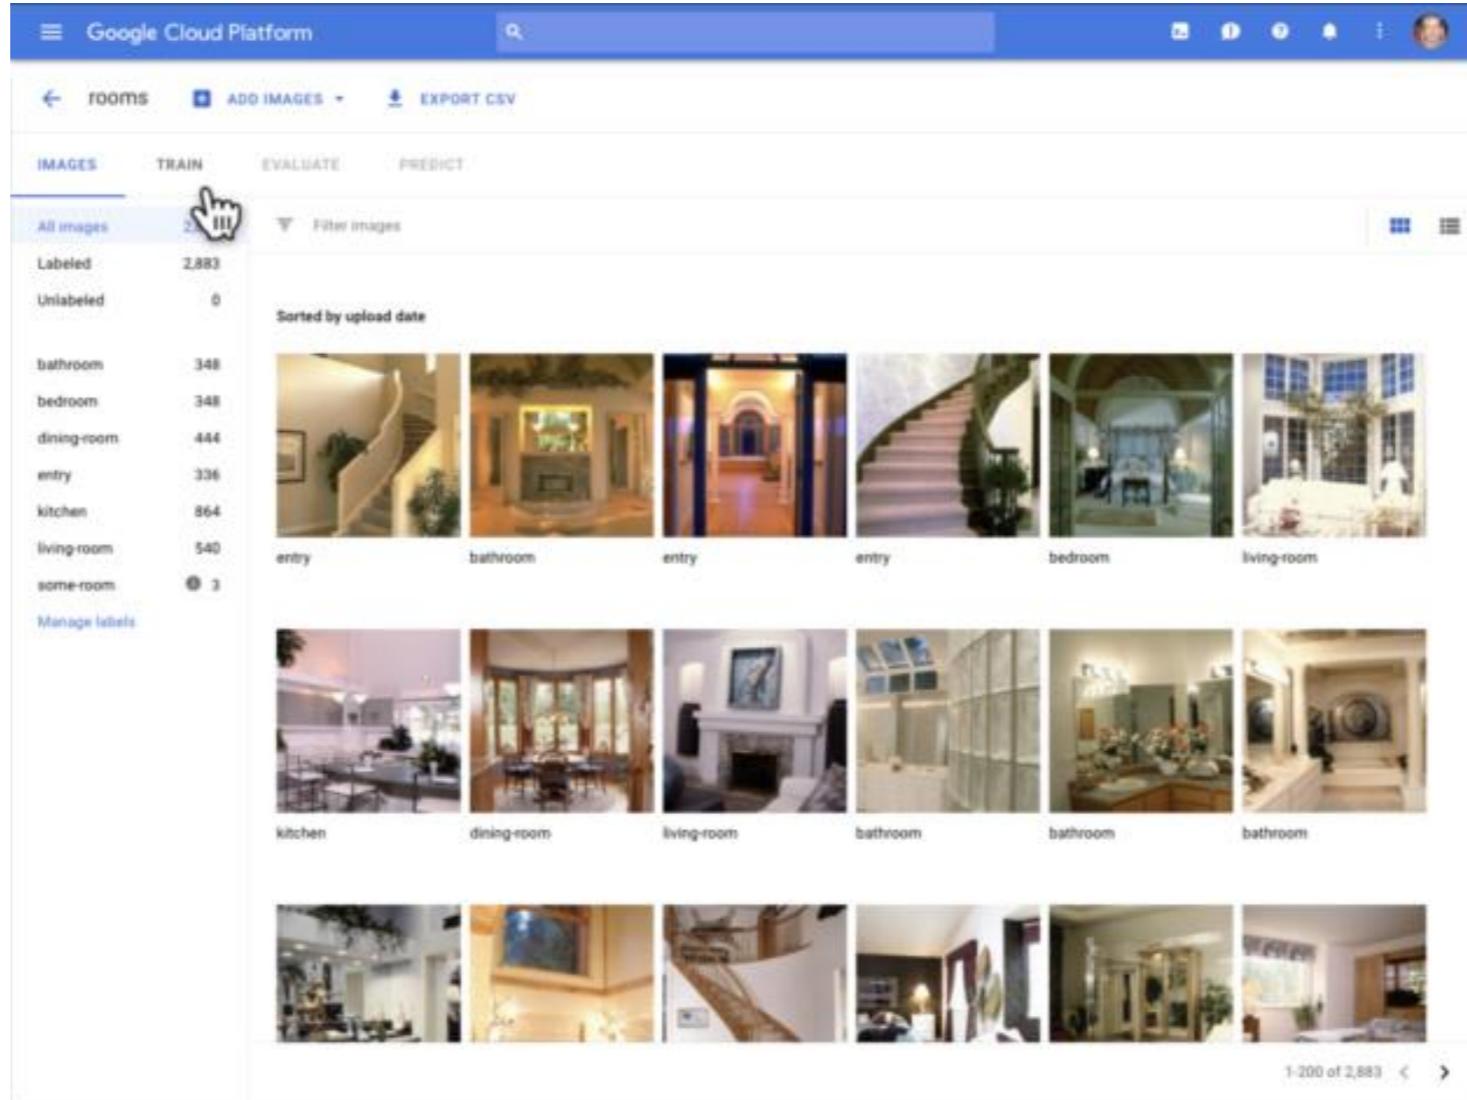


Reinforcement learning: 2D to 3D mountain car



Google Cloud AutoML

- Combines Neural Architecture Search with Transfer Learning



https://cloud.google.com/vision/automl/docs/samples/automl-vision-classification-create-model?authuser=1#automl_vision_classification_create_model-python

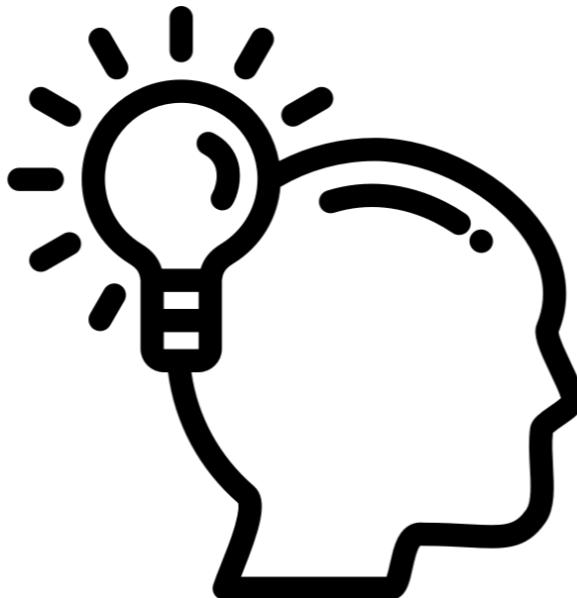
<https://www.fast.ai/2018/07/23/auto-ml-3/>

Human in the loop

AutoML

AutoML challenges

- Requires programming background
- Black box on top of Black box (hard to explain)
- No user feedback (how to adapt the models based on human knowledge)



Visual Analytics tools were created to tackle some of these problems

Visus

Interface between AutoML and user.
Model explanation / Score inspection

A

The screenshot shows a list of datasets available for selection:

- 100_baseball_dataset
- 22_handgeometry_dataset
- 26_radon_neut_dataset
- 31_urbansound_dataset
- AUDATA_taxi_vehicle_collisions_daily_dataset
- usd_SPECT_dataset

Each dataset entry includes a brief description and a "Select" button.

B

Select Task
Users choose one of the suggested problems or specify a new one.

C

Define Problem

Target: 0:n. collisions
Type: Regression
Sub-type: [empty]

Data Profiling Data Augmentation

Tabular

0: Table
Num Rows: 365
Datetime x N. trips x N. collisions x

Display only: Datetime Integer

datetime: Period: 1/1/2014 to 12/31/2014
n. trips: Distinct Values: 19
Count: Distinct Values: 18

D

Search Settings

Max Solutions: 200
Max Time: 60 Minutes
Evaluation Metrics: Accuracy, Precision, Recall, F1, F1-Micro, F1-Macro, Mean Squared Error

> Start Solutions Search

C1

101281.mp3 0: 00:00:00 / 01:27:066

Audio: Waveform plot

Time Series: Line plot showing series data over time

Images: Three small images labeled 001_HandPhoto_left_01.jpg, 002_HandPhoto_left_01.jpg, and 003_HandPhoto_left_01.jpg, each showing a hand interacting with a circular object.

Visus



Visus

Welcome!

What would you like to know? The price of a 780 ft², 1 bedroom apartment in your neighborhood? Your dog's life expectancy? Weather on next Monday?

Visus helps you to answer your questions using existing datasets and state-of-the-art machine learning techniques.

Visus presents step-by-step guidance which enables you to build your own Machine Learning model without knowing how machine learning works.

[» USE VISUS](#)



Visus

Challenges Addressed

- Requires programming background
- Black box on top of Black box (hard to explain)
- No user feedback (how to adapt the models based on human knowledge)

Limitations:

- User cannot edit the pipelines
- Linear flow

Two Ravens

AutoML user interface Feature selection

The screenshot shows the Two Ravens AutoML user interface. At the top, there is a navigation bar with tabs for Dataset Description, config_2019-12-20_10-41-34_ydf4 (workspace: config_2019-12-20_10-41-34_ydf4 (id: 254)), problem 52, Model, Explore, Results, and dev_admin.

Data Selection: This panel contains a list of variables under 'Variables' tab. Variables highlighted in orange include: forest_ih, goldplacer_y, goldvein_y, goldsurface_y, grass_ih, gwarea, gwno, irrig_max, irrig_min, irrig_sd, irrig_sum, nlights_calib_mean, pasture_ih, petroleum_y, pop_gpw_sum, pop_hyd_sum, prec_gpcc, prec_gpcp, savanna_ih, shrub_ih, temp, and urban_ih. Variables highlighted in light blue include: gwno. A summary statistics table for nlights_calib_mean is shown, listing Mean (0.03545), Median (0.03548), Mode Values (0.0327921), Mode Frequency (448), Least Freq (0.03320171), and Least Freq Occurrences (1).

Model Selection: This panel includes sections for Problem (set to regression), Task Type (multivariate), Task Subtype (multivariate), Primary Performance Metric (meanSquaredError), Secondary Performance Metrics (rSquared, rootMeanSquaredError, meanAbsoluteError), Split Options, Search Options, and Score Options.

Network Graph: A network graph visualization in the center shows nodes connected by arrows. Nodes include civconf (red circle), bdist3 (purple circle), cmr_mean (grey circle), and gwno (orange circle). An arrow points from civconf to bdist3.

Legend: The legend indicates that light blue circles represent Dep Var, orange circles represent Nominal, and teal circles represent Predictors.

D'Orazio, Vito, Marcus Deng, and Michael Shoemate. "Tworavens for event data."

Two Ravens



Problem Specification and Model Construction

Two Ravens

Challenges Addressed:

- Requires programming background
- Black box on top of Black box (hard to explain)
- No user feedback (how to adapt the models based on human knowledge)

Limitations:

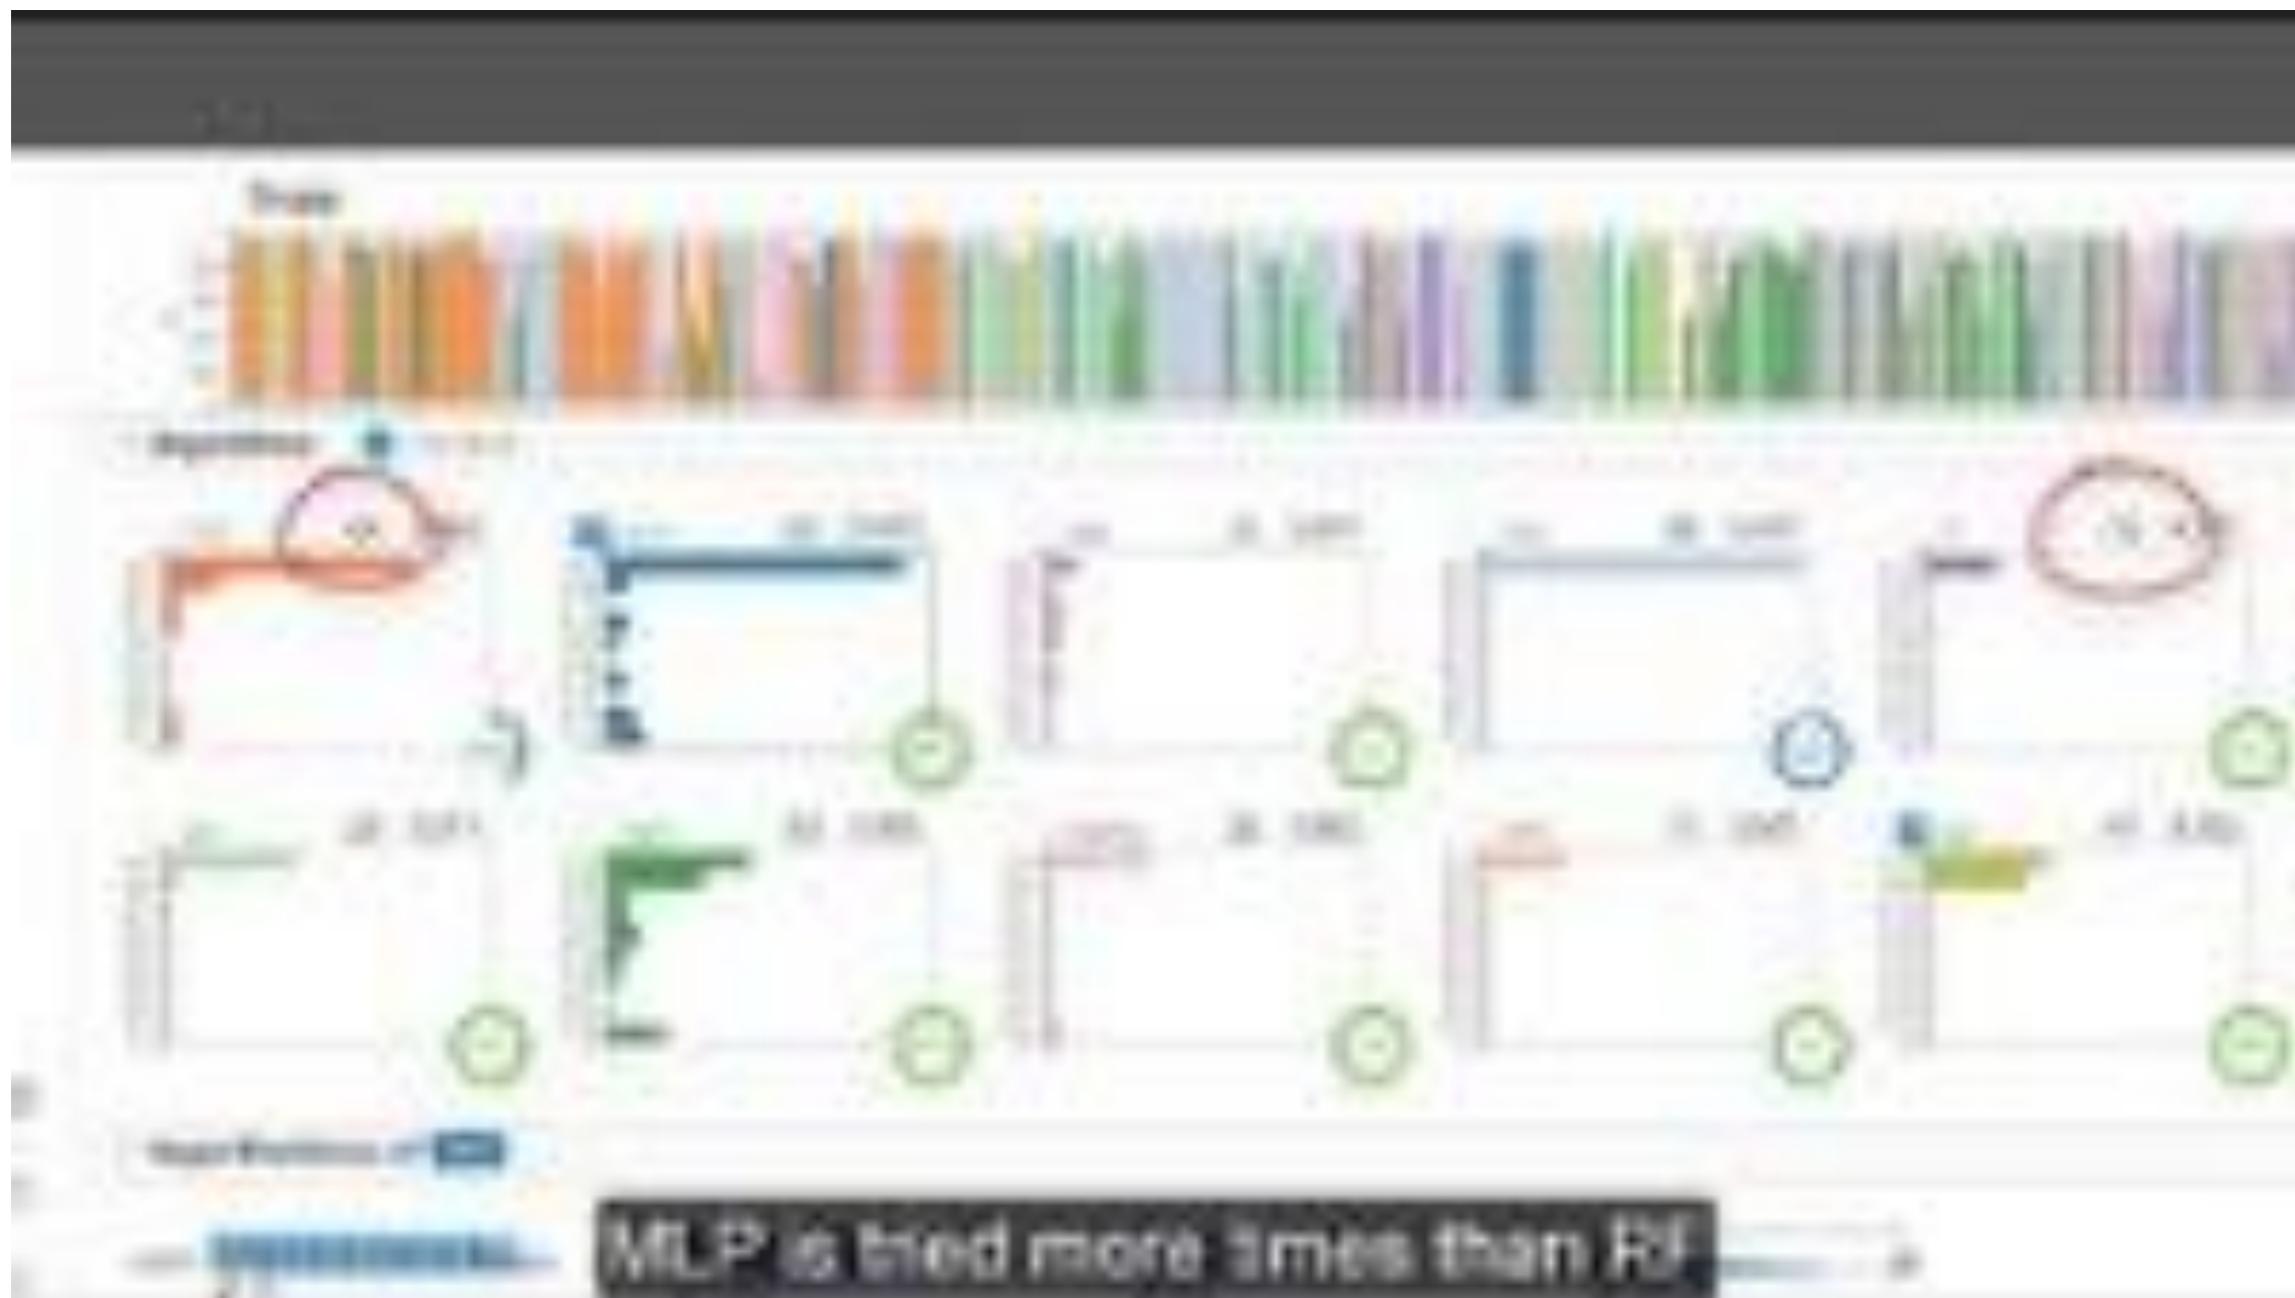
- No model explanation features
- Does not scale well for multiple features

ATMSeer

Interface between AutoML and user.
Lets the user select new hyperparameters/algorithms to try



ATMSeer



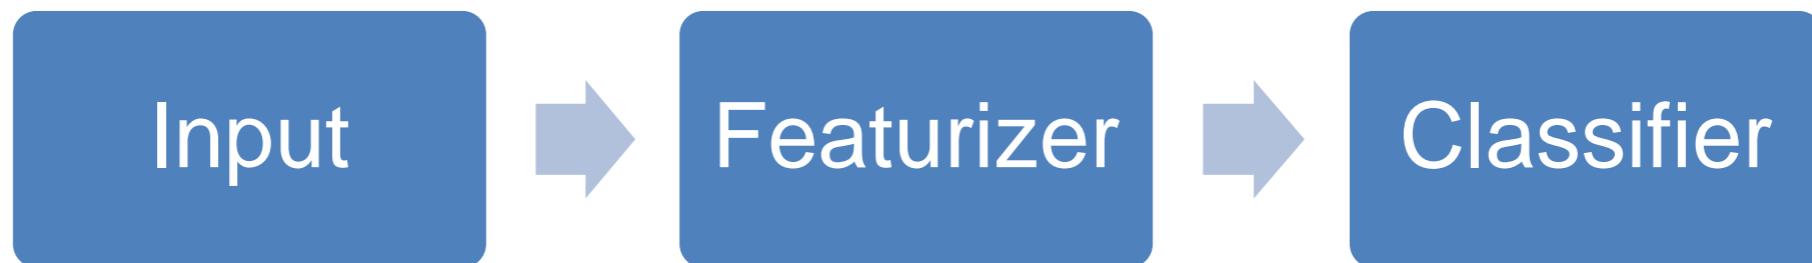
ATMSeer

Challenges Addressed

- Requires programming background
- Black box on top of Black box (hard to explain)
- No user feedback (how to adapt the models based on human knowledge)

Limitations:

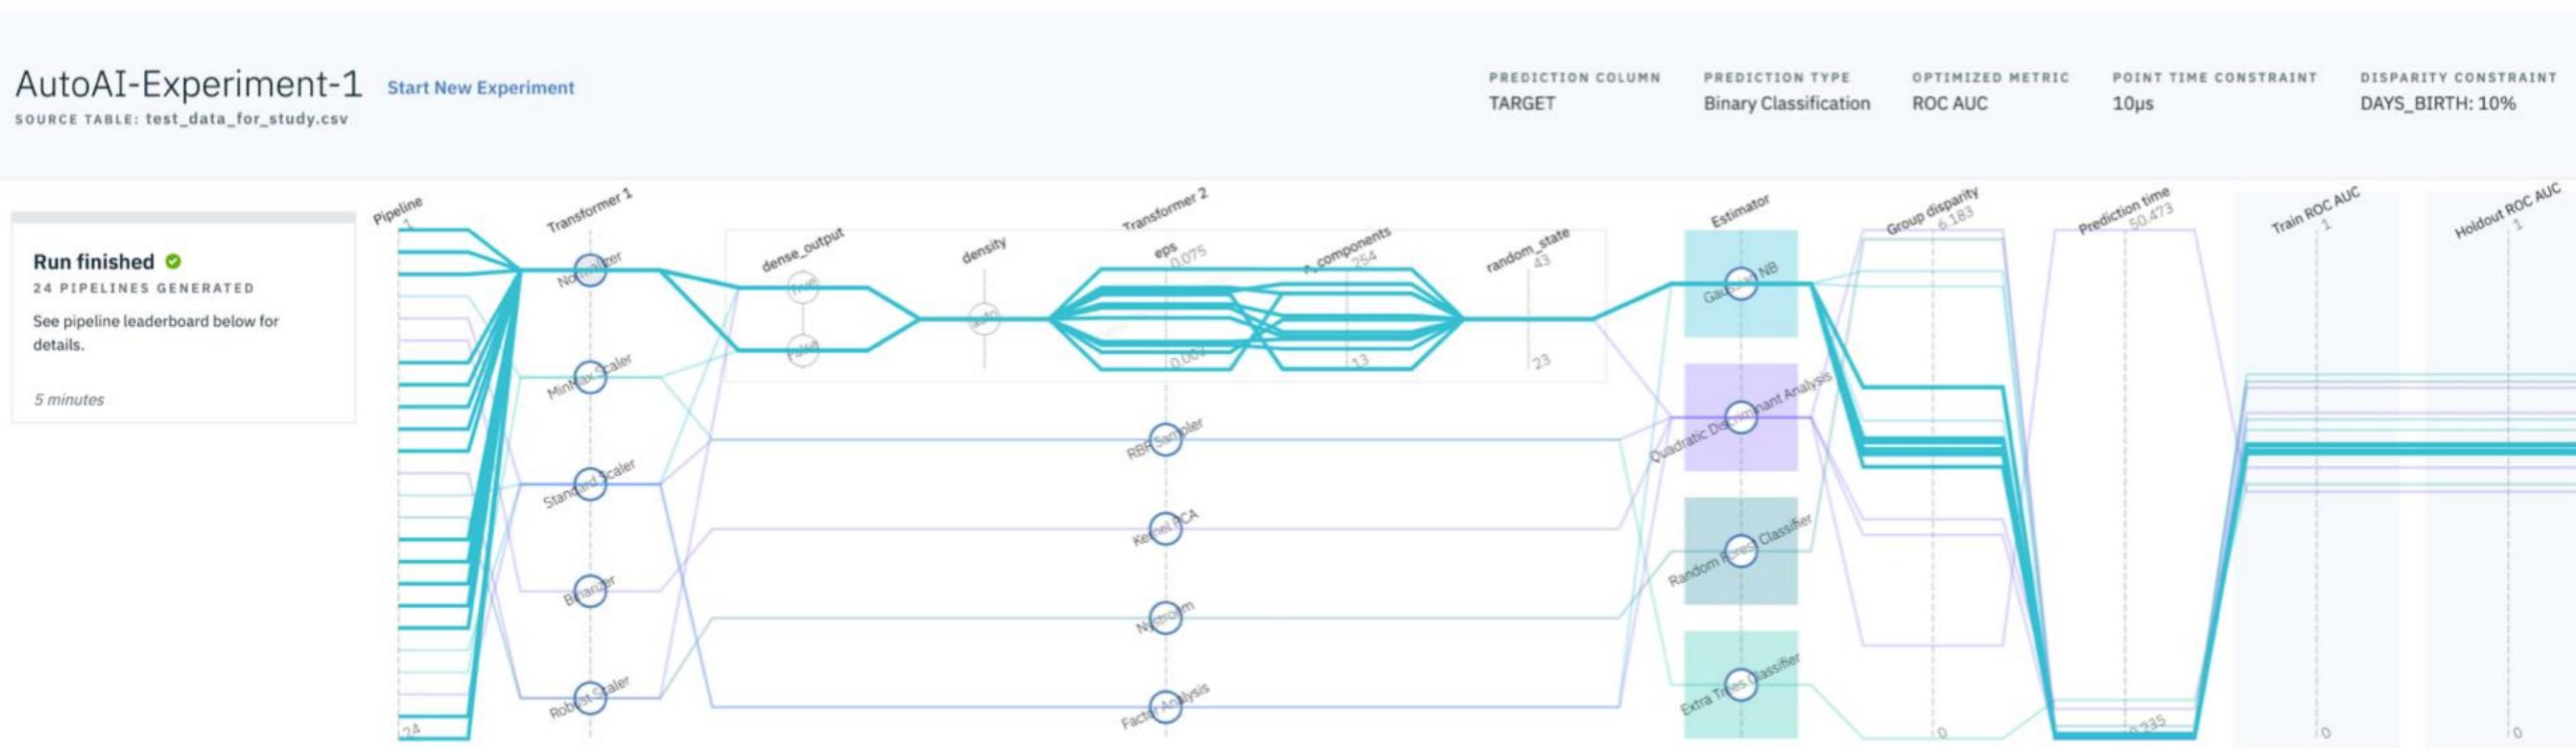
- Fixed pipeline



- Support for a single AutoML (ATM)

AutoAIvz

Debugging/Comparing AutoML Pipelines



Pipeline leaderboard

Ranking based on: [ROC AUC](#) ▾

RANK	ROC AUC	GROUP DISPARITY	PREDICTION TIME	PIPELINE INFORMATION	View details	Open notebook
> 1	0.716	55%	6μs	P4 - Gaussian nb estimator Transformers (3): MinMaxScalerTransformer > SparseRandomProjectionTransformer > Gaussian nb estimator	View details	Open notebook
> 2	0.702	61%	15μs	P14 - Random forest classifier estimator Transformers (3): RobustScalerTransformer > NystroemTransformer > Random forest classifier estimator	View details	Open notebook
> 3	0.690	25%	31μs	P12 - Quadratic discriminant analysis estimator	View details	Open notebook

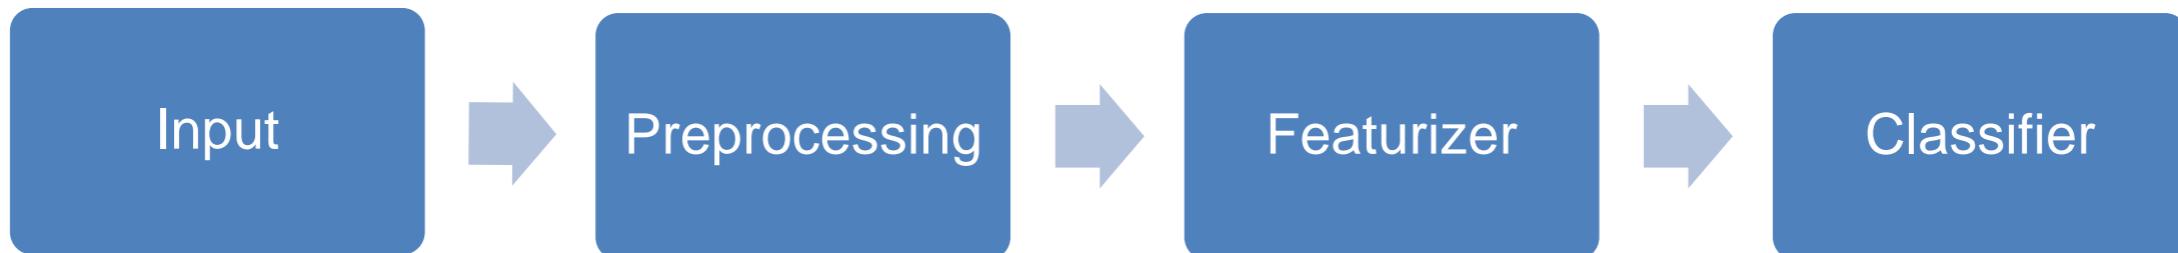
AutoAIviz

Challenges Addressed

- Requires programming background
- Black box on top of Black box (hard to explain)
- No user feedback (how to adapt the models based on human knowledge)

Limitations:

- Fixed pipeline

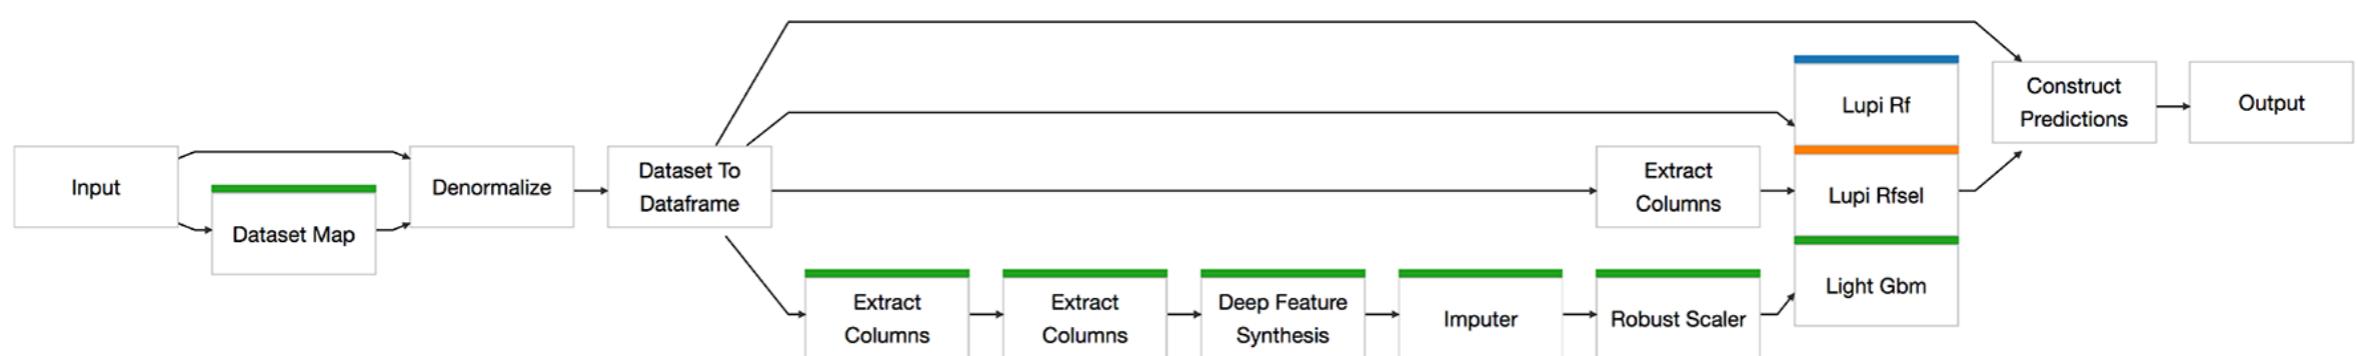
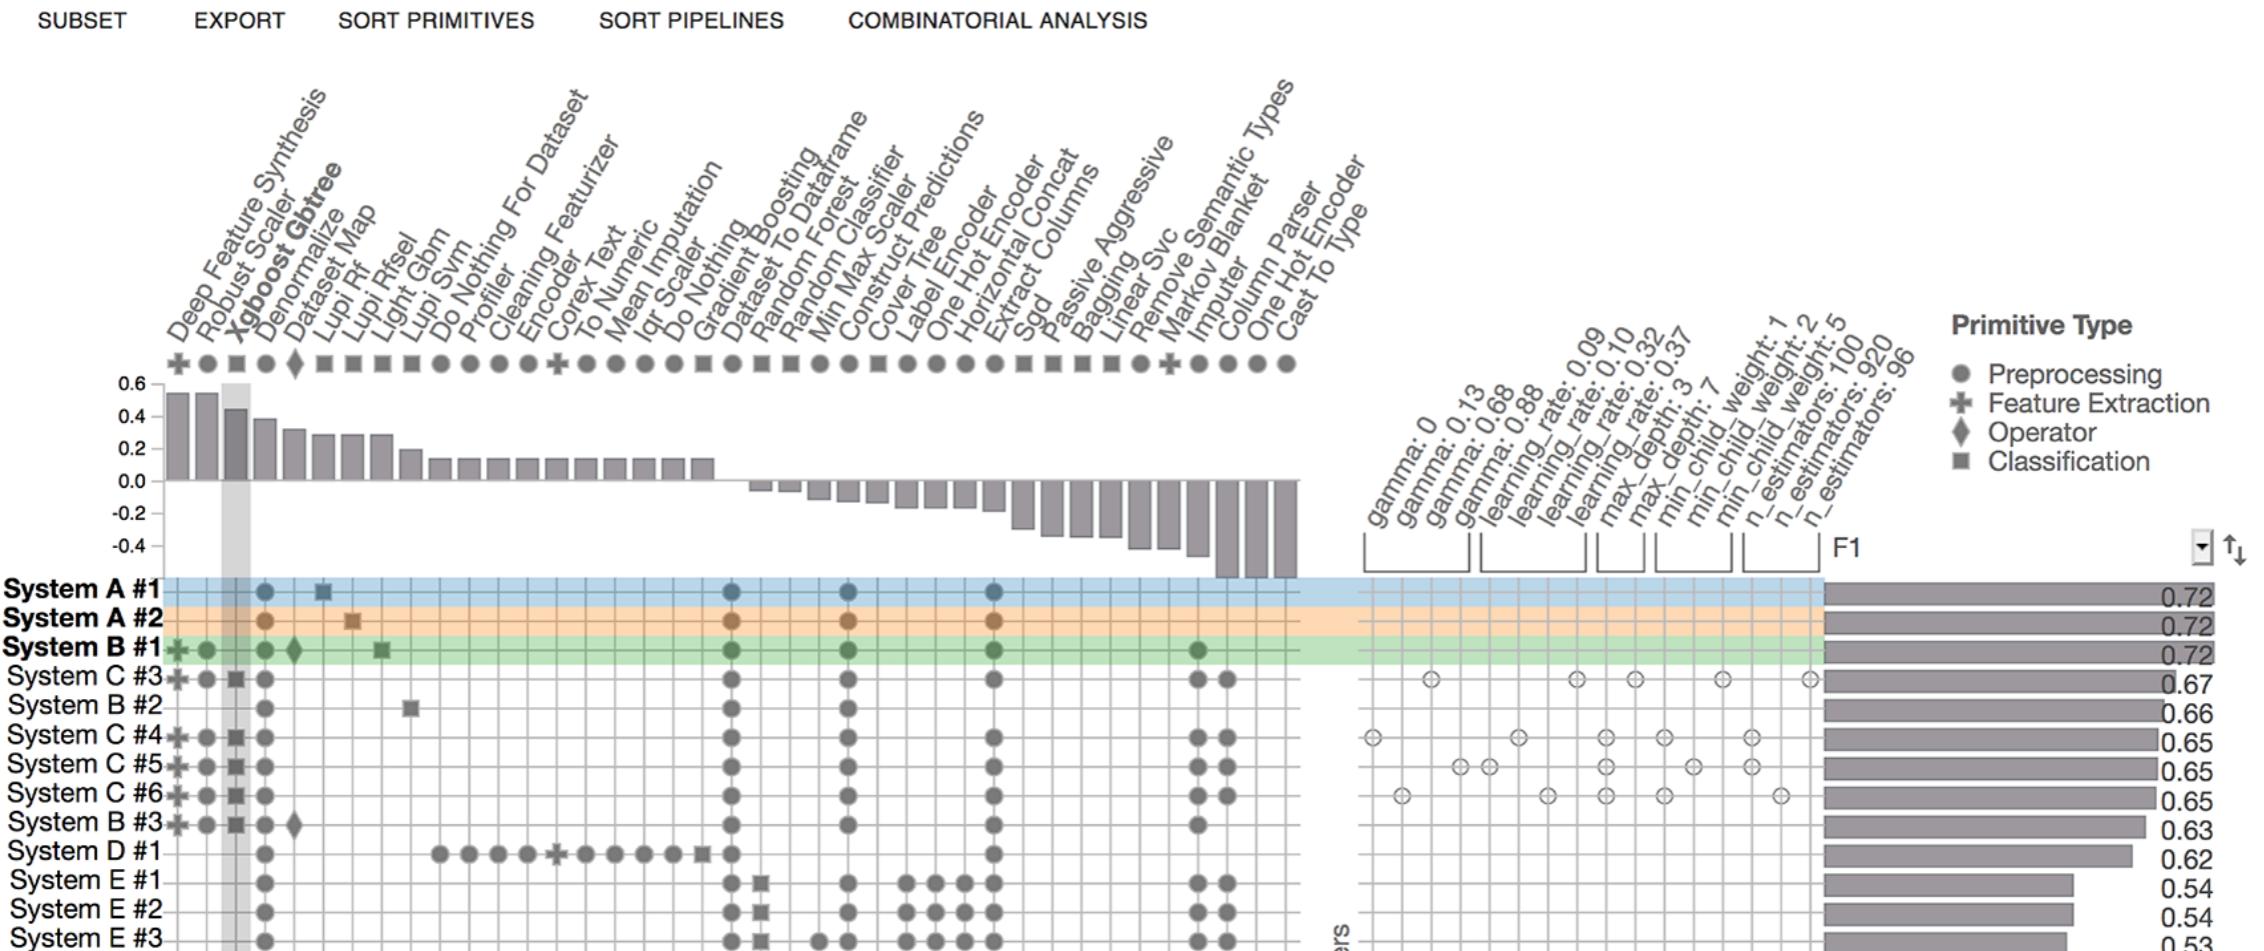


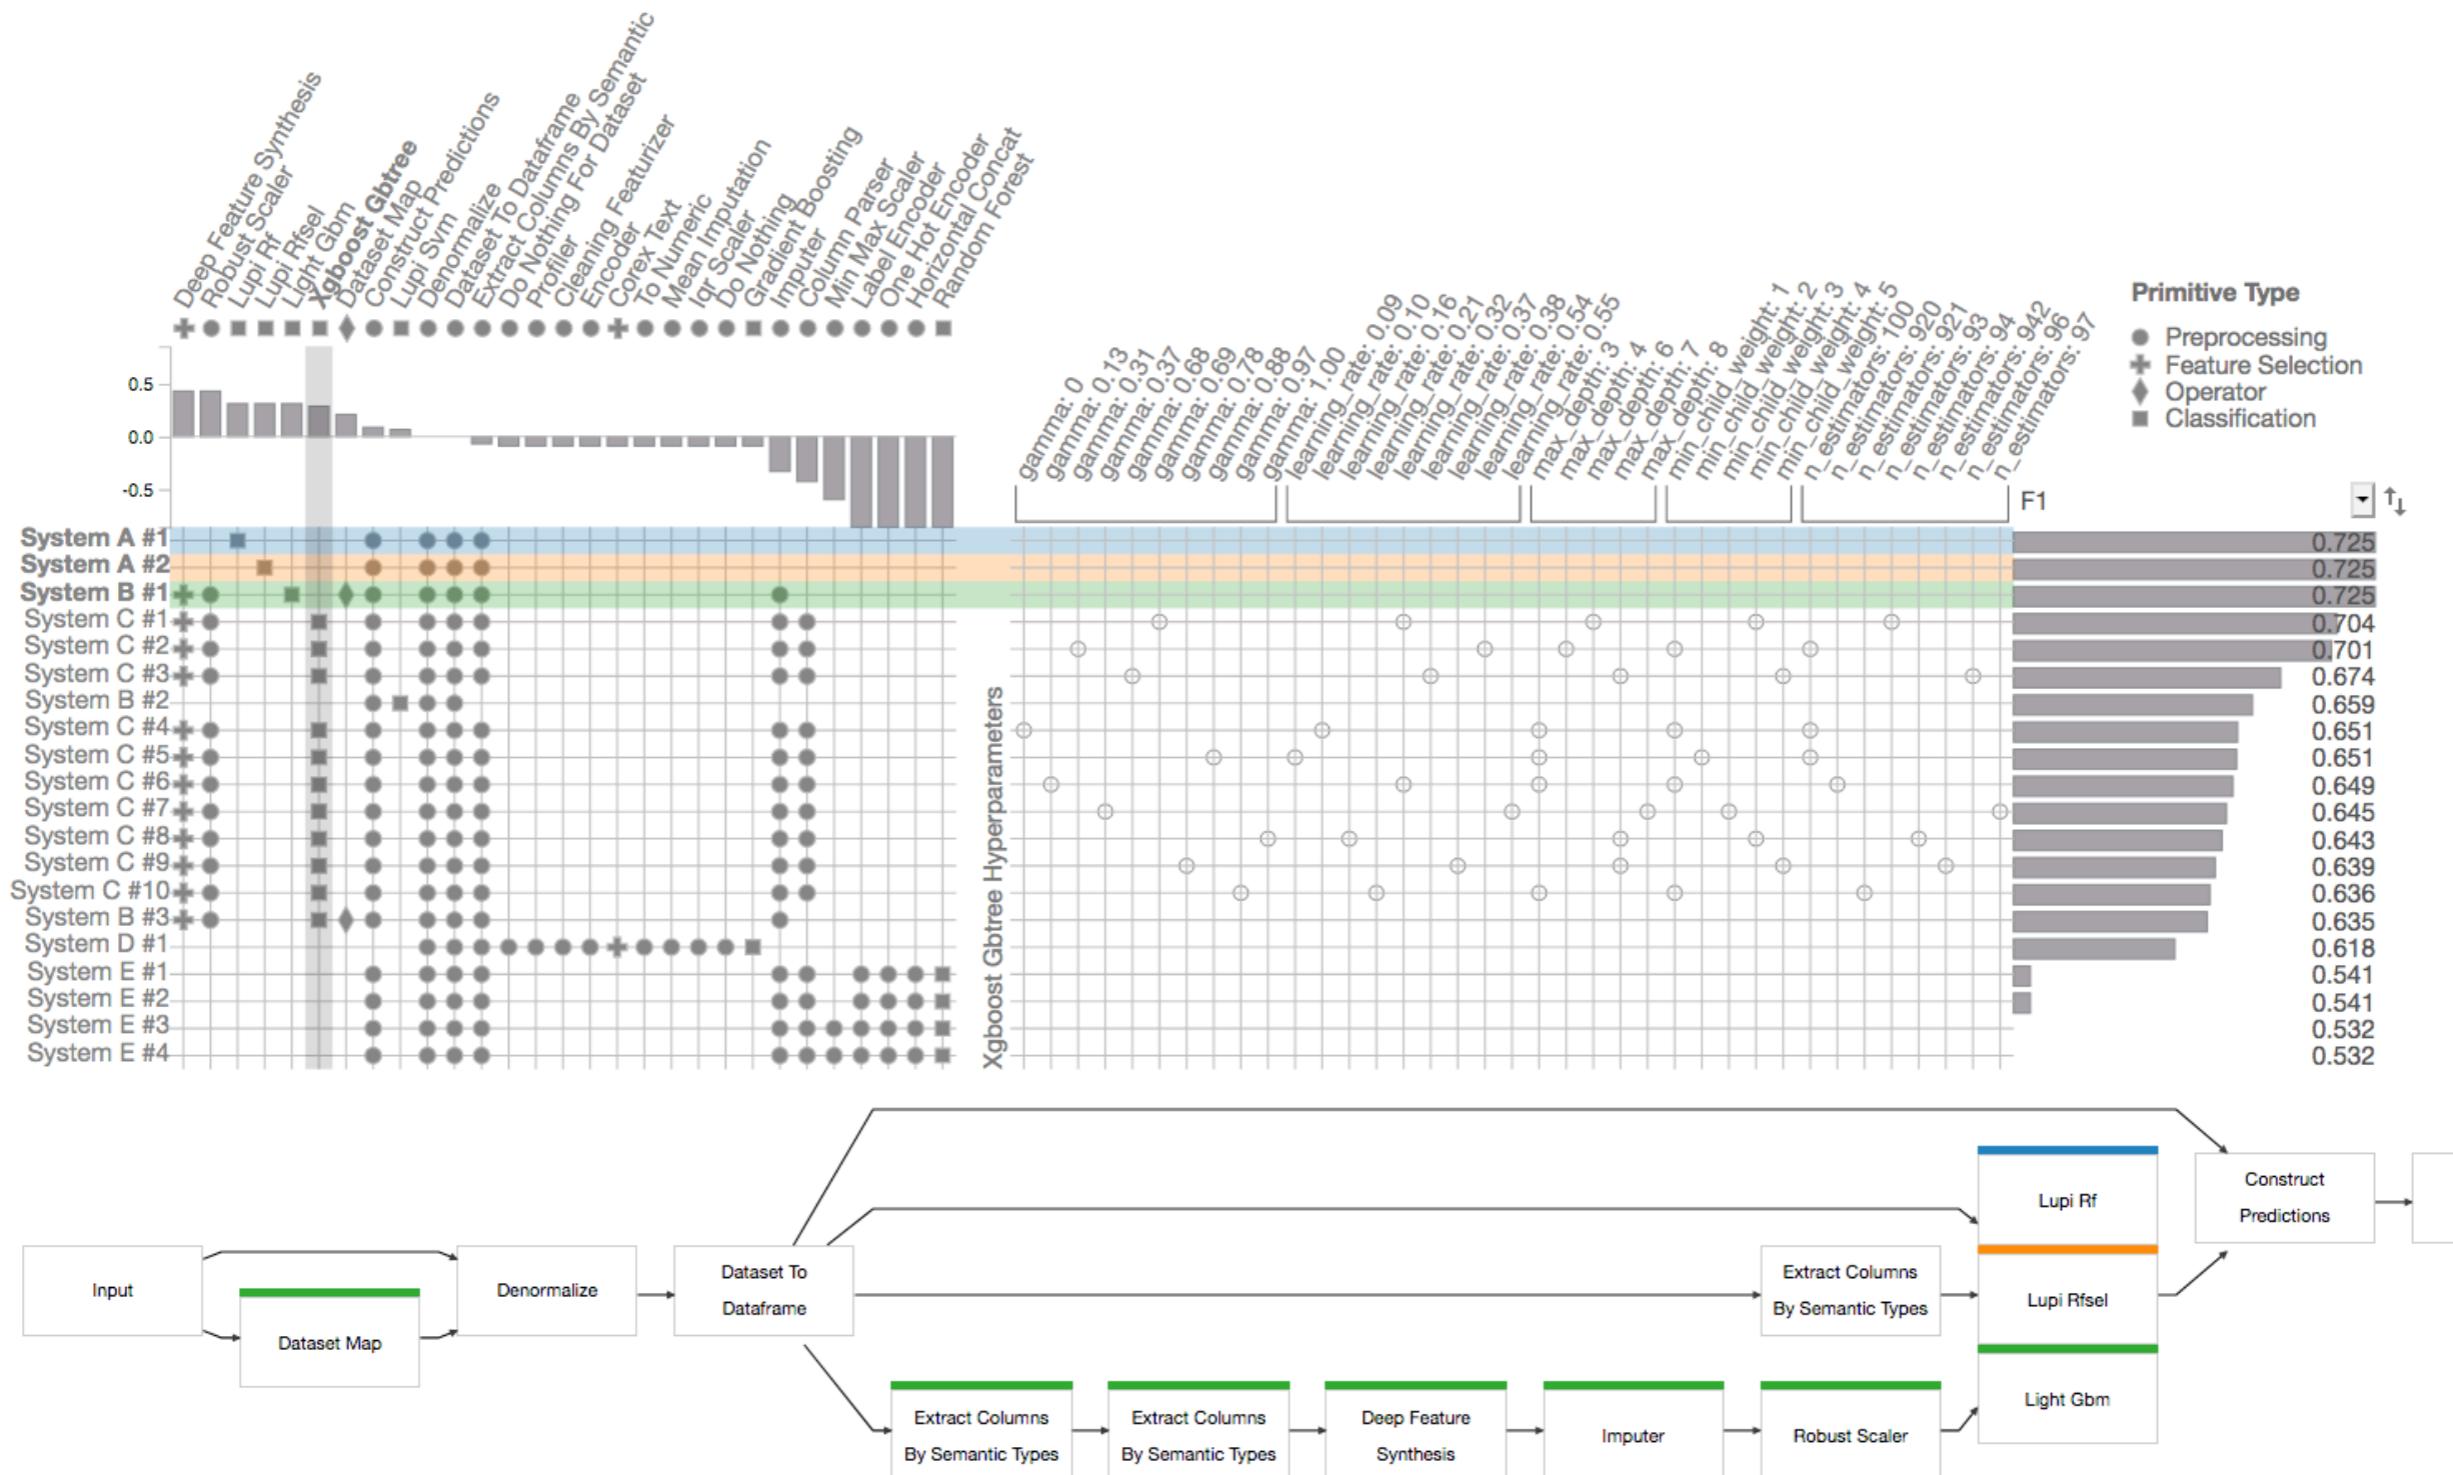
- Support for a single AutoML (ATM)

PipelineProfiler

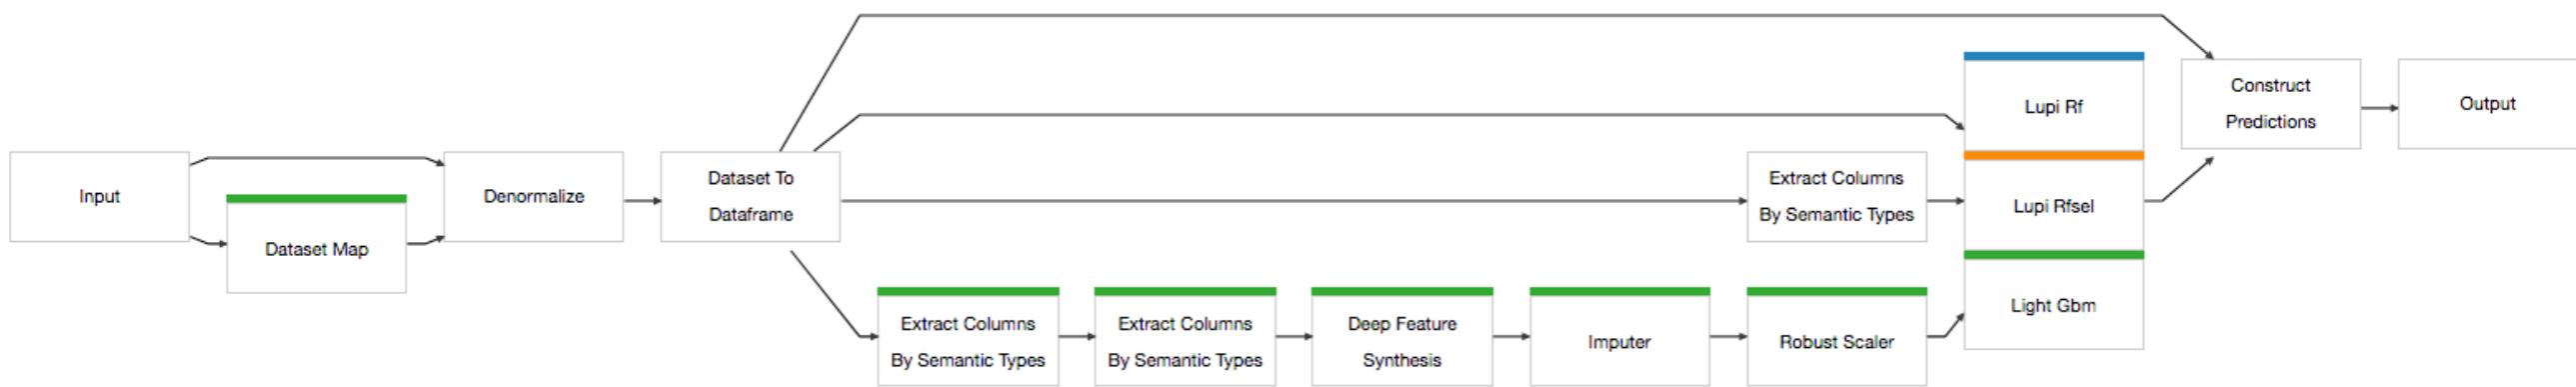
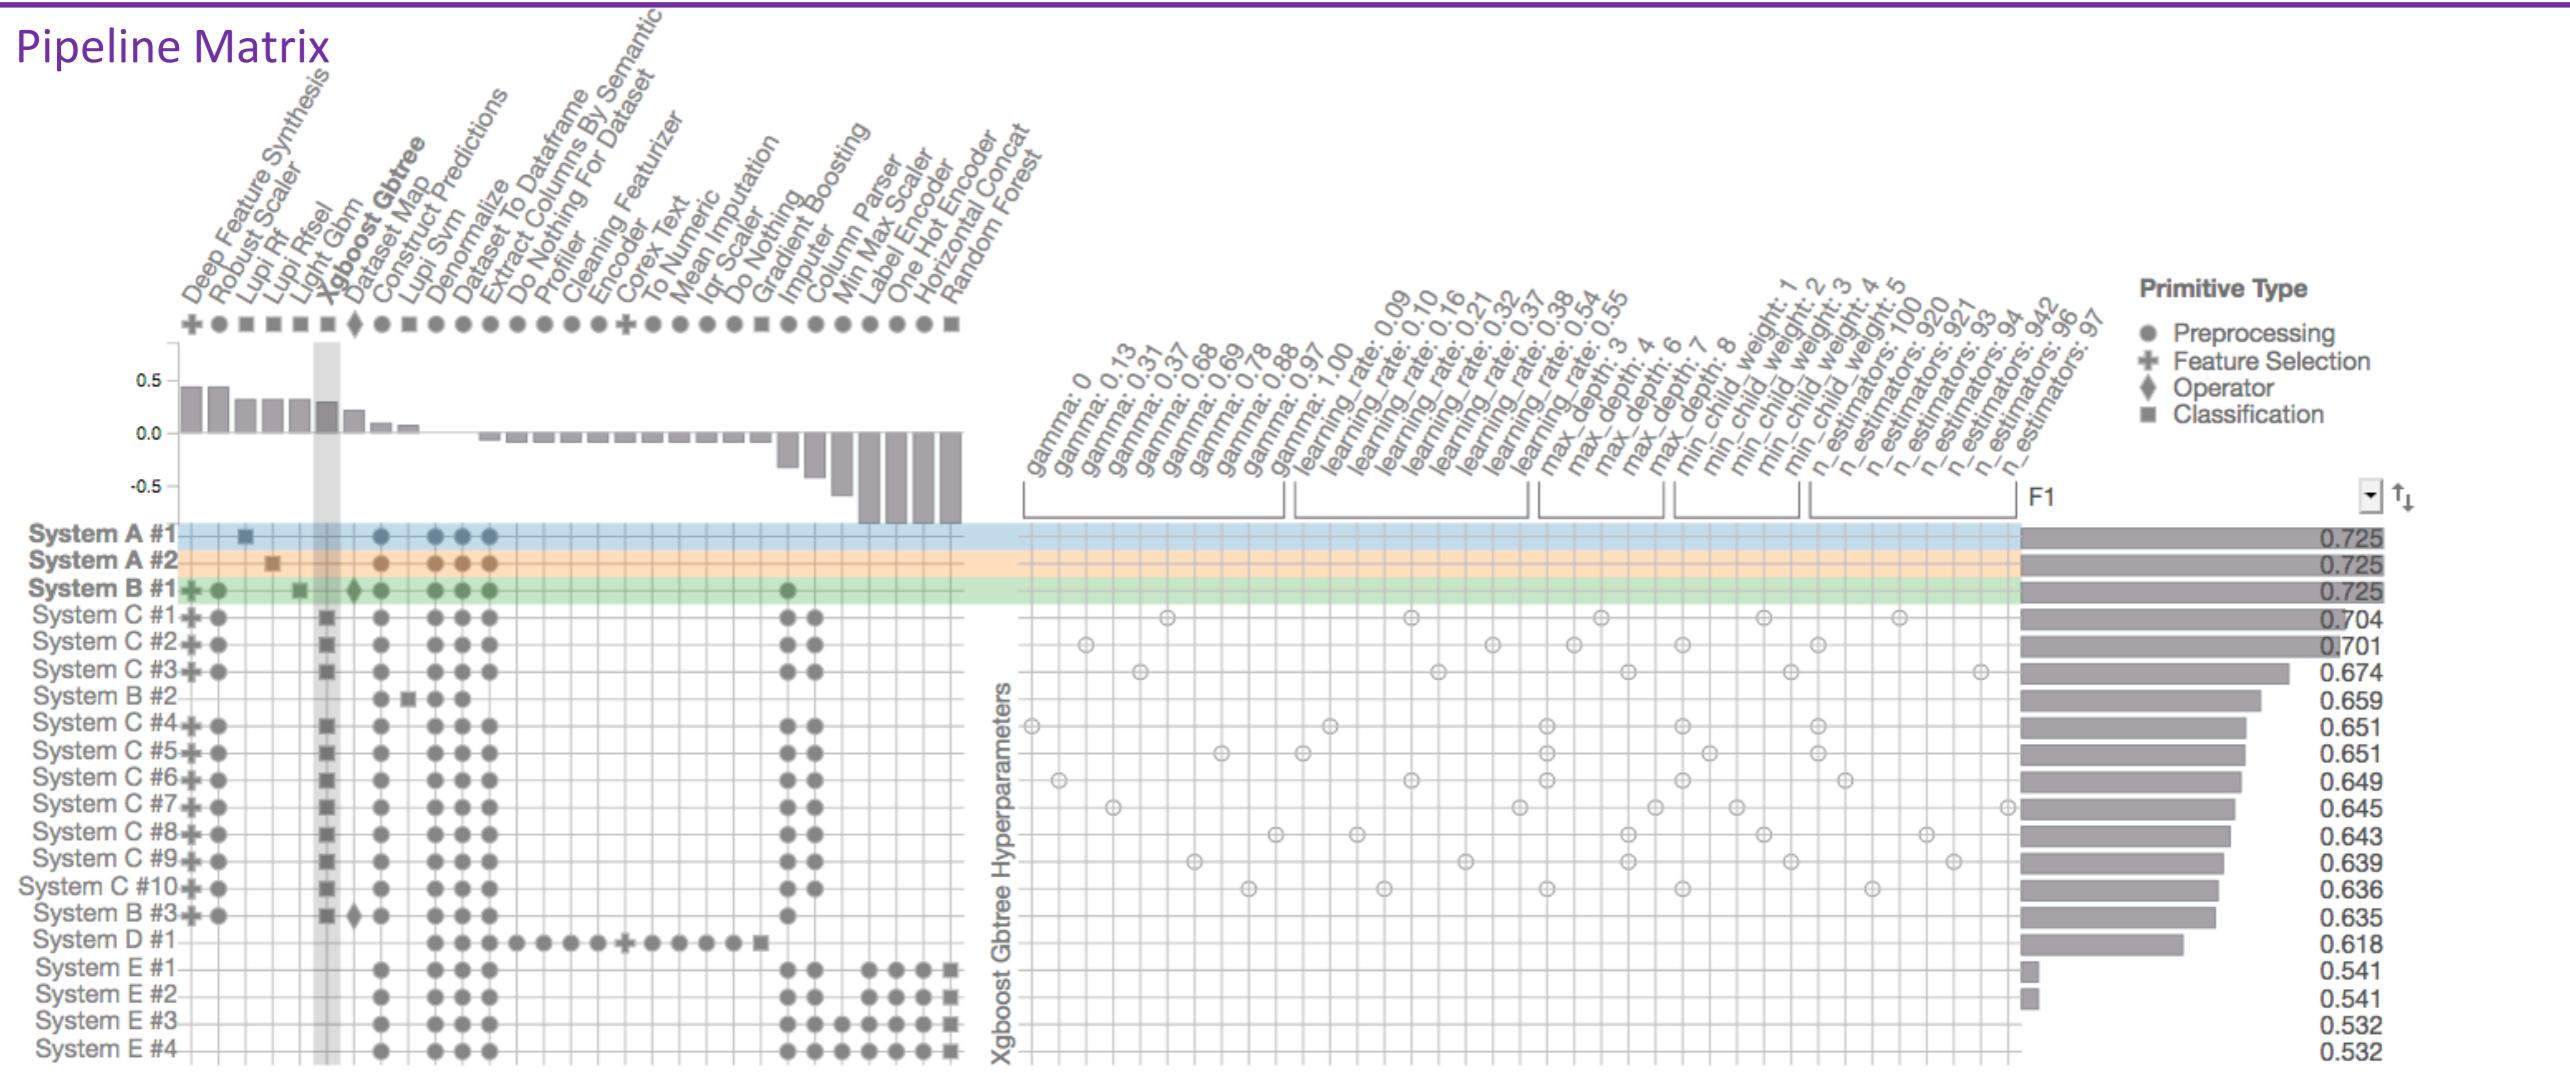
Debugging/comparing AutoML pipelines produced by multiple systems

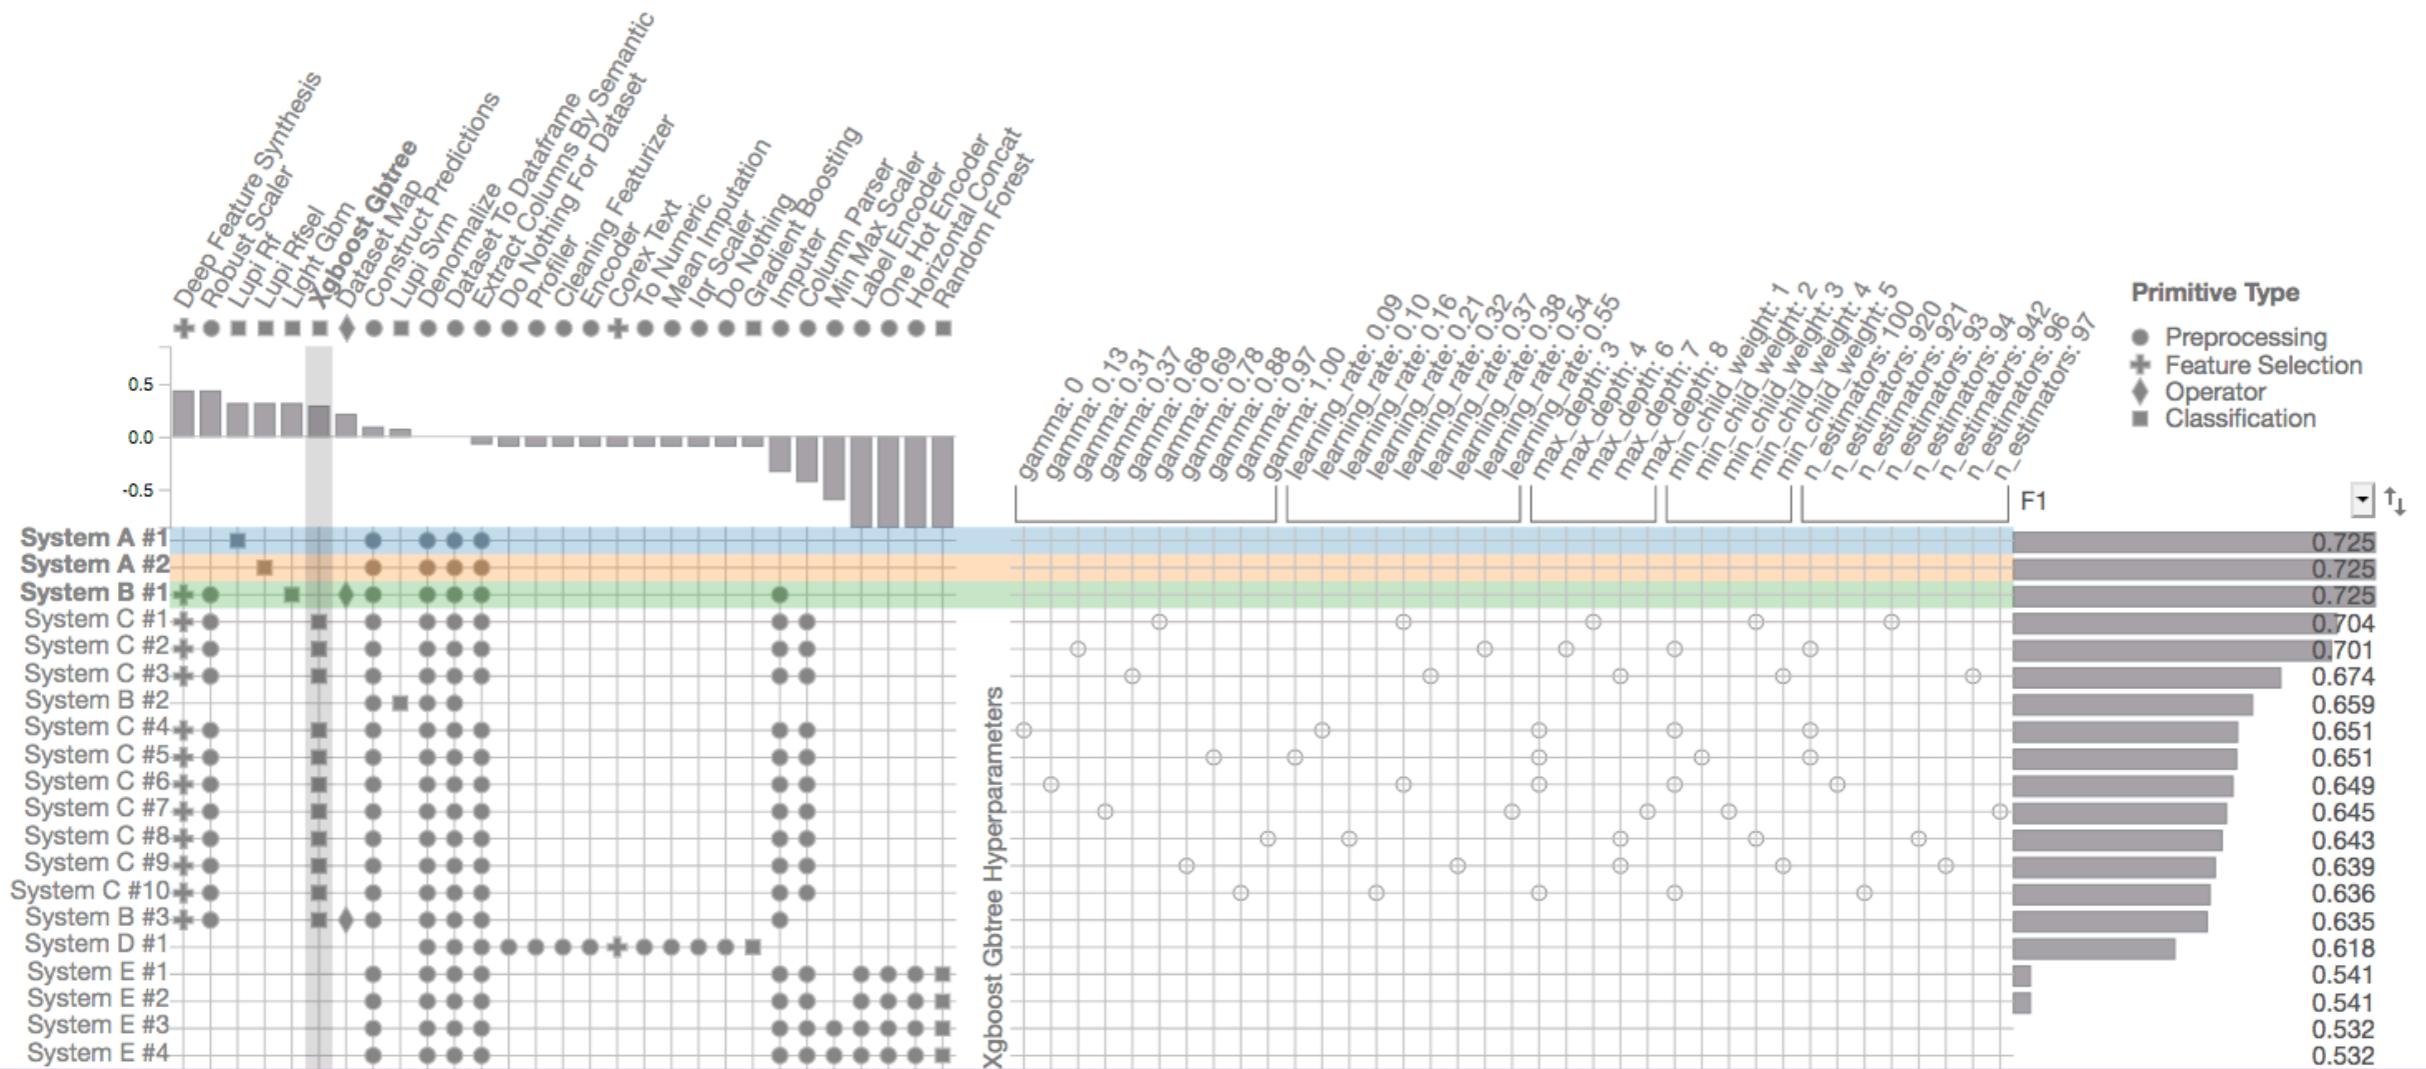
```
PipelineProfiler.plot_pipeline_matrix(heartstatlog)
```



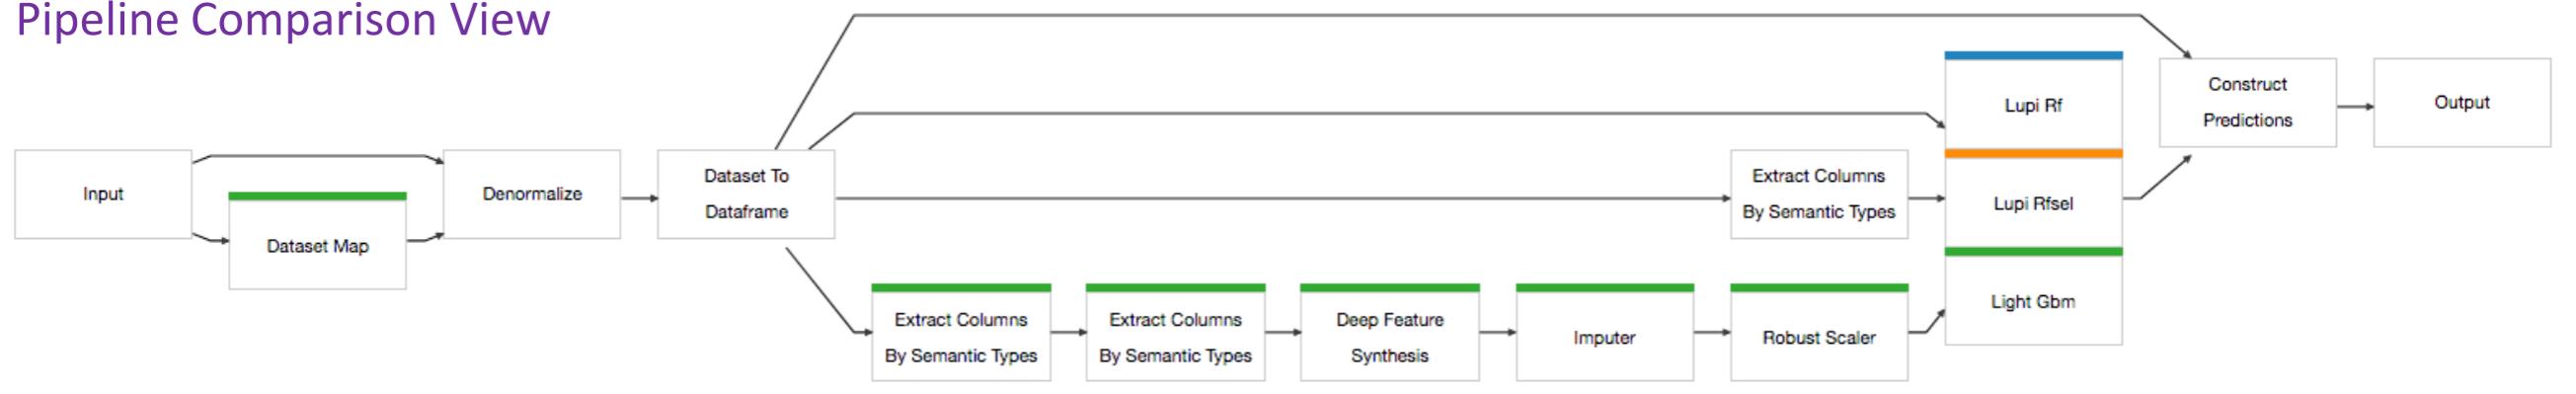


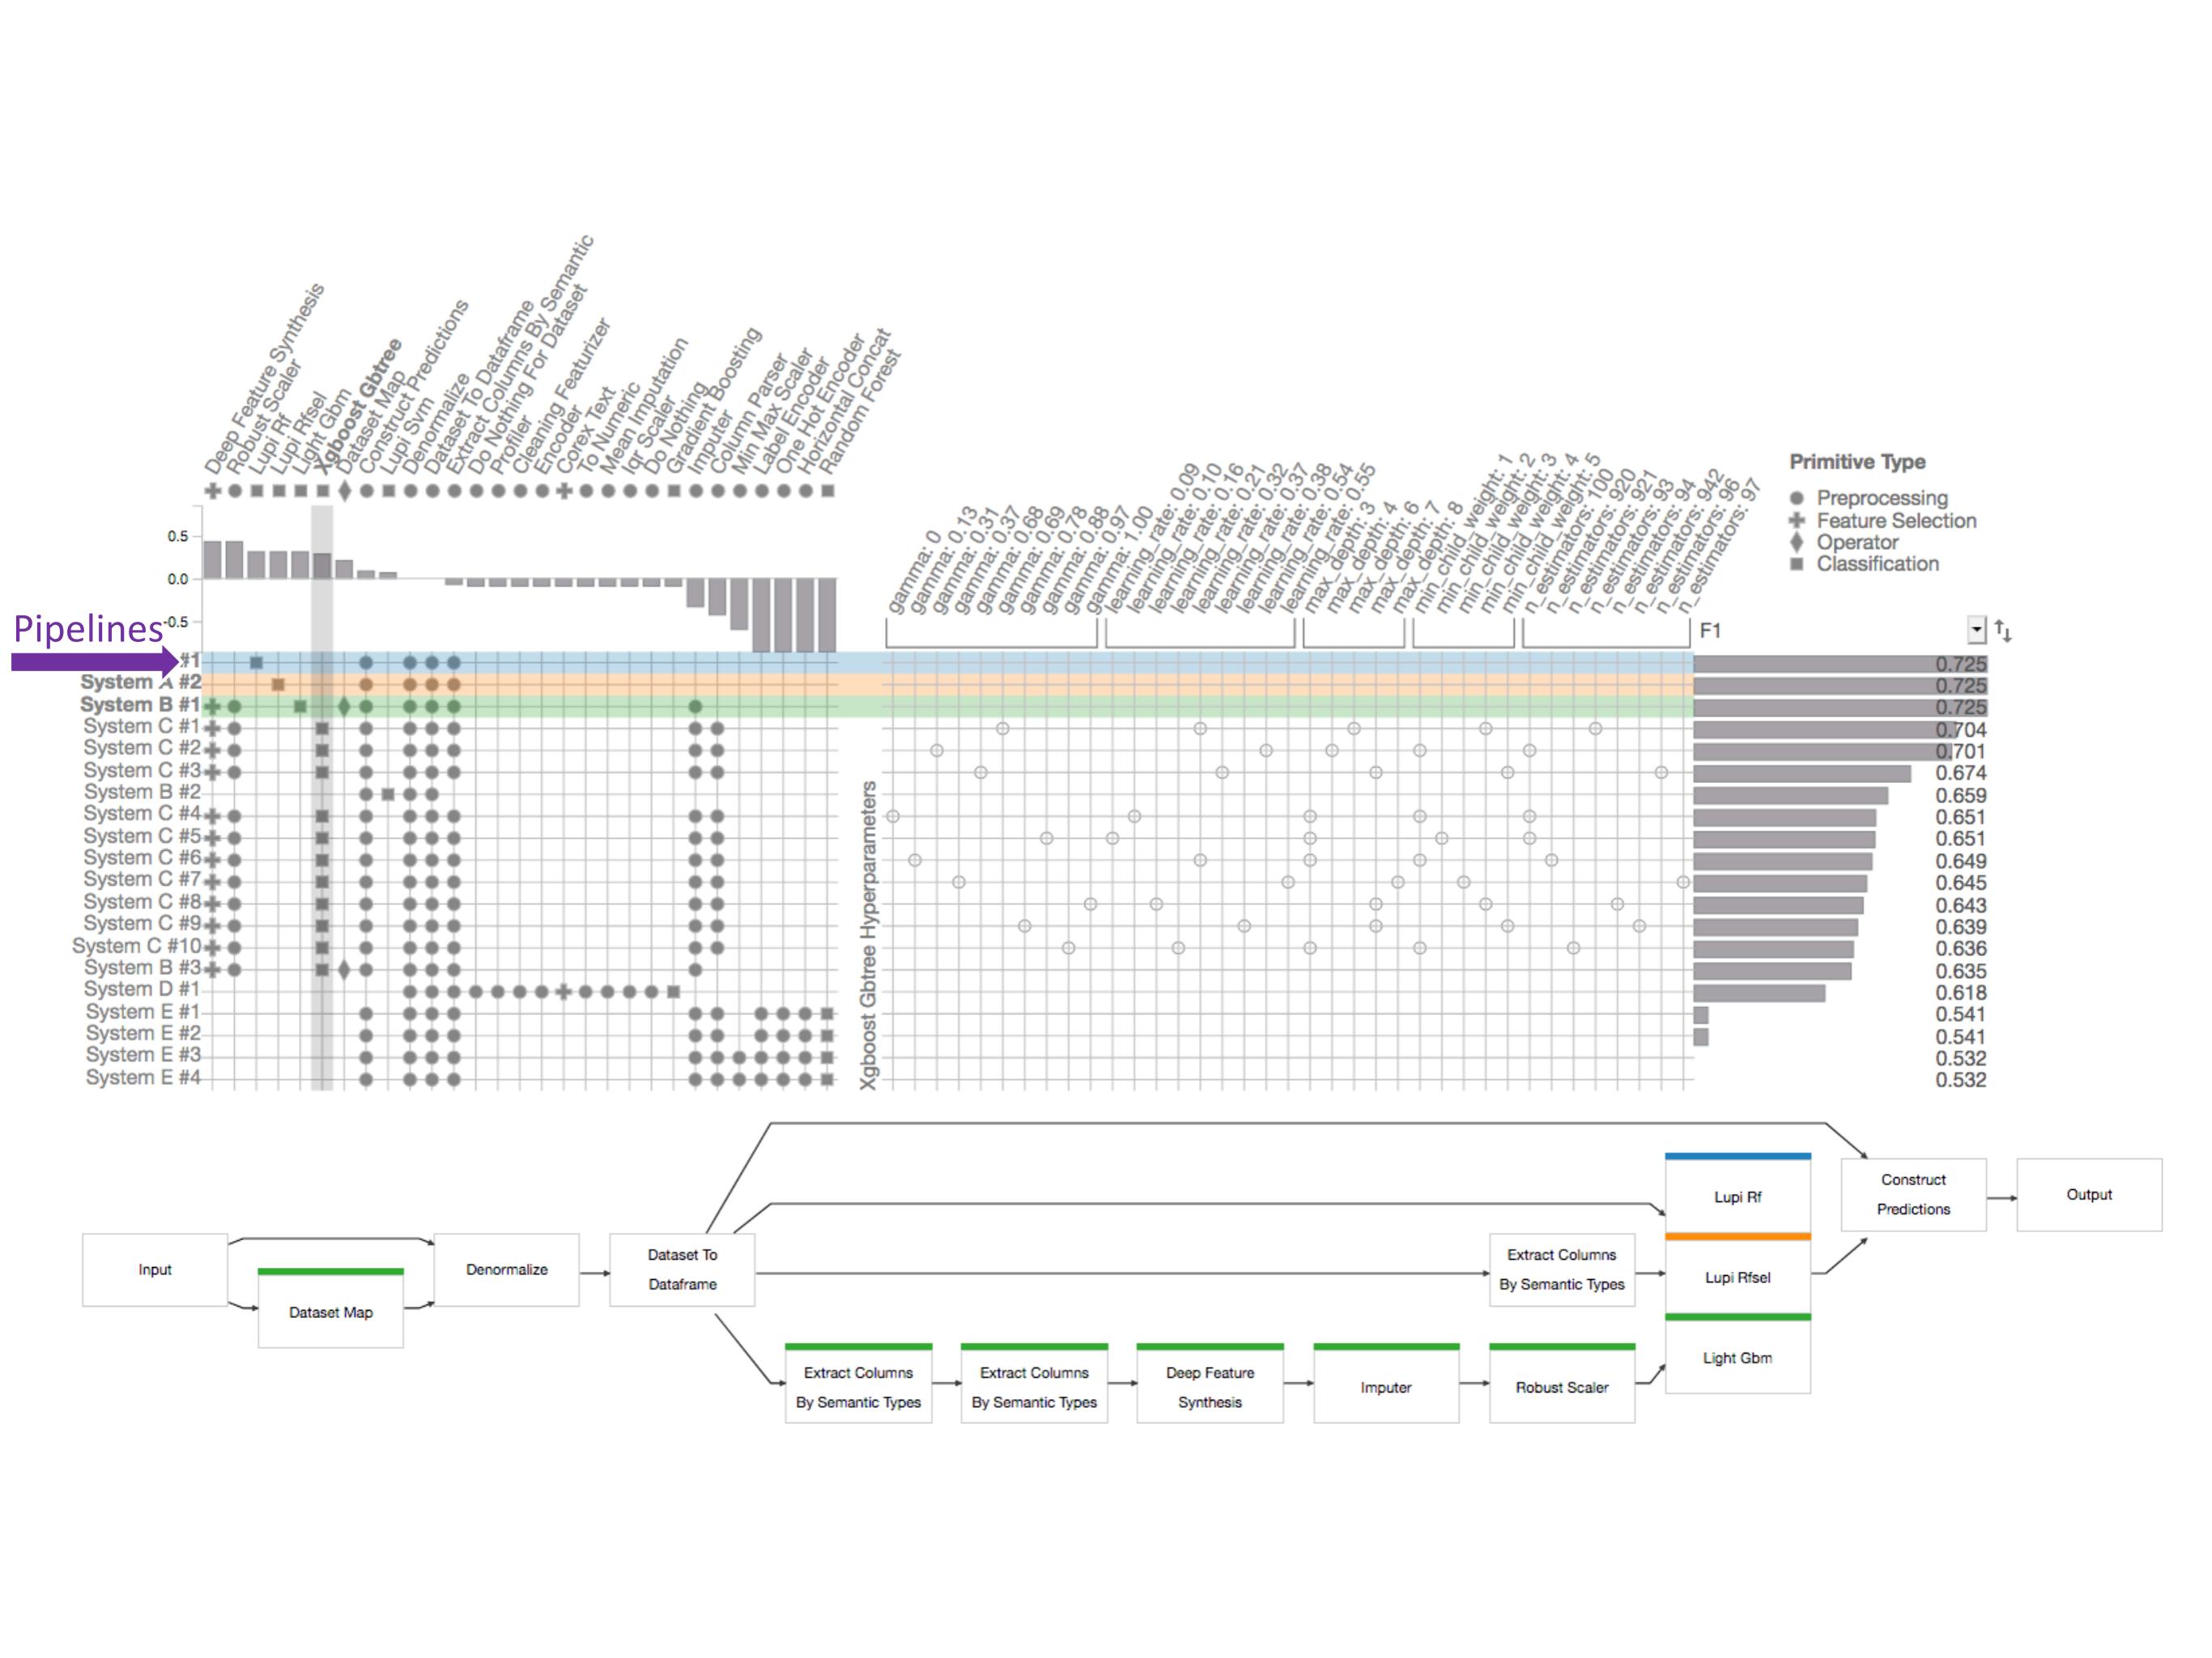
Pipeline Matrix

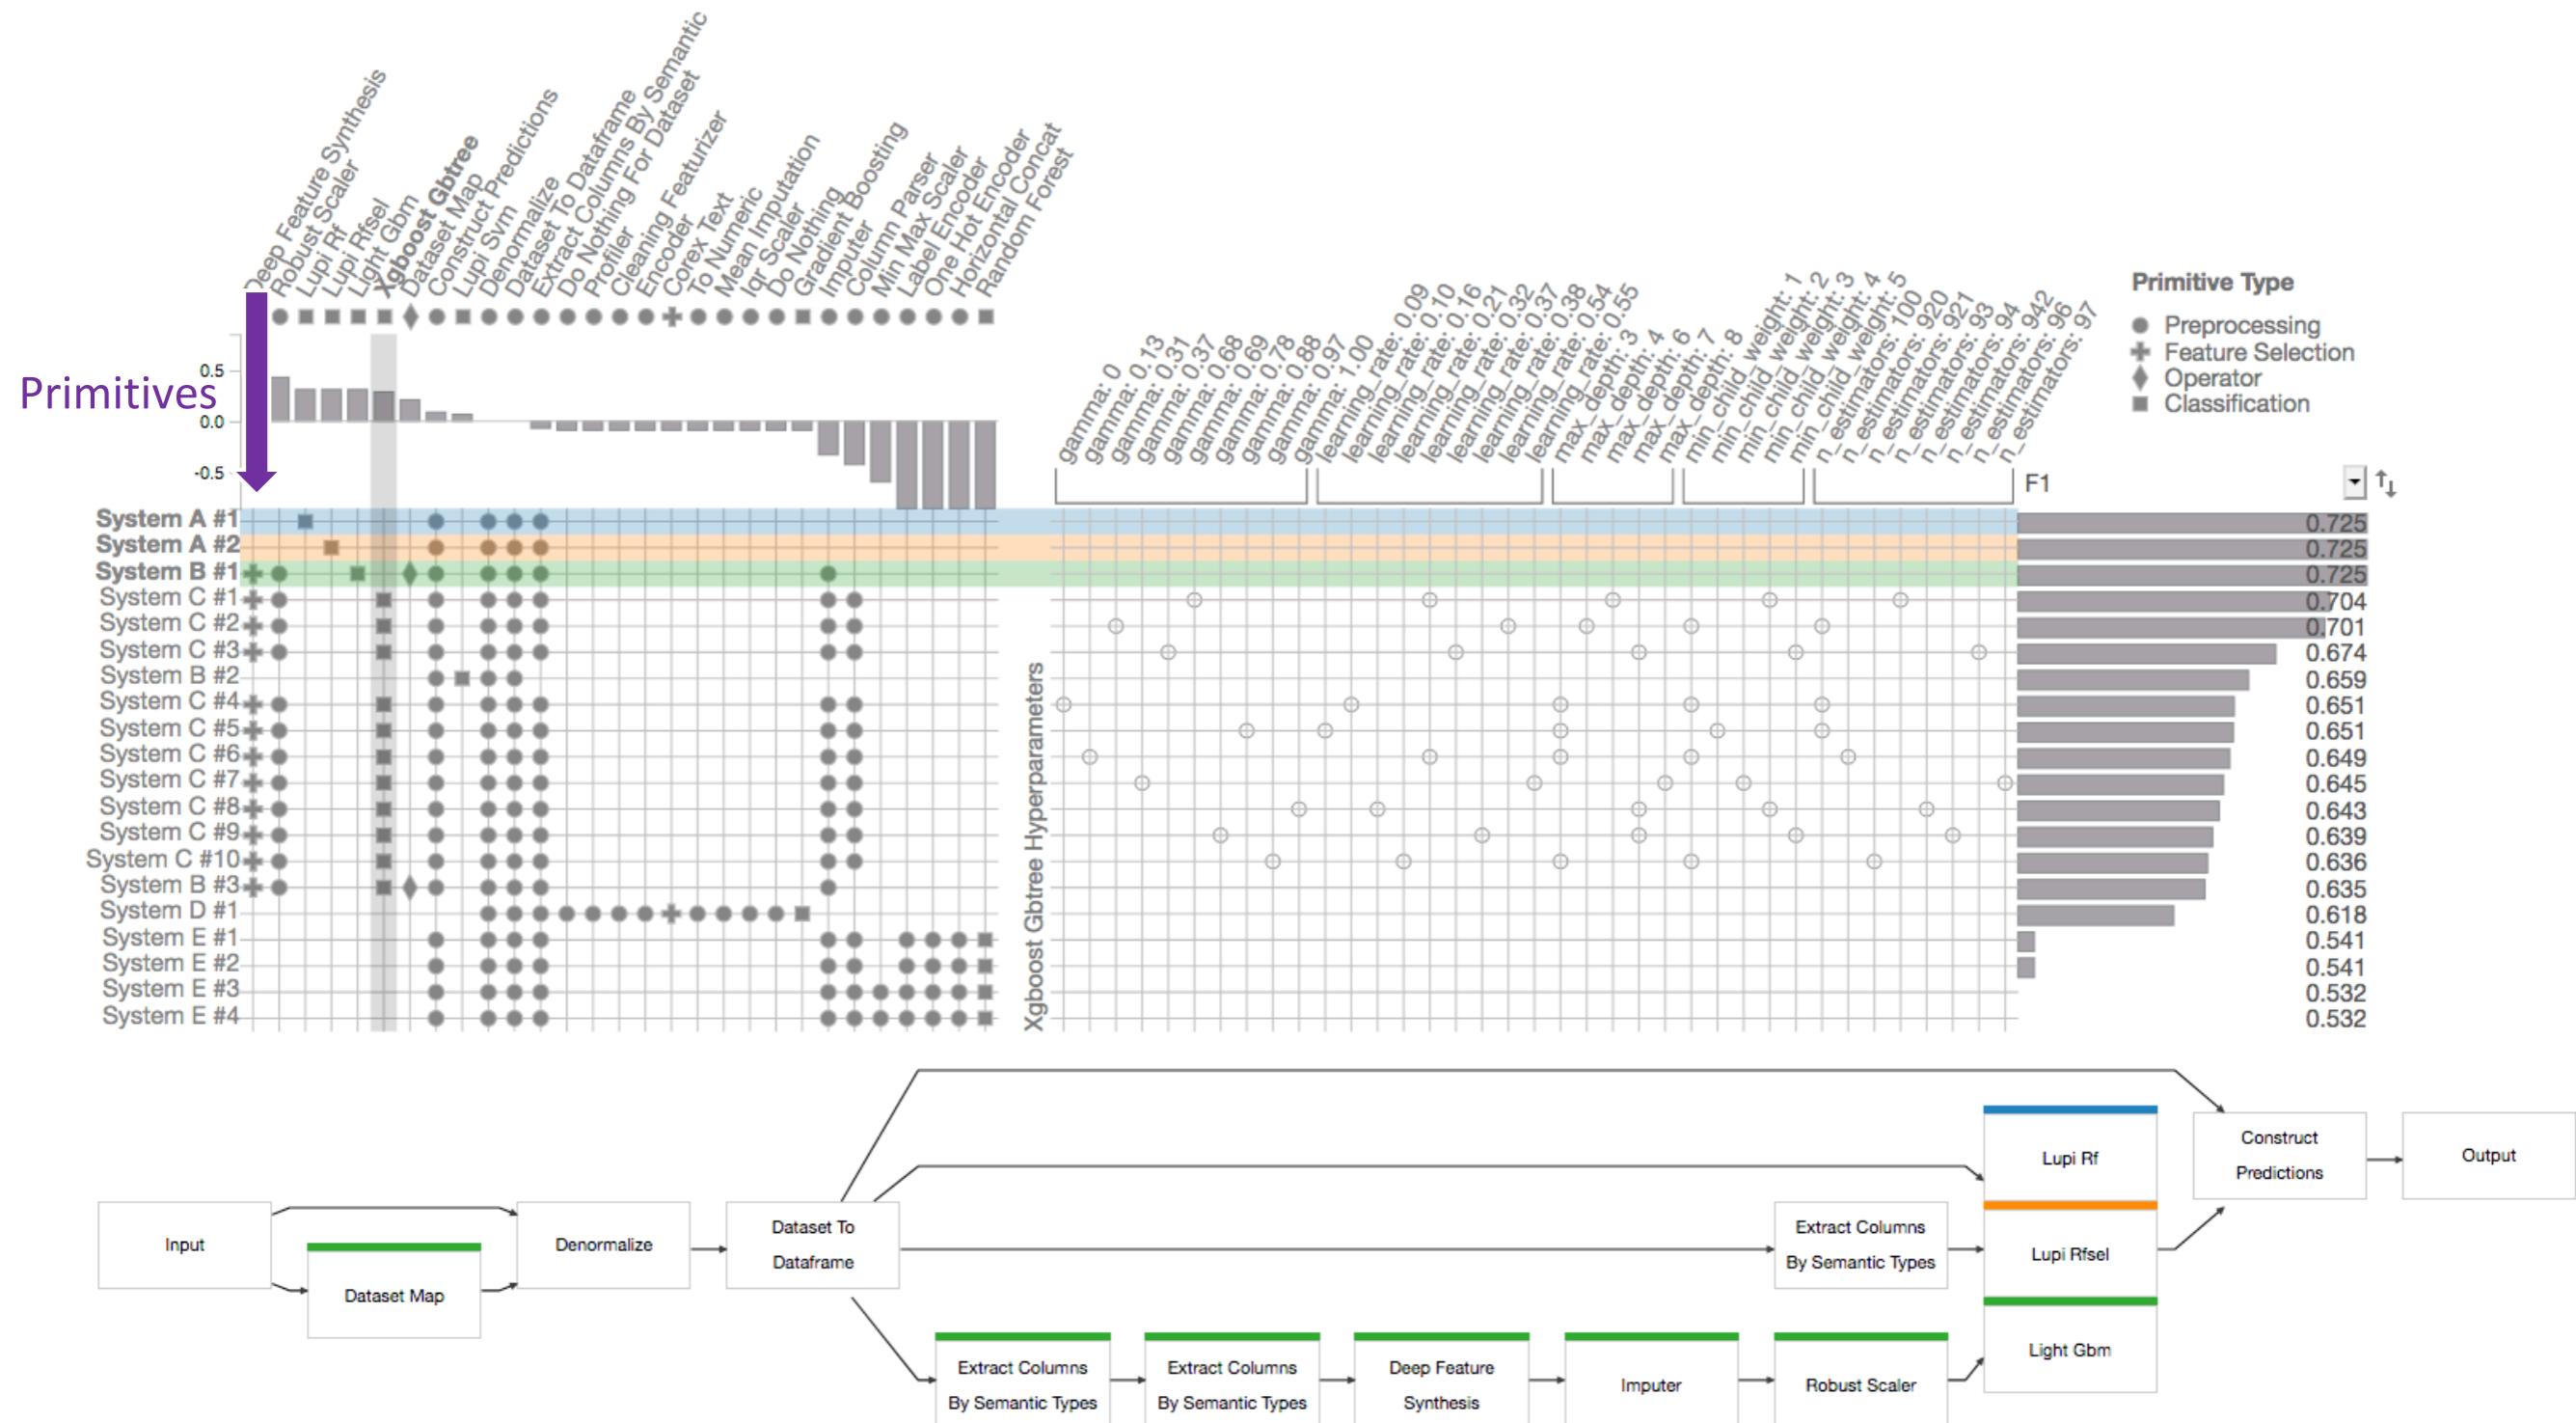


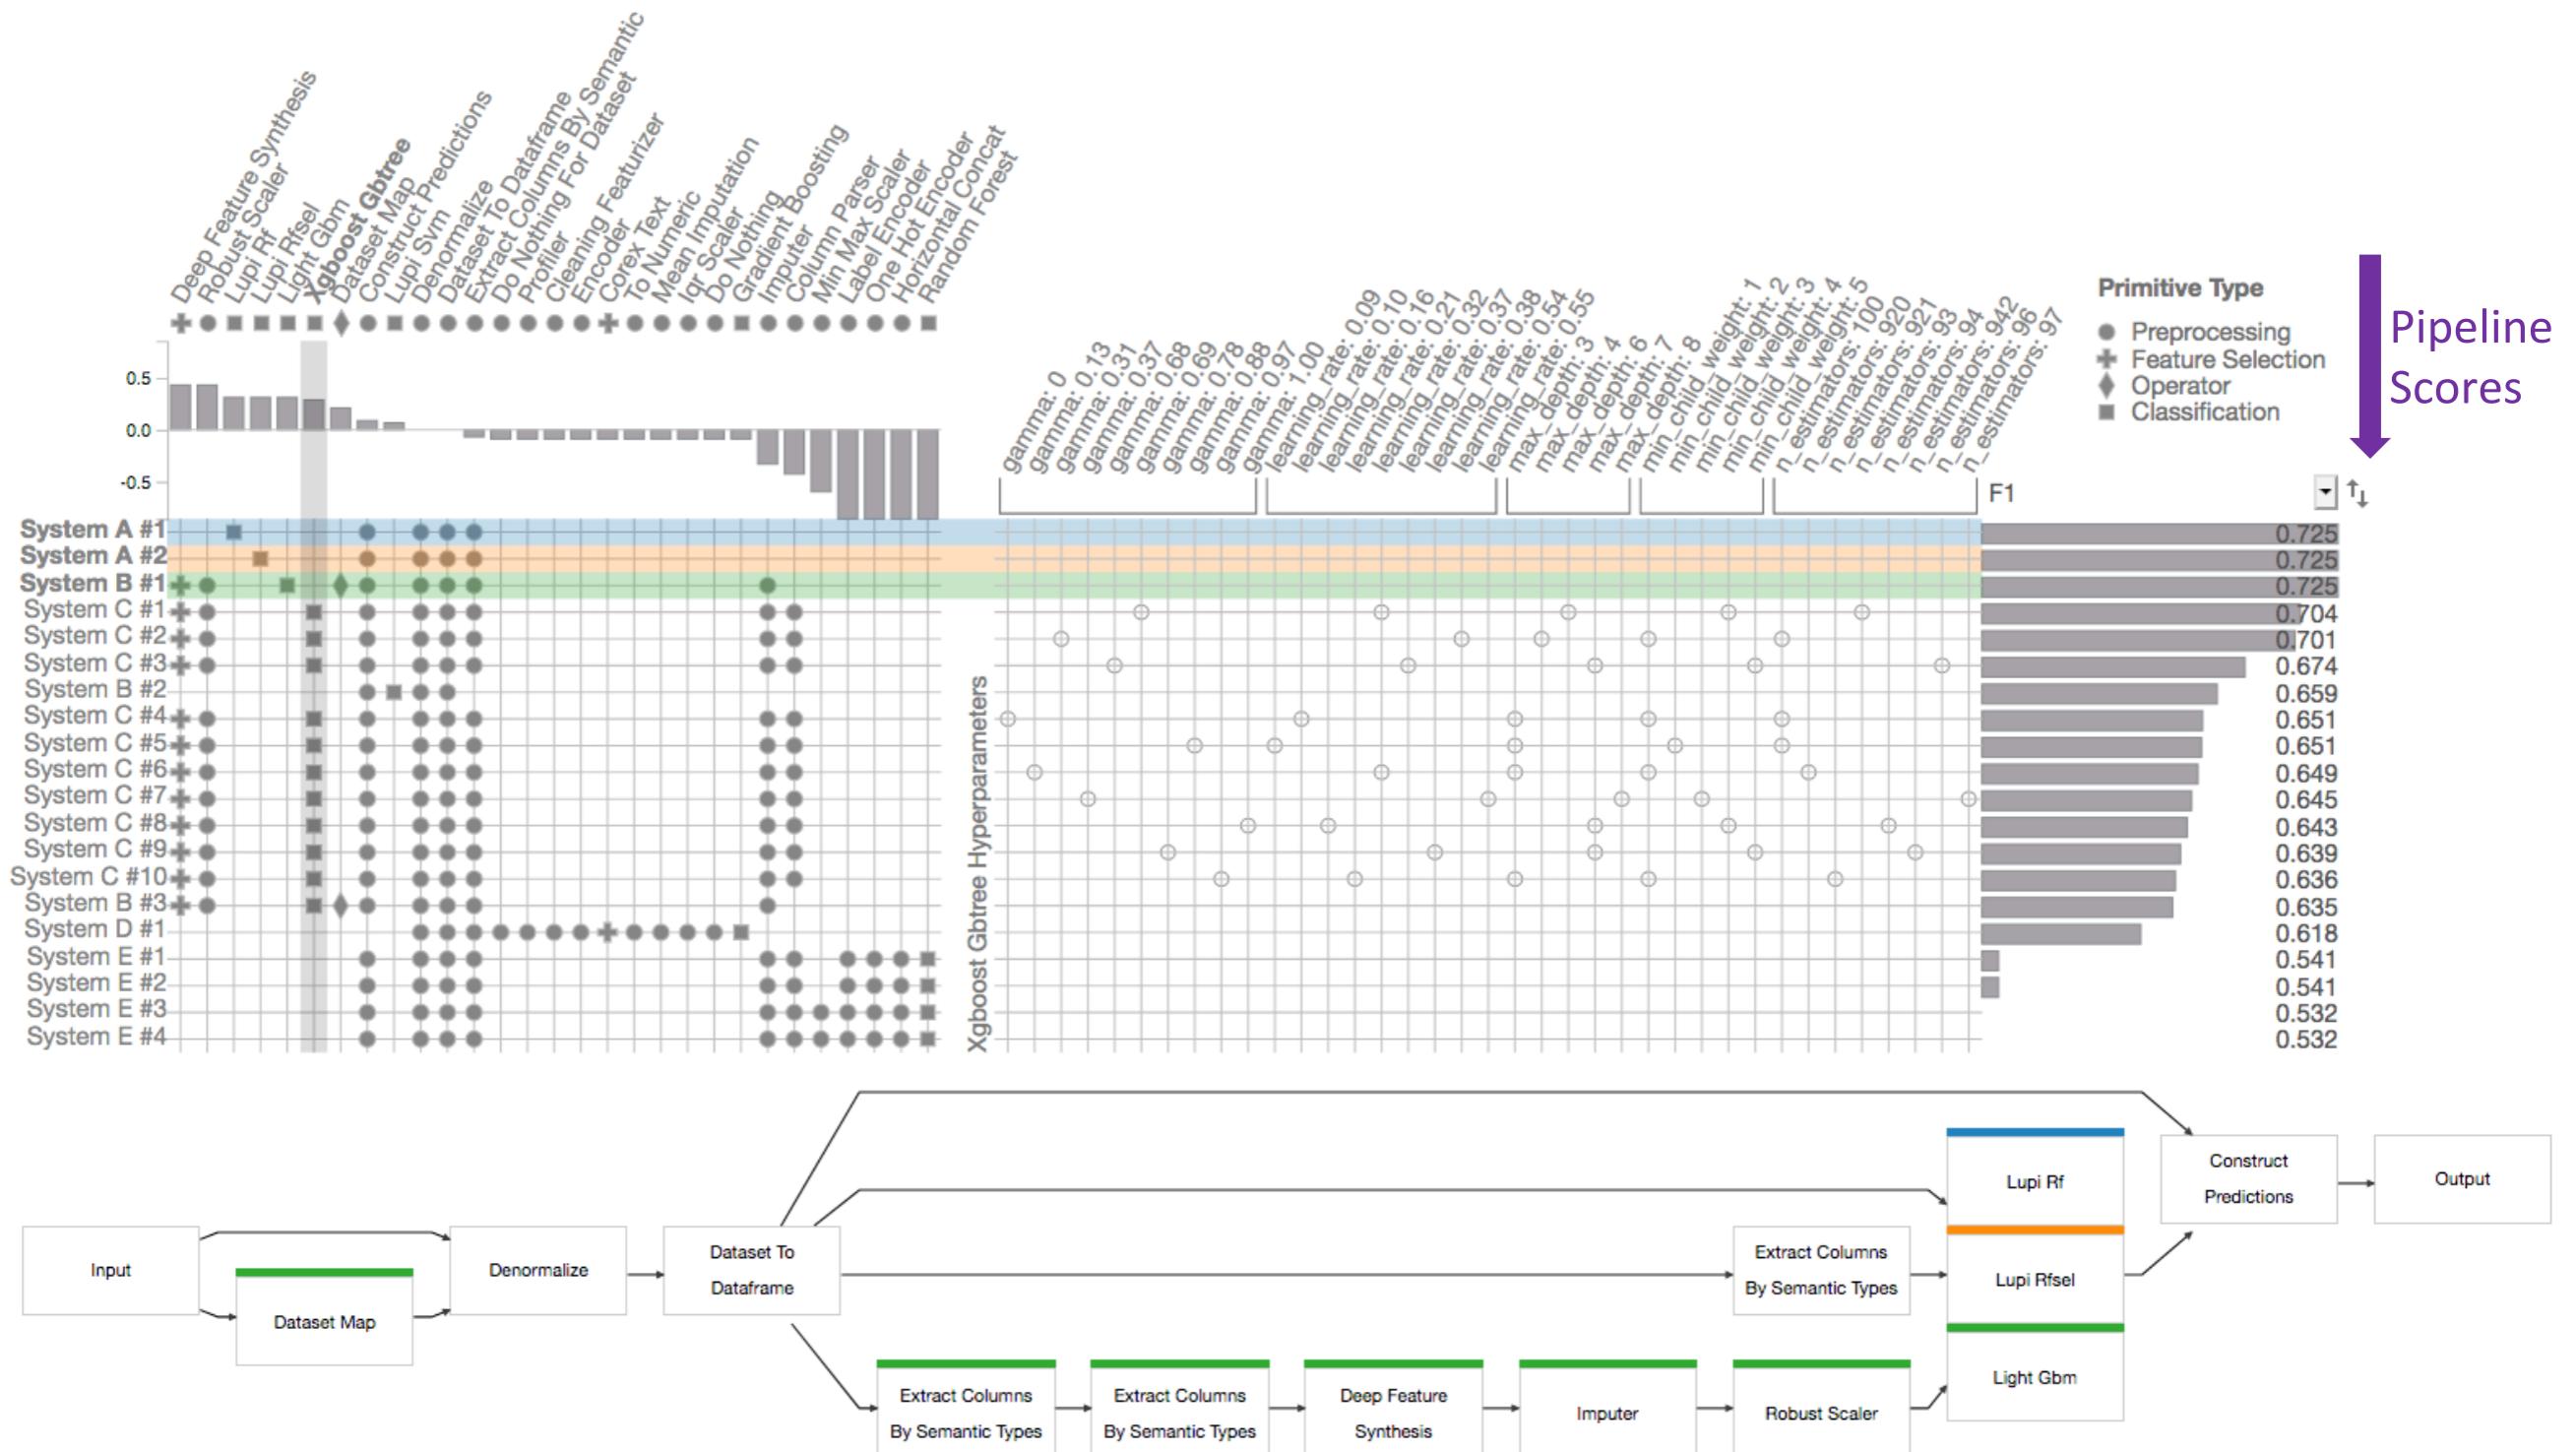


Pipeline Comparison View

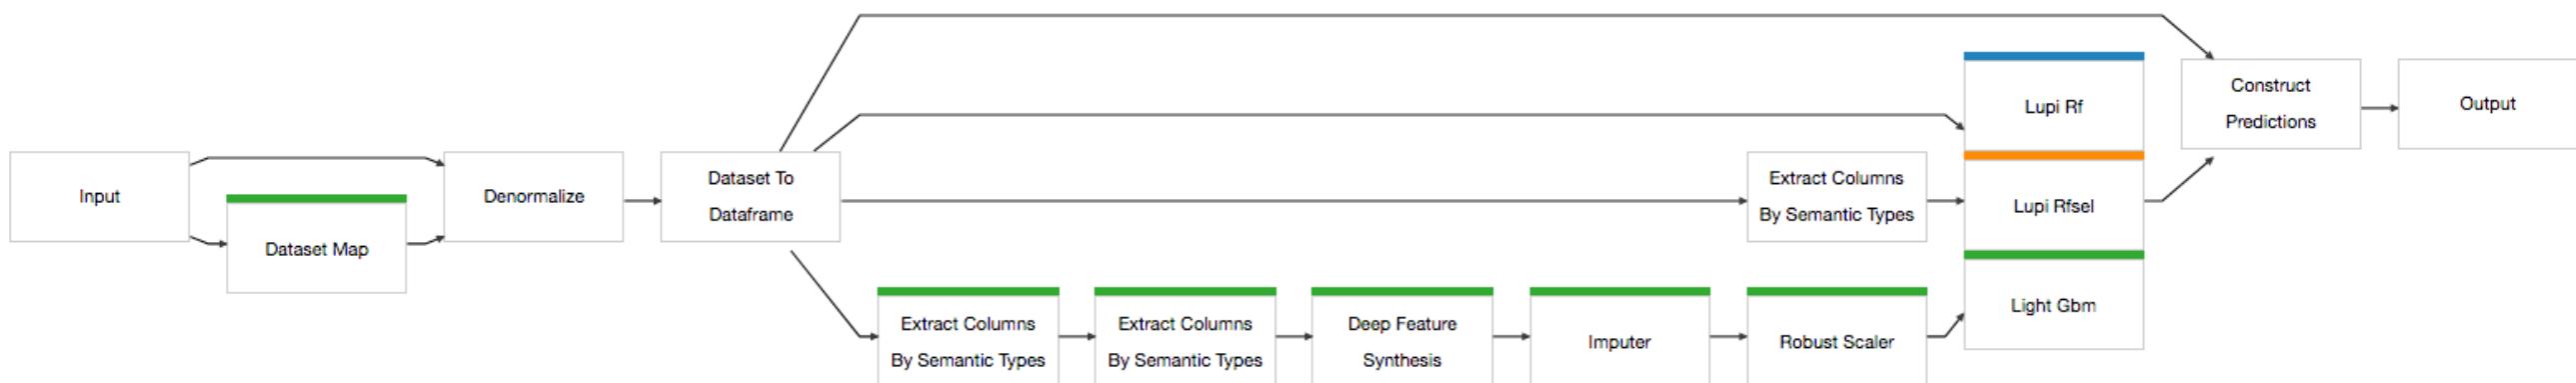
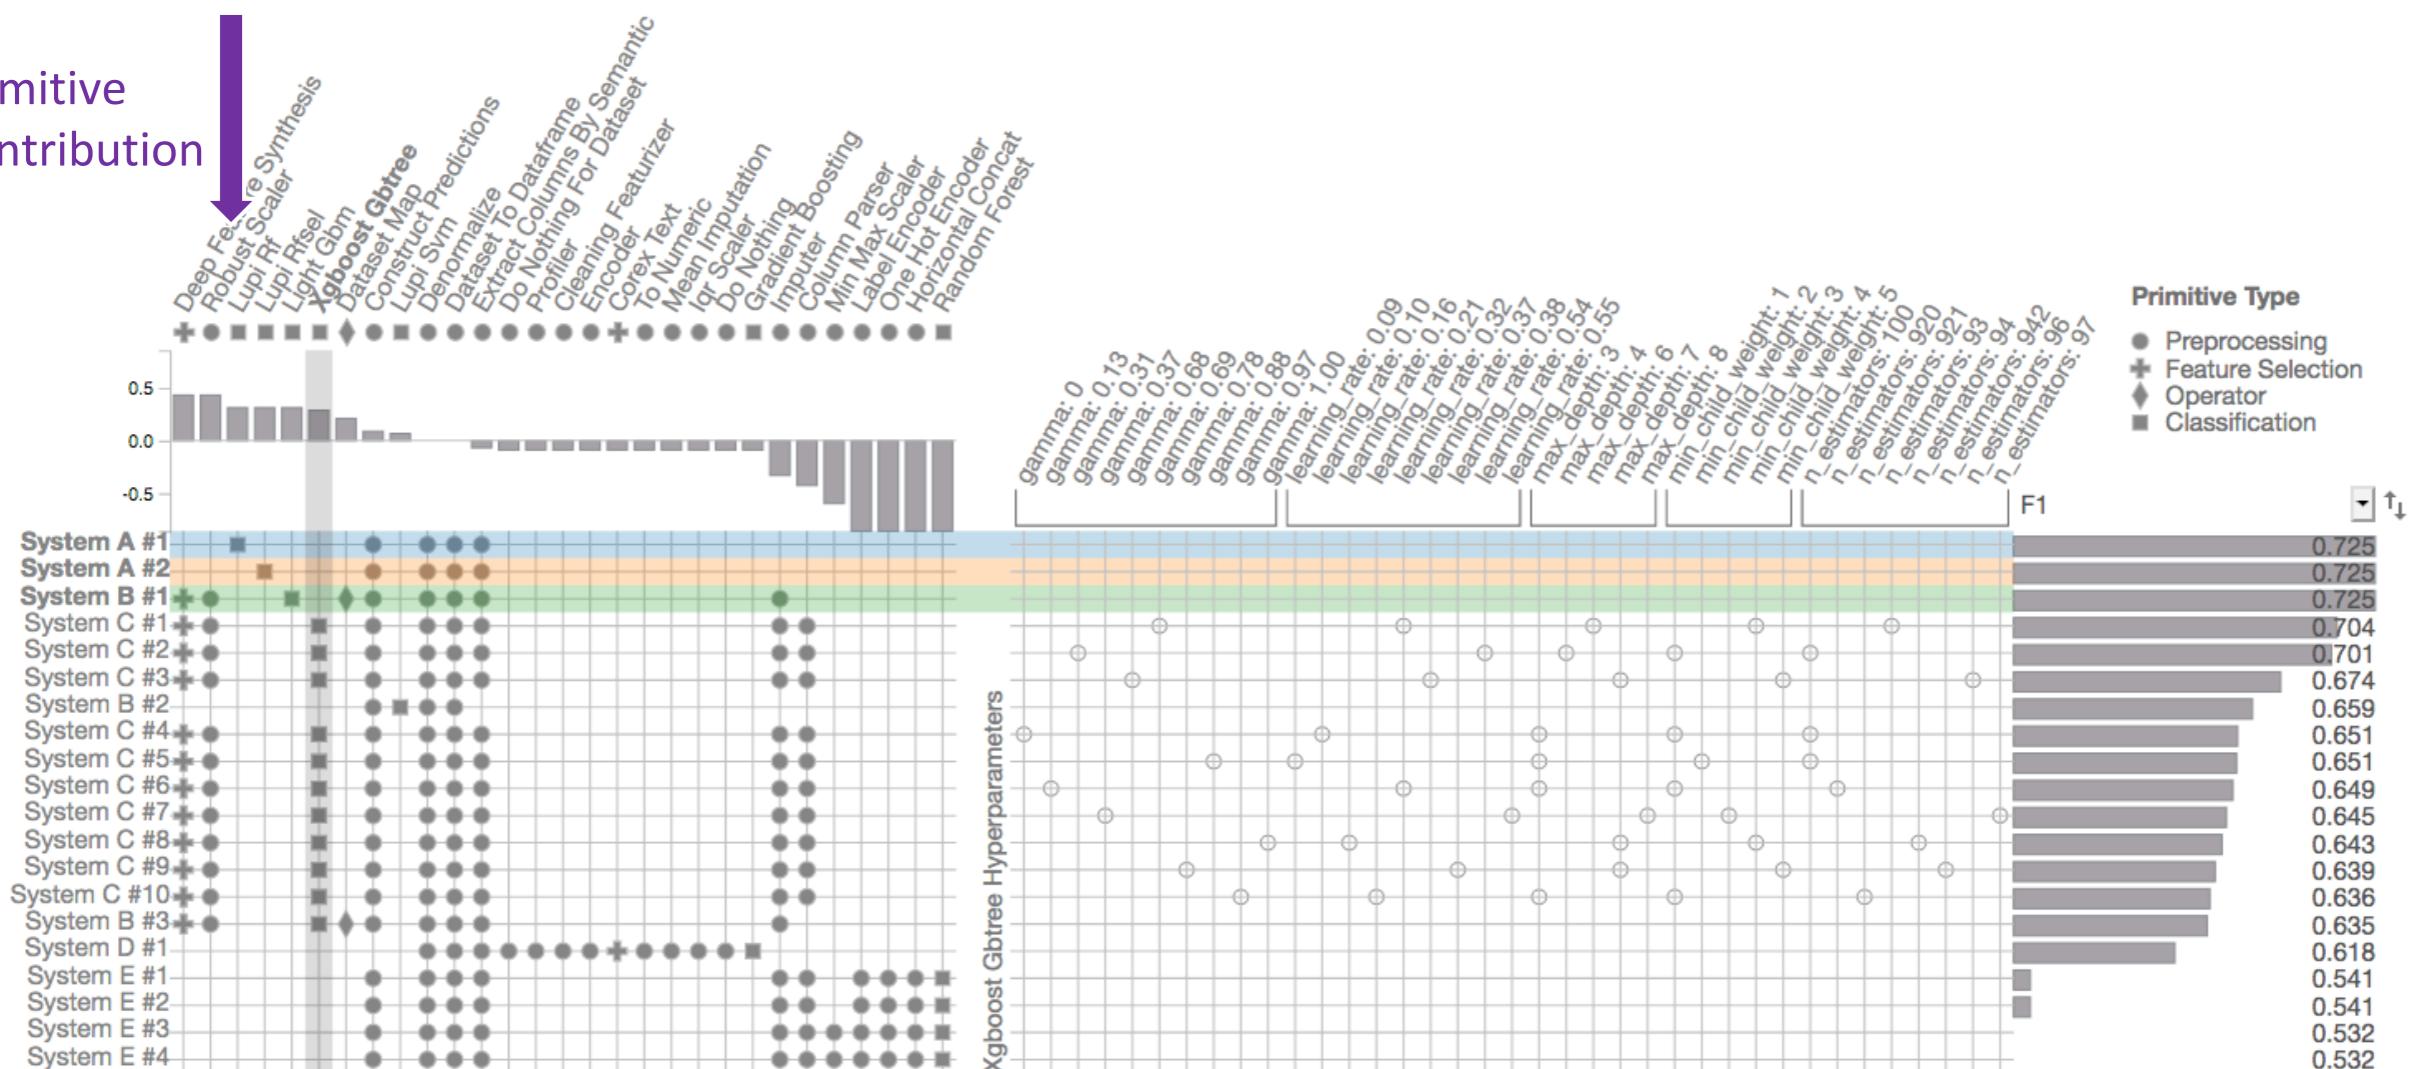


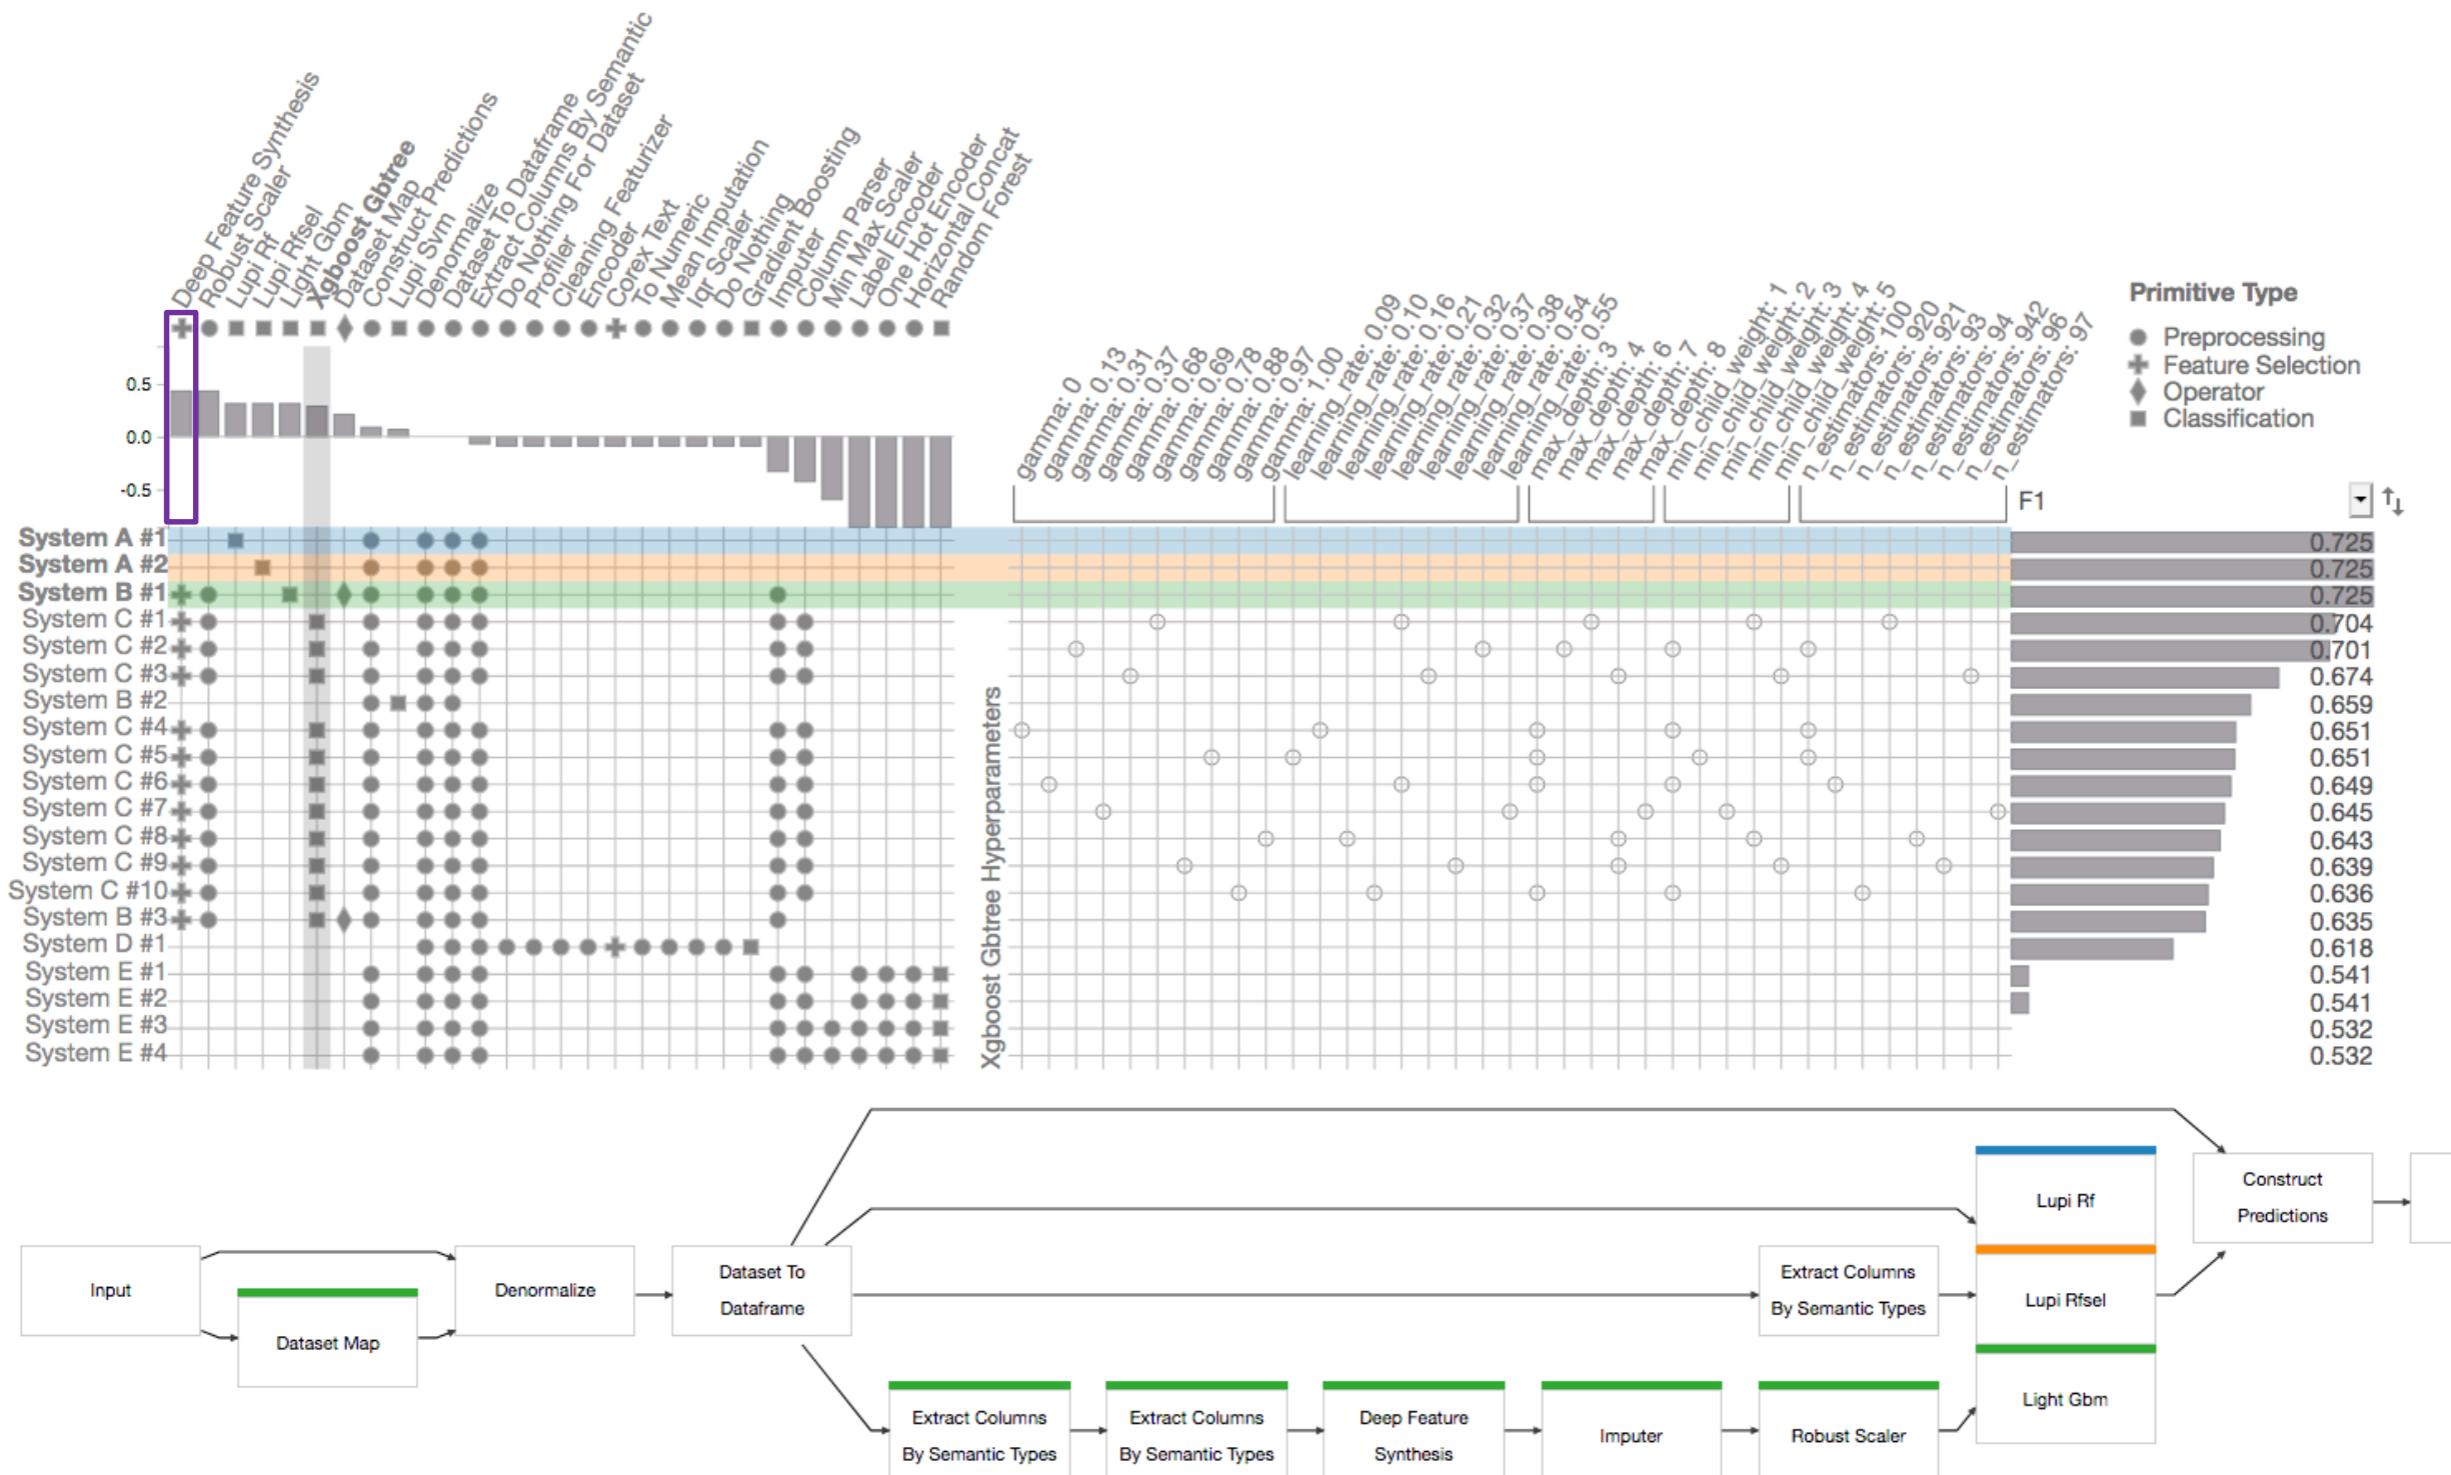


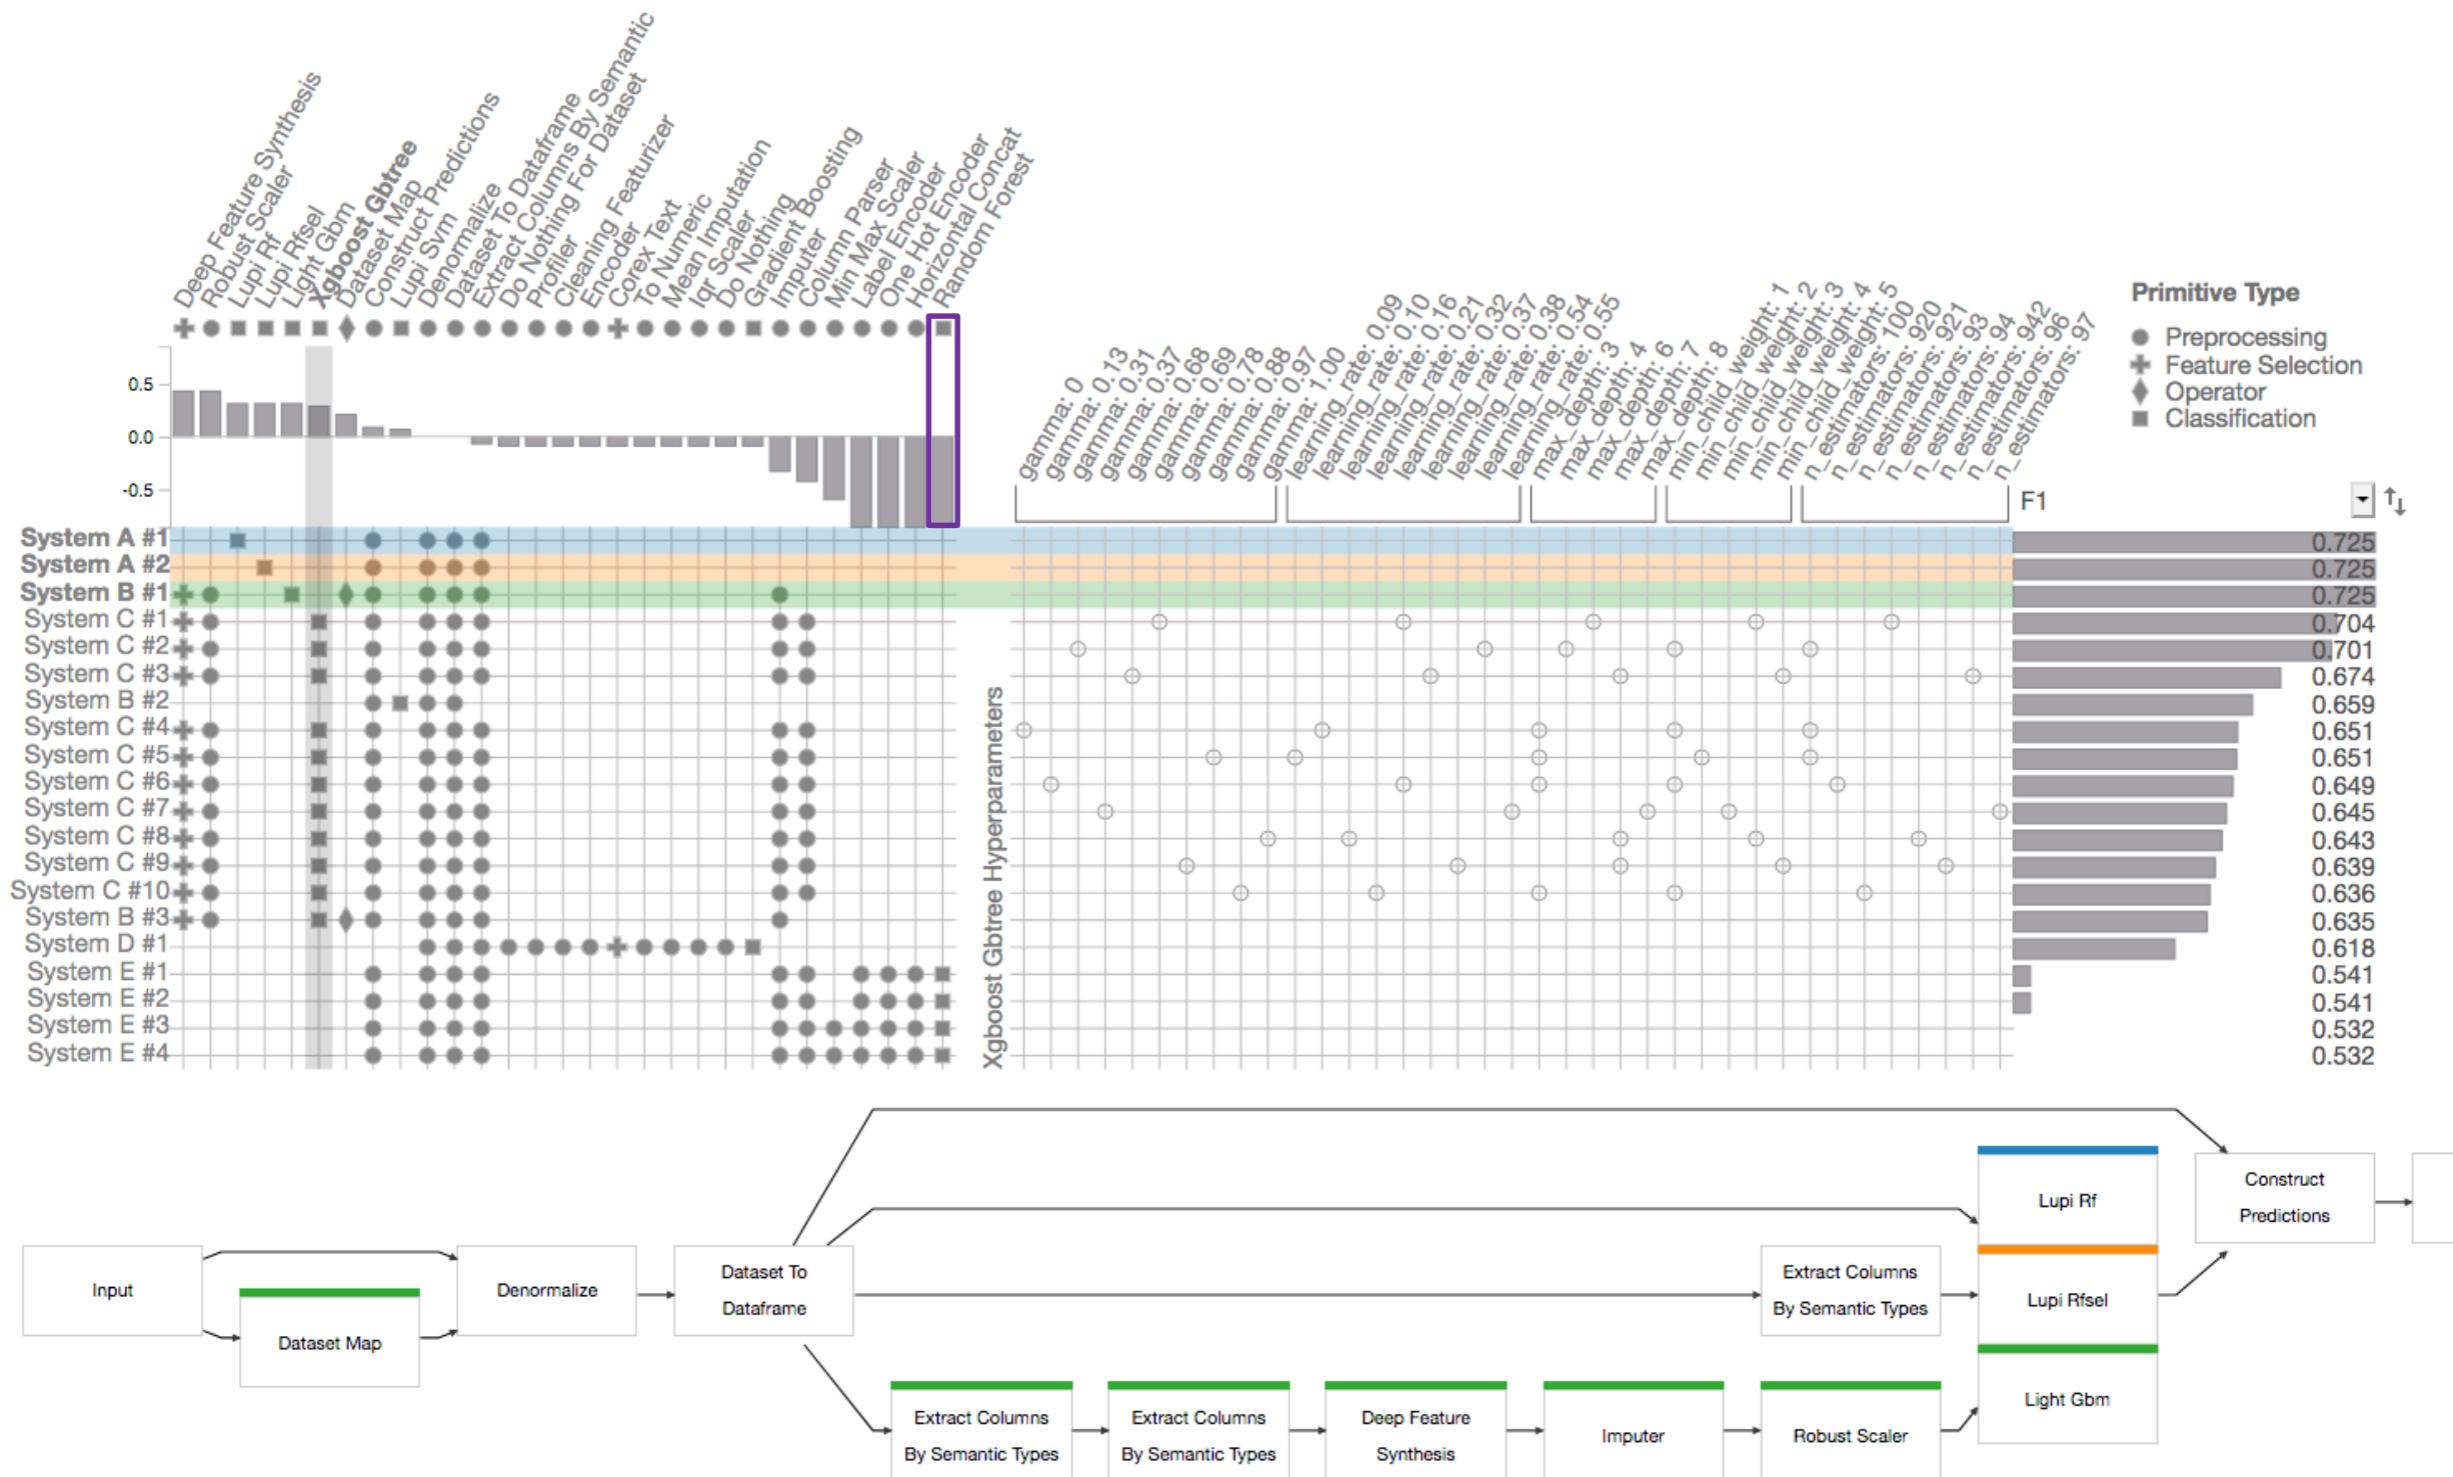




Primitive Contribution







Primitive Contribution

Given a primitive p , and the pipeline scores S :

1. We compute the primitive indicator vector P , where $P_i = 1$ if p is used in pipeline i , and 0 otherwise.
2. The contribution of primitive p to the set of pipelines is given by

$$\text{Contribution} = \text{Corr}(P, S)$$

71

Primitive Contribution is useful for guiding the pipeline analysis.
However, it does not convey primitive interactions.

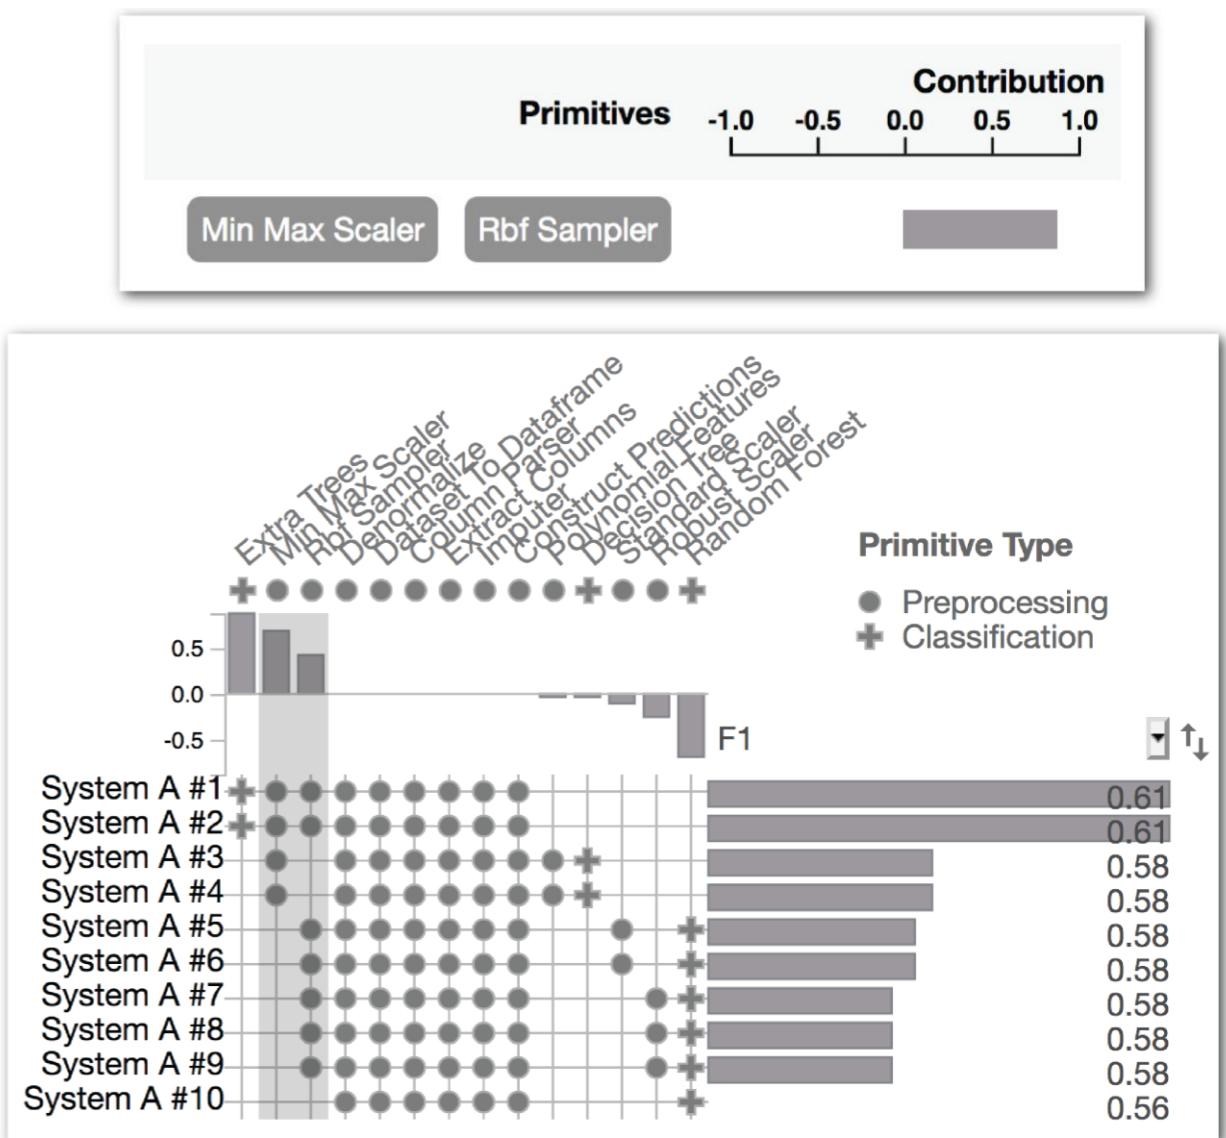
Combined Primitive Contribution

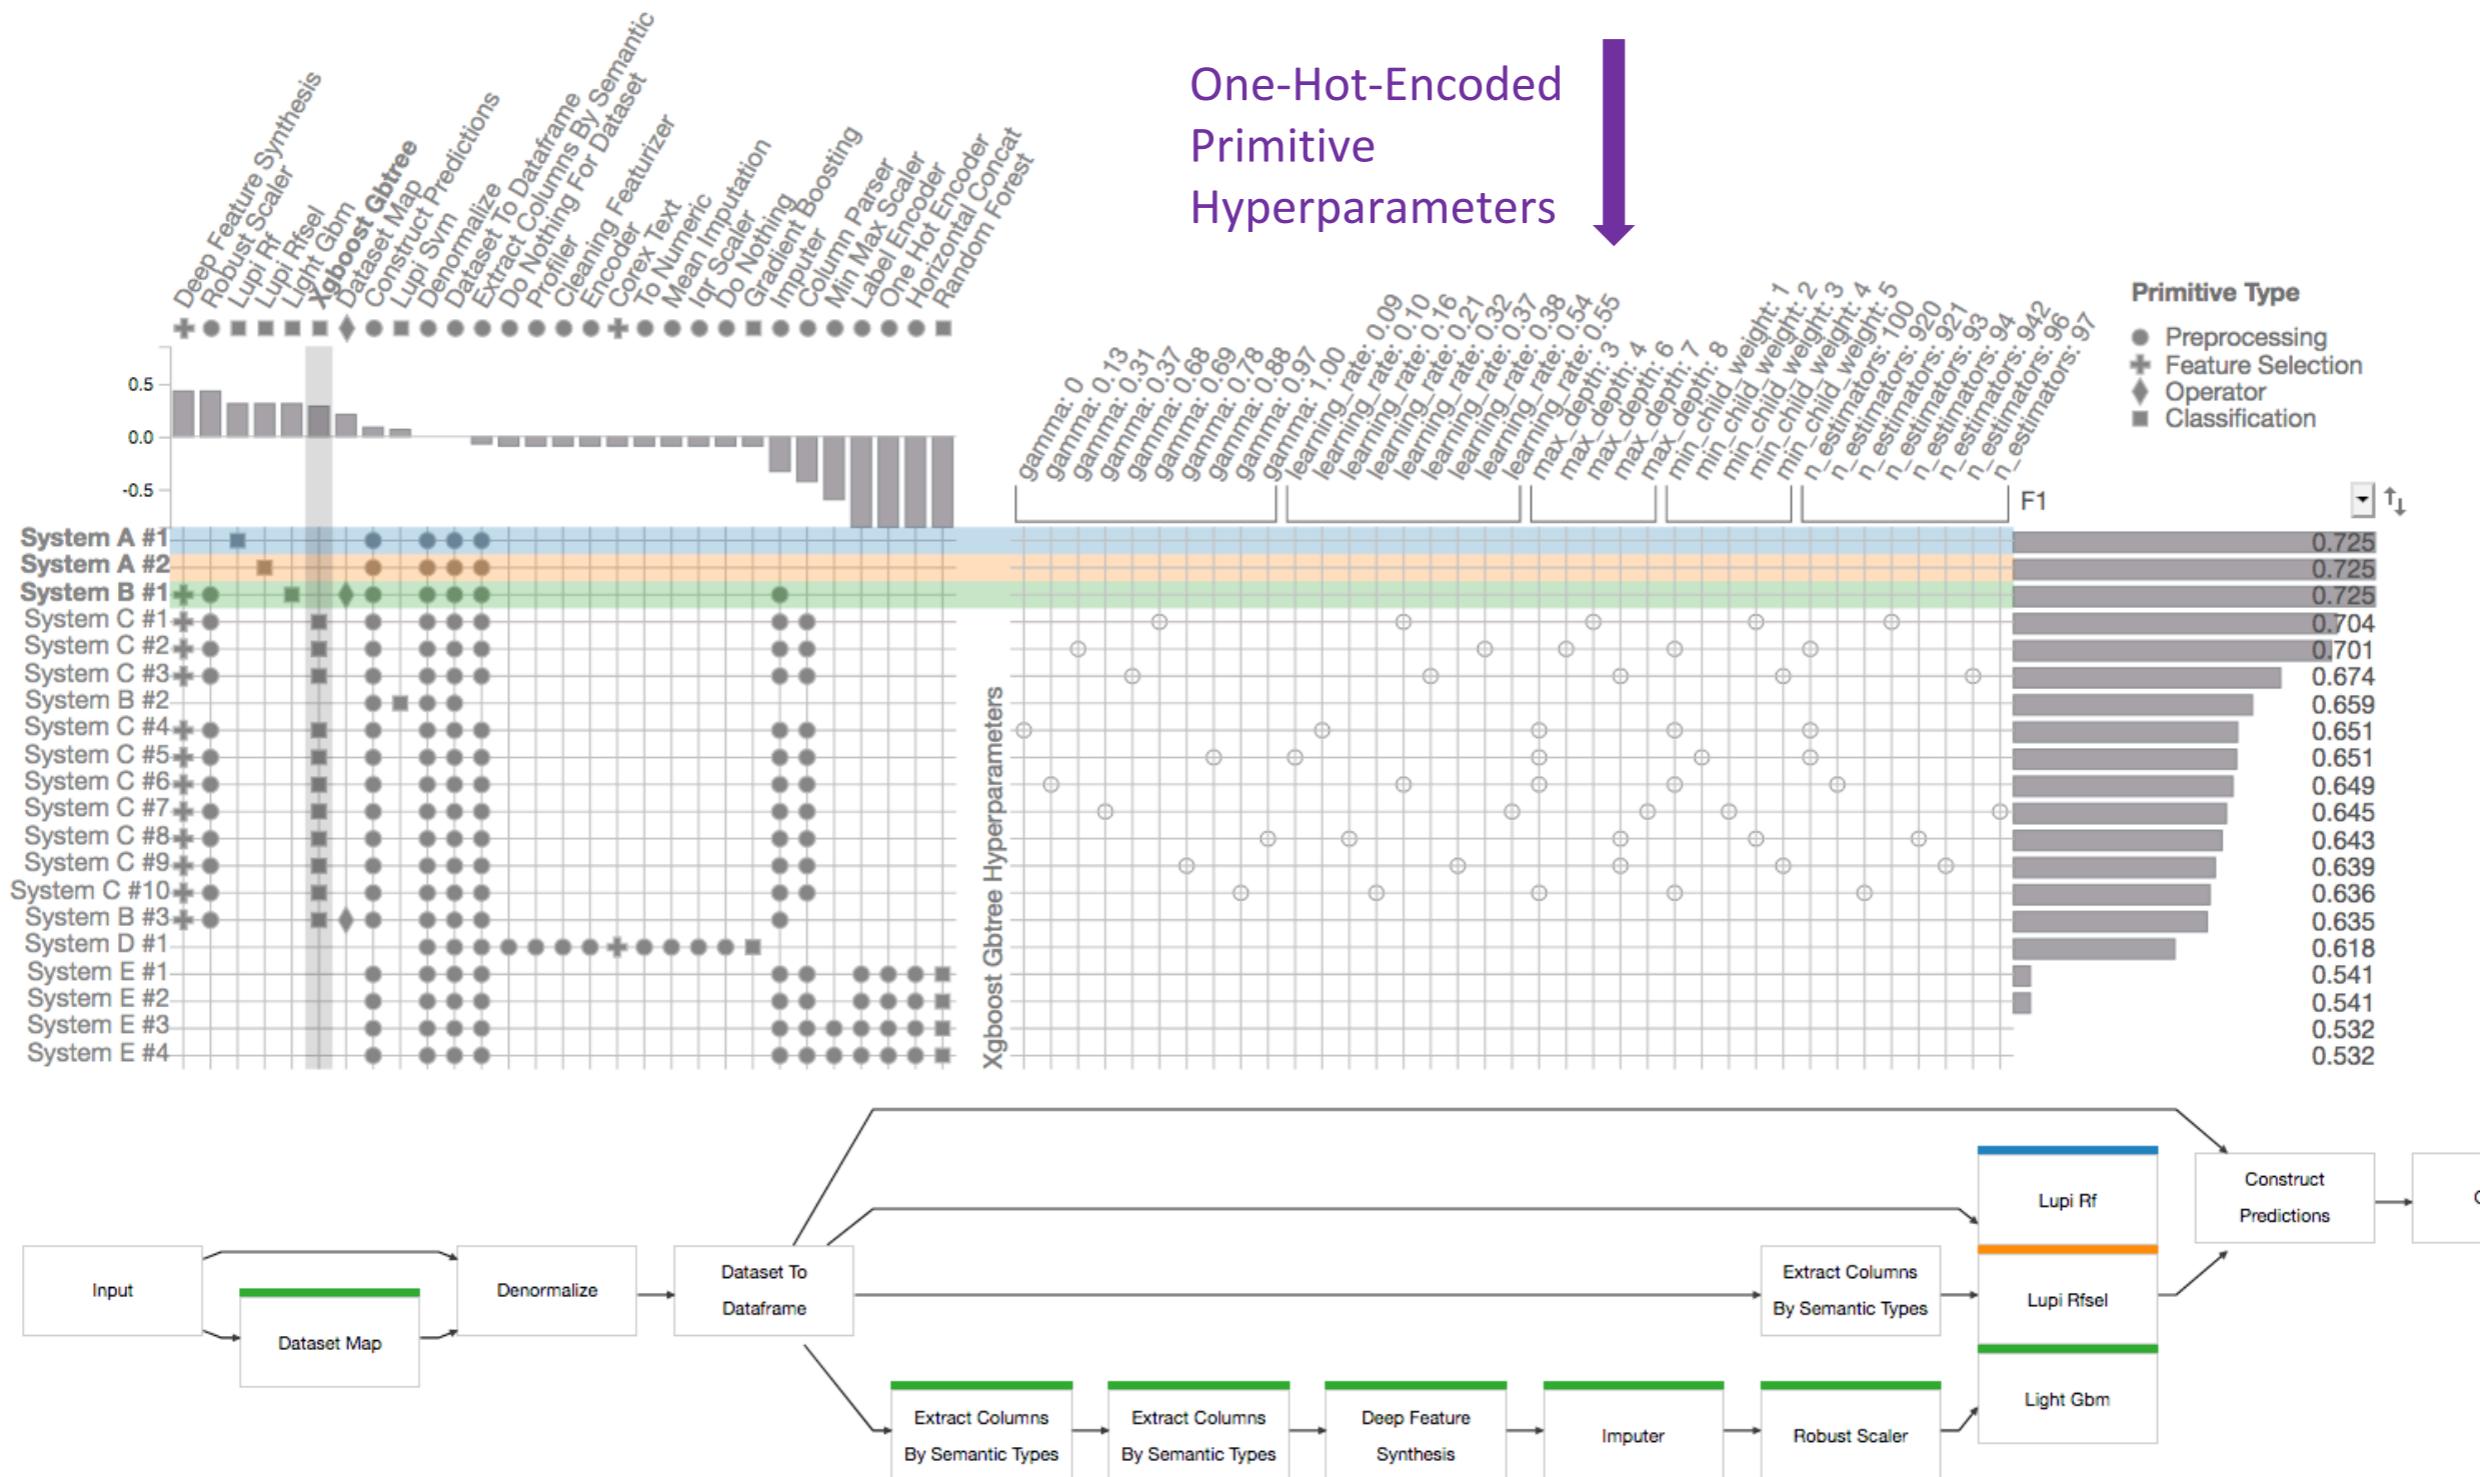
Idea: Investigate correlations of scores with the powerset of primitives

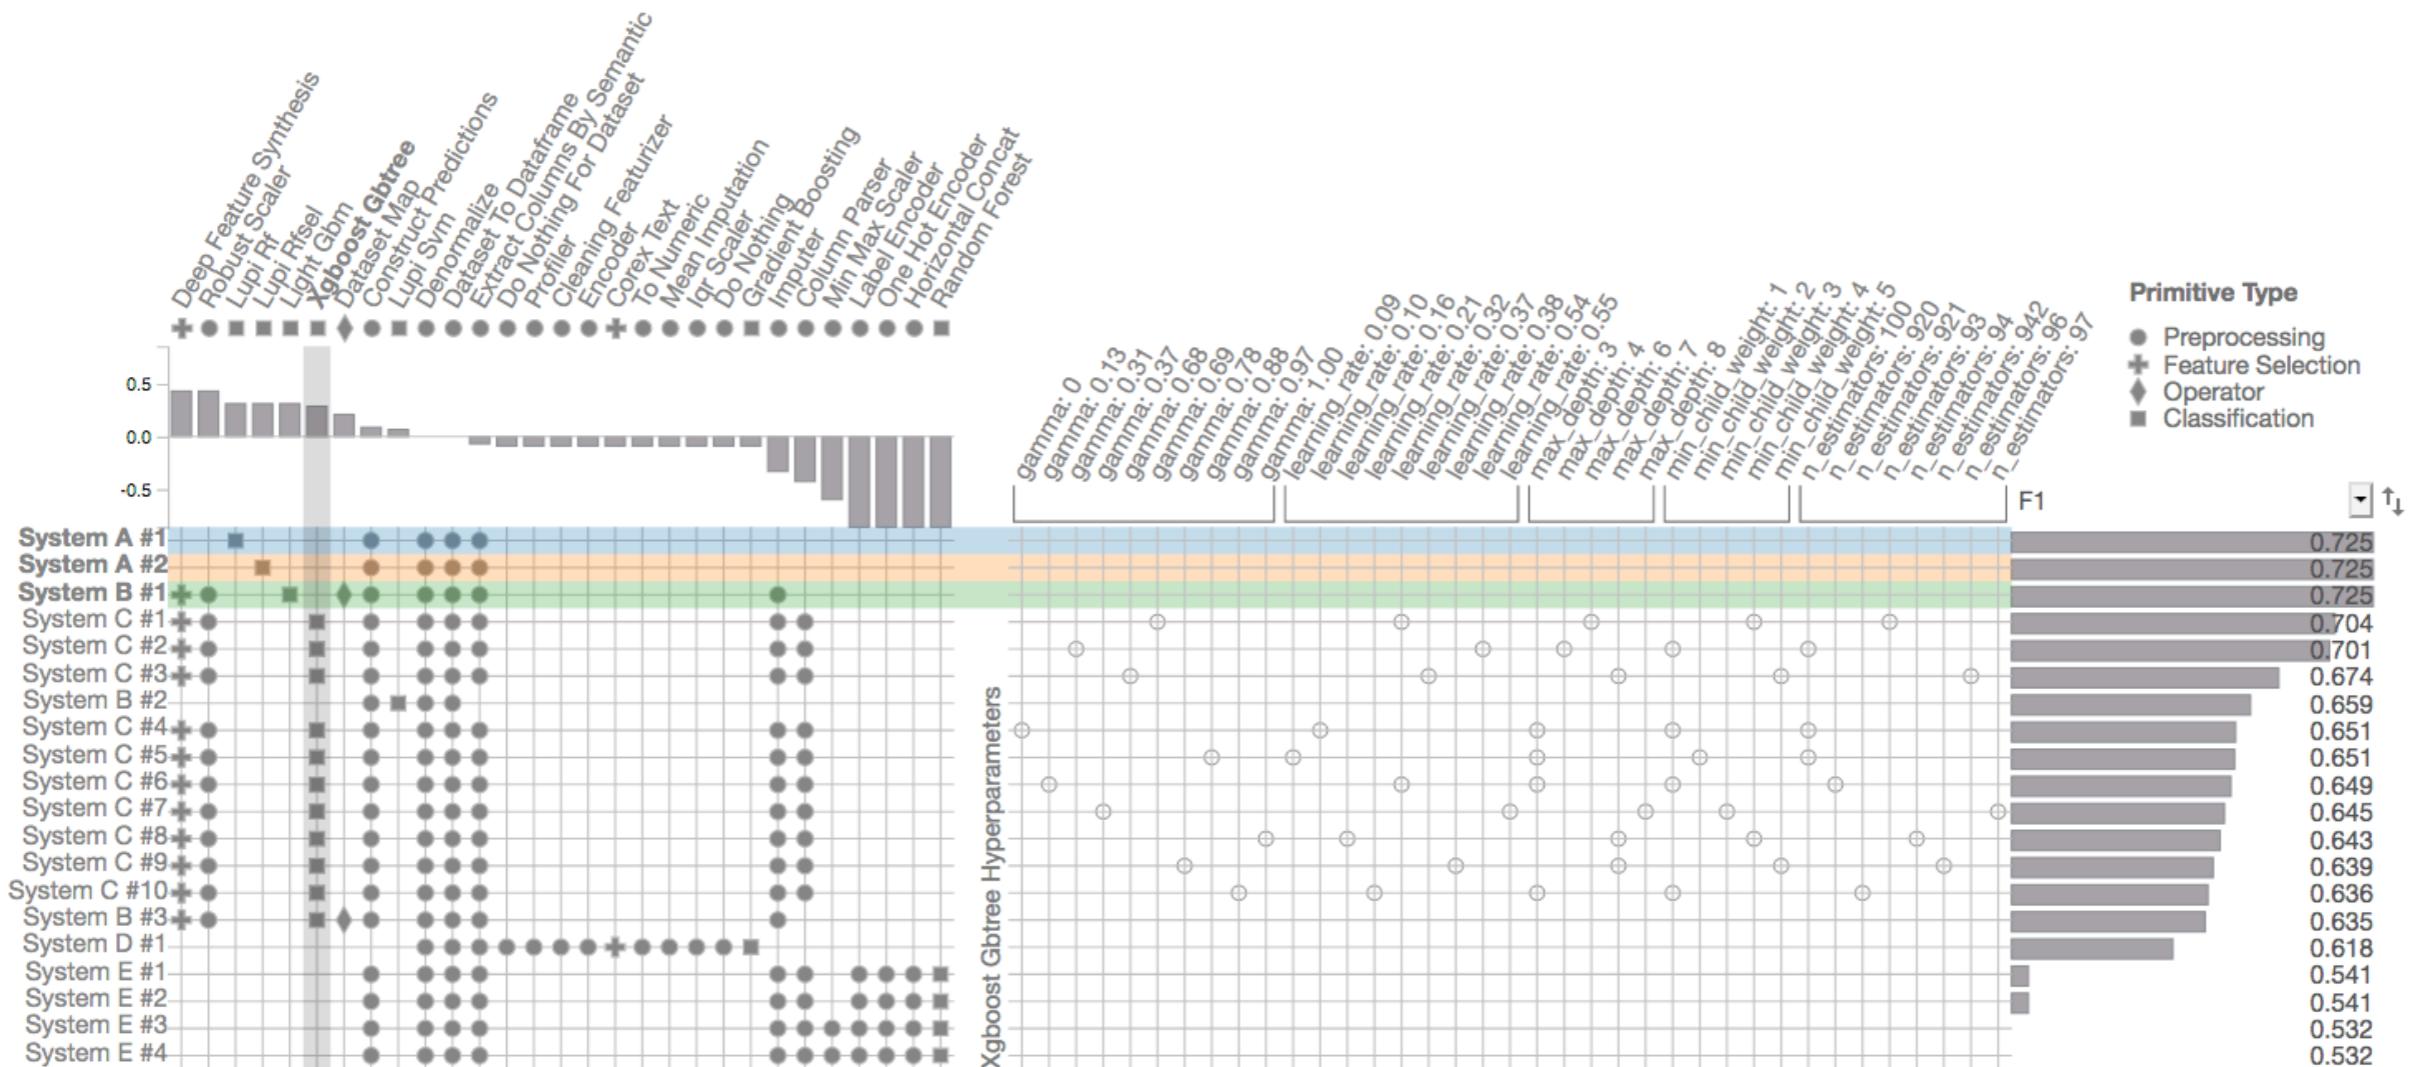
Only show surprising results.

Given $g \in 2^P$

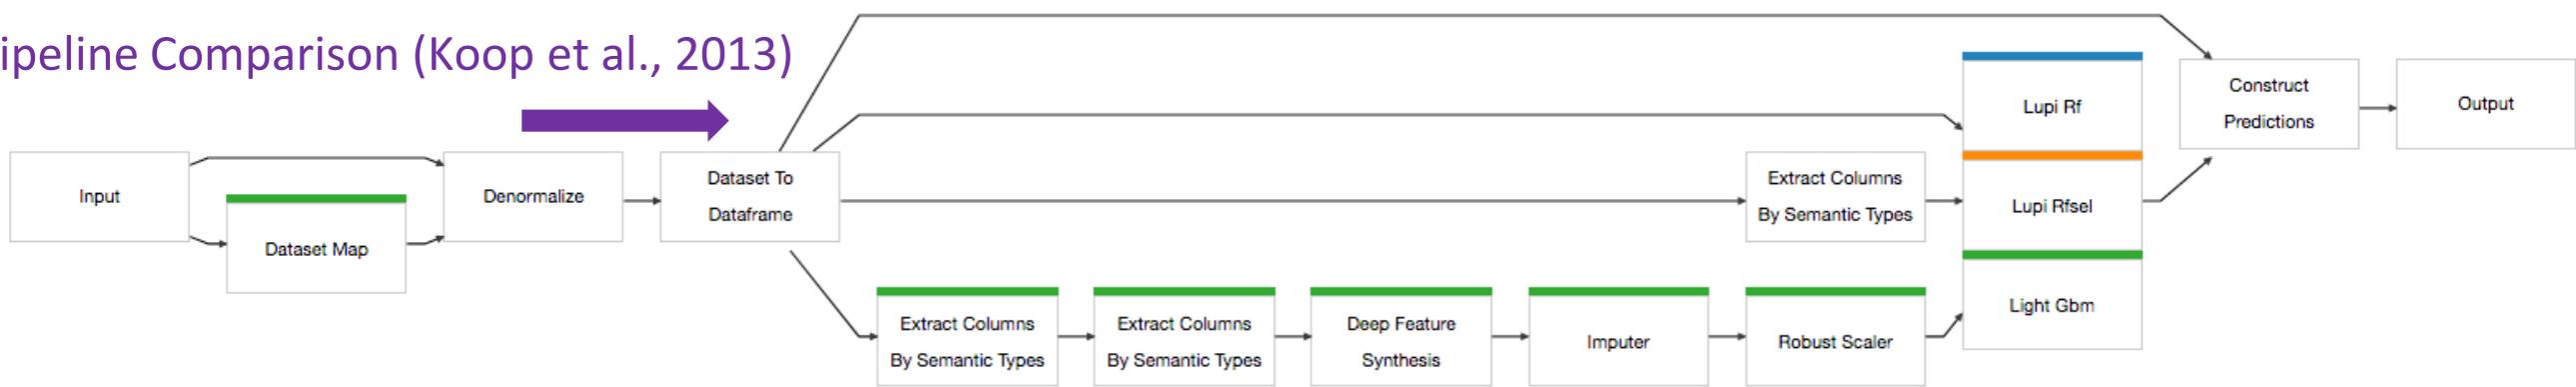
$$\text{Corr}(g, s) > \text{Corr}(g', s) \forall g' \in 2^g$$







Pipeline Comparison (Koop et al., 2013)



PipelineProfiler

Challenges Addressed:

- Requires programming background
- Black box on top of Black box (hard to explain)
- No user feedback (how to adapt the models based on human knowledge)

Limitation:

- User needs Python background

Lab – Pipeline Profiler & Auto-Sklearn

Material

https://colab.research.google.com/drive/1_2FRIkHNFGOiiJt-n_3zuh8vpSMLhwzx?usp=sharing